



Document Identifier: DSP0266

Date: 2025-05-01

Version: 1.22.1

# Redfish Specification

Supersedes: 1.22.0

Document Class: Normative

Document Status: Published

Document Language: en-US

**Copyright Notice**

Copyright © 2015-2025 DMTF. All rights reserved.

DMTF is a not-for-profit association of industry members dedicated to promoting enterprise and systems management and interoperability. Members and non-members may reproduce DMTF specifications and documents, provided that correct attribution is given. As DMTF specifications may be revised from time to time, the particular version and release date should always be noted.

Implementation of certain elements of this standard or proposed standard may be subject to third party patent rights, including provisional patent rights (herein "patent rights"). DMTF makes no representations to users of the standard as to the existence of such rights, and is not responsible to recognize, disclose, or identify any or all such third party patent right, owners or claimants, nor for any incomplete or inaccurate identification or disclosure of such rights, owners or claimants. DMTF shall have no liability to any party, in any manner or circumstance, under any legal theory whatsoever, for failure to recognize, disclose, or identify any such third party patent rights, or for such party's reliance on the standard or incorporation thereof in its product, protocols or testing procedures. DMTF shall have no liability to any party implementing such standard, whether such implementation is foreseeable or not, nor to any patent owner or claimant, and shall have no liability or responsibility for costs or losses incurred if a standard is withdrawn or modified after publication, and shall be indemnified and held harmless by any party implementing the standard from any and all claims of infringement by a patent owner for such implementations.

For information about patents held by third-parties which have notified DMTF that, in their opinion, such patent may relate to or impact implementations of DMTF standards, visit <http://www.dmtf.org/about/policies/disclosures.php>.

This document's normative language is English. Translation into other languages is permitted.

## CONTENTS

Foreword . . . . .	12
Acknowledgments . . . . .	12
Introduction . . . . .	15
1 Scope . . . . .	16
2 Normative references . . . . .	17
3 Terms, definitions, symbols, and abbreviated terms . . . . .	19
3.1 Hardware terms . . . . .	19
3.1.1 baseboard management controller (BMC) . . . . .	19
3.1.2 IPMI . . . . .	19
3.1.3 KVM-IP . . . . .	20
3.1.4 NIC . . . . .	20
3.1.5 PCI . . . . .	20
3.1.6 PCIe . . . . .	20
3.2 Web development terms . . . . .	20
3.2.1 CORS . . . . .	20
3.2.2 CRUD . . . . .	20
3.2.3 CSRF . . . . .	20
3.2.4 HTTP . . . . .	20
3.2.5 HTTPS . . . . .	20
3.2.6 hypermedia API . . . . .	21
3.2.7 IP . . . . .	21
3.2.8 JSON . . . . .	21
3.2.9 message . . . . .	21
3.2.10 OData . . . . .	21
3.2.11 OData service document . . . . .	21
3.2.12 operation . . . . .	21
3.2.13 parent resource . . . . .	21
3.2.14 property . . . . .	21
3.2.15 request . . . . .	22
3.2.16 response . . . . .	22
3.2.17 subscription . . . . .	22
3.2.18 TCP . . . . .	22
3.2.19 TLS . . . . .	22
3.2.20 XSS . . . . .	22
3.3 Redfish terms . . . . .	22
3.3.1 action . . . . .	22
3.3.2 collection . . . . .	22
3.3.3 event . . . . .	22
3.3.4 excerpt . . . . .	23
3.3.5 member . . . . .	23
3.3.6 Redfish client . . . . .	23

3.3.7 Redfish protocol . . . . .	23
3.3.8 Redfish schema . . . . .	23
3.3.9 Redfish service . . . . .	23
3.3.10 resource . . . . .	23
3.3.11 resource collection . . . . .	23
3.3.12 resource tree . . . . .	24
3.3.13 resource type . . . . .	24
3.3.14 service root . . . . .	24
3.3.15 subordinate resource . . . . .	24
3.3.16 task . . . . .	24
3.3.17 task monitor . . . . .	24
4 Typographical conventions . . . . .	25
5 Overview . . . . .	26
5.1 Goals . . . . .	26
5.2 Design tenets . . . . .	27
5.3 Limitations . . . . .	27
5.4 Additional design background and rationale . . . . .	28
5.4.1 REST-based interface . . . . .	28
5.4.2 Data-oriented . . . . .	28
5.4.3 Separation of protocol from data model . . . . .	28
5.4.4 Hypermedia API service root . . . . .	28
5.4.5 OpenAPI v3.0 support . . . . .	28
5.4.6 OData conventions . . . . .	29
5.5 Service elements . . . . .	29
5.5.1 Synchronous and asynchronous operation support . . . . .	29
5.5.2 Eventing mechanism . . . . .	29
5.5.3 Actions . . . . .	30
5.5.4 Service discovery . . . . .	30
5.5.5 Remote access support . . . . .	30
5.6 Security . . . . .	30
6 Protocol details . . . . .	31
6.1 Universal Resource Identifiers . . . . .	31
6.2 HTTP methods . . . . .	33
6.3 HTTP redirect . . . . .	34
6.4 Media types . . . . .	34
6.5 ETags . . . . .	34
6.6 Protocol version . . . . .	36
6.7 Redfish-defined URIs and relative reference rules . . . . .	37
7 Service requests . . . . .	39
7.1 Request headers . . . . .	39
7.2 GET (read requests) . . . . .	43
7.2.1 GET (read requests) overview . . . . .	43
7.2.2 Resource collection requests . . . . .	43

7.2.3 Service root request .....	44
7.2.4 OData service and metadata document requests.....	44
7.3 Query parameters .....	44
7.3.1 Query parameter overview .....	44
7.3.2 The \$expand query parameter .....	48
7.3.3 The \$select query parameter .....	50
7.3.4 The \$filter query parameter.....	51
7.4 HEAD.....	52
7.5 Data modification requests .....	53
7.5.1 Data modification requests overview .....	53
7.5.2 Modification success responses .....	53
7.5.3 Modification error responses.....	54
7.6 PATCH (update) .....	54
7.7 PATCH on array properties .....	55
7.8 PUT (replace) .....	56
7.9 POST (create) .....	56
7.10 DELETE (delete).....	57
7.11 POST (action) .....	57
7.12 Operation apply time .....	60
7.13 Deep operations .....	62
8 Service responses .....	68
8.1 Response headers .....	68
8.2 Link header .....	70
8.3 Status codes .....	71
8.4 OData metadata responses.....	73
8.4.1 OData metadata responses overview .....	73
8.4.2 OData \$metadata .....	73
8.4.2.1 Referencing other schemas .....	74
8.4.2.2 Referencing OEM extensions.....	75
8.4.3 OData service document.....	75
8.5 Resource responses .....	76
8.6 Error responses .....	76
9 Data model .....	78
9.1 Resources .....	78
9.2 Resource types .....	78
9.3 Resource collections .....	79
9.4 OEM resources .....	79
9.5 Common data types .....	80
9.5.1 Primitive types.....	80
9.5.2 Enumerations .....	80
9.5.3 Empty string values.....	81
9.5.4 GUID and UUID values.....	81
9.5.5 Date-Time values .....	81

9.5.6 Duration values . . . . .	82
9.5.7 Reference properties . . . . .	83
9.5.8 Non-resource reference properties . . . . .	83
9.5.9 Array properties . . . . .	83
9.5.10 Structured properties . . . . .	84
9.5.11 Message object . . . . .	84
9.5.11.1 Overview . . . . .	84
9.5.11.2 MessageId format . . . . .	85
9.6 Properties . . . . .	86
9.6.1 Properties overview . . . . .	86
9.6.2 Resource identifier (@odata.id) property . . . . .	87
9.6.3 Resource type (@odata.type) property . . . . .	87
9.6.4 Resource ETag (@odata.etag) property . . . . .	88
9.6.5 Resource context (@odata.context) property . . . . .	88
9.6.6 Id . . . . .	88
9.6.7 Name . . . . .	88
9.6.8 Description . . . . .	89
9.6.9 MemberId . . . . .	89
9.6.10 Count (Members@odata.count) property . . . . .	89
9.6.11 Members . . . . .	89
9.6.12 Next link (Members@odata.nextLink) property . . . . .	89
9.6.13 Links . . . . .	90
9.6.13.1 Reference to a related resource . . . . .	90
9.6.13.2 References to multiple related resources . . . . .	90
9.6.14 Actions property . . . . .	91
9.6.14.1 Action representation . . . . .	91
9.6.14.2 Action responses . . . . .	92
9.6.15 Oem . . . . .	92
9.6.16 Status . . . . .	93
9.7 Naming conventions . . . . .	93
9.7.1 Naming rules . . . . .	93
9.7.2 URI naming rules . . . . .	94
9.8 Extending standard resources . . . . .	95
9.8.1 Extending standard resources overview . . . . .	95
9.8.2 OEM property format and content . . . . .	95
9.8.3 OEM-specified object naming . . . . .	96
9.8.4 OEM resource types . . . . .	97
9.8.5 OEM registries . . . . .	97
9.8.6 OEM URIs . . . . .	97
9.8.7 OEM property examples . . . . .	98
9.8.8 OEM actions . . . . .	99
9.9 Payload annotations . . . . .	100
9.9.1 Payload annotations overview . . . . .	100

9.9.2 Allowable values for strings . . . . .	101
9.9.3 Allowable patterns for string values . . . . .	101
9.9.4 Allowable values for numbers and durations . . . . .	101
9.9.5 Extended information . . . . .	102
9.9.5.1 Extended object information . . . . .	102
9.9.5.2 Extended property information . . . . .	103
9.9.5.3 Extended information implementation notes . . . . .	104
9.9.6 Action info annotation . . . . .	104
9.9.7 Settings and settings apply time annotations . . . . .	105
9.9.8 Operation apply time and operation apply time support annotations . . . . .	105
9.9.9 Maintenance window annotation . . . . .	105
9.9.10 Collection capabilities annotation . . . . .	106
9.9.11 Requested count and allow over-provisioning annotations . . . . .	108
9.9.12 Zone affinity annotation . . . . .	109
9.9.13 Supported certificates annotation . . . . .	109
9.9.14 Deprecated annotation . . . . .	109
9.9.15 Writable properties annotation . . . . .	110
9.10 Settings resource . . . . .	111
9.11 Special resource situations . . . . .	113
9.11.1 Overview . . . . .	113
9.11.2 Absent resources . . . . .	113
9.11.3 Transiently unavailable resources . . . . .	114
9.12 Registries . . . . .	114
9.13 Schema annotations . . . . .	115
9.13.1 Schema annotations overview . . . . .	115
9.13.2 Description annotation . . . . .	115
9.13.3 Long description annotation . . . . .	115
9.13.4 Resource capabilities annotation . . . . .	116
9.13.5 Resource URI patterns annotation . . . . .	116
9.13.6 Additional properties annotation . . . . .	117
9.13.7 Permissions annotation . . . . .	117
9.13.8 Required annotation . . . . .	117
9.13.9 Required on create annotation . . . . .	118
9.13.10 Units of measure annotation . . . . .	118
9.13.11 Expanded resource annotation . . . . .	118
9.13.12 Owning entity annotation . . . . .	118
9.13.13 Deprecated annotation . . . . .	118
9.13.14 URI segment annotation . . . . .	119
9.13.15 URI annotation . . . . .	119
9.14 Versioning . . . . .	119
9.15 Localization . . . . .	120
10 File naming and publication . . . . .	121
10.1 Registry file naming . . . . .	121

10.2 Profile file naming .....	121
10.3 Dictionary file naming .....	121
10.4 Localized file naming.....	122
10.5 DMTF Redfish file repository.....	122
11 Schema definition languages.....	124
11.1 OData Common Schema Definition Language .....	124
11.1.1 OData Common Schema Definition Language overview .....	124
11.1.2 File naming conventions for CSDL .....	124
11.1.3 Core CSDL files .....	124
11.1.4 CSDL format .....	125
11.1.4.1 Referencing other CSDL files .....	125
11.1.4.2 CSDL data services .....	126
11.1.5 Elements of CSDL namespaces .....	126
11.1.5.1 Qualified names .....	127
11.1.5.2 Entity type and complex type elements .....	127
11.1.5.3 Action element .....	128
11.1.5.4 Action element for OEM actions .....	129
11.1.5.5 Action with a response body .....	129
11.1.5.6 Property element .....	130
11.1.5.7 Navigation property element.....	131
11.1.5.8 Enum type element.....	131
11.1.5.9 Annotation element.....	132
11.2 JSON Schema.....	135
11.2.1 JSON Schema overview .....	135
11.2.2 File naming conventions for JSON Schema .....	135
11.2.3 Core JSON Schema files .....	136
11.2.4 JSON Schema format .....	136
11.2.5 JSON Schema definitions body .....	137
11.2.5.1 Resource definitions in JSON Schema.....	137
11.2.5.2 Enumerations in JSON Schema.....	138
11.2.5.3 Actions in JSON Schema .....	139
11.2.5.4 OEM actions in JSON Schema.....	140
11.2.5.5 Action with a response body .....	141
11.2.6 JSON Schema terms .....	141
11.3 OpenAPI .....	142
11.3.1 OpenAPI overview .....	142
11.3.2 File naming conventions for OpenAPI schema .....	142
11.3.3 Core OpenAPI schema files .....	143
11.3.4 openapi.yaml.....	143
11.3.5 OpenAPI file format.....	145
11.3.6 OpenAPI components body .....	146
11.3.6.1 Resource definitions in OpenAPI .....	146
11.3.6.2 Enumerations in OpenAPI .....	146

11.3.6.3 Actions in OpenAPI.....	147
11.3.6.4 OEM actions in OpenAPI .....	148
11.3.7 OpenAPI terms used by Redfish.....	149
11.4 Schema modification rules .....	150
12 Service details .....	151
12.1 Eventing.....	151
12.1.1 Eventing overview.....	151
12.1.2 POST to subscription collection .....	151
12.1.3 Open an SSE connection .....	152
12.1.4 EventType-based eventing .....	153
12.1.5 Subscribing to events .....	153
12.1.6 Event formats .....	155
12.1.7 OEM extensions .....	156
12.2 Asynchronous operations .....	156
12.3 Resource tree stability .....	158
12.4 Discovery.....	159
12.4.1 Discovery overview.....	159
12.4.2 UPnP compatibility .....	159
12.4.3 USN format .....	159
12.4.4 M-SEARCH response .....	159
12.4.5 Notify, alive, and shutdown messages .....	160
12.5 Server-sent events .....	160
12.5.1 General .....	160
12.5.2 Event service.....	161
12.5.2.1 Event message SSE stream.....	164
12.5.2.2 Metric report SSE stream .....	164
12.6 Update service.....	165
12.6.1 Overview .....	165
12.6.2 Software update types .....	165
12.6.2.1 Simple updates.....	166
12.6.2.2 Multipart HTTP push updates .....	166
12.7 Outbound connections.....	168
12.7.1 Overview .....	168
12.7.2 Establishing an outbound connection .....	168
12.7.3 MTLS for an outbound connection .....	169
12.7.4 Handling Redfish requests over an outbound connection.....	169
12.7.5 Closing an outbound connection.....	170
13 Security details.....	171
13.1 Transport Layer Security (TLS) protocol .....	171
13.1.1 Transport Layer Security (TLS) protocol overview .....	171
13.1.2 Cipher suites .....	171
13.1.3 Certificates .....	172
13.2 Sensitive data .....	172

13.3 Authentication . . . . .	172
13.3.1 Authentication overview . . . . .	172
13.3.2 Authentication requirements . . . . .	173
13.3.2.1 Resource and operation authentication requirements . . . . .	173
13.3.2.2 HTTP header authentication requirements . . . . .	173
13.3.2.3 Authentication failure requirements . . . . .	173
13.3.3 HTTP Basic authentication . . . . .	174
13.3.4 Redfish session login authentication . . . . .	174
13.3.4.1 Redfish login sessions . . . . .	174
13.3.4.2 Session login . . . . .	175
13.3.4.3 Session lifetime . . . . .	176
13.3.4.4 Session termination or logout . . . . .	176
13.3.5 Client certificate authentication . . . . .	177
13.4 Authorization . . . . .	177
13.4.1 Authorization overview . . . . .	177
13.4.2 Privilege model . . . . .	178
13.4.2.1 Roles . . . . .	178
13.4.2.2 Restricted roles and restricted privileges . . . . .	179
13.4.2.3 OEM privileges . . . . .	180
13.4.3 Redfish service operation-to-privilege mapping . . . . .	180
13.4.3.1 Why specify operation-to-privilege mapping? . . . . .	181
13.4.3.2 Representing operation-to-privilege mappings . . . . .	181
13.4.3.3 Operation map syntax . . . . .	181
13.4.3.4 Mapping overrides syntax . . . . .	182
13.4.3.5 Property override example . . . . .	183
13.4.3.6 Subordinate override . . . . .	184
13.4.3.7 Resource URI override . . . . .	185
13.4.3.8 Privilege AND and OR syntax . . . . .	186
13.4.4 Delegated authorization with OAuth 2.0 . . . . .	187
13.4.4.1 OAuth 2.0 overview . . . . .	187
13.4.4.2 OAuth 2.0 data model requirements . . . . .	187
13.4.4.3 OAuth 2.0 access tokens . . . . .	187
13.4.4.4 Redfish OAuth2.0 scope usage . . . . .	189
13.5 Account service . . . . .	190
13.5.1 Account service overview . . . . .	190
13.5.2 Password management . . . . .	190
13.5.3 Atomic password changes . . . . .	190
13.5.4 Password change required handling . . . . .	190
13.5.5 Time-based One-Time Password secret key handling . . . . .	191
13.6 Asynchronous tasks . . . . .	192
13.7 Event subscriptions . . . . .	192
14 Redfish Host Interface . . . . .	193
15 Redfish composability . . . . .	194

15.1 Composition requests . . . . .	195
15.1.1 Composition requests overview . . . . .	195
15.1.2 Specific composition . . . . .	195
15.1.3 Constrained composition . . . . .	196
15.1.4 Expandable resources . . . . .	197
15.2 Updating a composed resource. . . . .	197
16 Aggregation . . . . .	198
16.1 Classes of aggregators . . . . .	198
16.1.1 Implicit and complex aggregators . . . . .	198
16.1.2 Use cases . . . . .	199
16.2 Aggregation service . . . . .	199
16.2.1 Aggregation service overview . . . . .	199
16.2.2 Aggregator requirements . . . . .	199
16.2.3 Aggregates . . . . .	200
16.2.4 Aggregation sources and connection methods . . . . .	200
17 ANNEX A (informative) Change log . . . . .	202
18 Bibliography . . . . .	223

# Foreword

---

DMTF's Redfish Forum develops the Redfish standard.

DMTF is a not-for-profit association of industry members that promotes enterprise and systems management and interoperability. For information about DMTF, see [DMTF](#).

This version supersedes version 1.22.0. For a list of the changes, see [ANNEX A \(informative\) Change log](#).

## Acknowledgments

---

DMTF acknowledges the following individuals for their contributions to the Redfish standard, including this document and Redfish schemas, interoperability profiles, and message registries:

- Rafiq Ahamed — Hewlett Packard Enterprise
- Richelle Ahlvers — Intel Corporation, Broadcom Inc.
- Jeff Autor — Hewlett Packard Enterprise
- David Black — Dell Technologies
- Jeff Bobzin — Insyde Software Corp.
- Patrick Boyd — Dell Technologies
- David Brockhaus — Vertiv
- Richard Brunner — VMware Inc.
- Sean Byland — Hewlett Packard Enterprise
- Lee Calcote — Seagate Technology
- Keith Campbell — Lenovo
- Derek Chan — Google LLC
- P Chandrasekhar — Dell Technologies
- Barbara Craig — Hewlett Packard Enterprise
- Chris Davenport — Hewlett Packard Enterprise
- Gamma Dean — Vertiv
- Michael Du — Huawei Technologies Co., Ltd.
- Daniel Dufresne — Dell Technologies
- Samer El-Haj-Mahmoud — Arm Limited, Lenovo, Hewlett Packard Enterprise
- George Ericson — Dell Technologies
- Wassim Fayed — Microsoft Corporation
- Kevin Ferguson — Vertiv
- Mike Garrett — Hewlett Packard Enterprise
- Steve Geffin — Vertiv

- Martin Halstead — Hewlett Packard Enterprise
- Joe Handzik — Hewlett Packard Enterprise
- Jon Hass — Dell Technologies
- Jeff Hilland — Hewlett Packard Enterprise
- Blake Hilliard — Hewlett Packard Enterprise
- Chris Hoffman — Vertiv
- Cactus Jiang — Vertiv
- Barry Kittner — Intel Corporation
- Steven Krig — Intel Corporation
- Jennifer Lee — Intel Corporation
- John Leung — Intel Corporation
- Magnus Lundmark — Ericsson AB
- Steve Lyle — Hewlett Packard Enterprise
- Gunnar Mills — IBM
- Jagan Molleti — Dell Technologies
- Milena Natanov — Microsoft Corporation
- Balaji Natrajan — Microchip Technology Inc., Hewlett Packard Enterprise
- Scott Phuong — Cisco Systems, Inc.
- Michael Pizzo — Microsoft Corporation
- Chris Poblete — Dell Technologies
- Slawek Putyrski — Intel Corporation
- Michael Raineri — Dell Technologies
- Joseph Reynolds — IBM
- Irina Salvan — Microsoft Corporation
- Joseph-Jonathan Salzano — HP, Inc
- Bill Scherer — Hewlett Packard Enterprise
- Geoff Schunicht — Hewlett Packard Enterprise
- Abhirup Seal — Dell Technologies
- Hemal Shah — Broadcom Inc.
- Jim Shelton — Vertiv
- Tom Slight — Intel Corporation
- Josiah Smith — Eaton
- Donnie Sturgeon — Vertiv
- Pawel Szymanski — Intel Corporation
- Ed Tanous — Google LLC
- Willy Tu — Google LLC
- Paul Vancil — Dell Technologies

- Ganesh Viswanathan — Dell Technologies
- Claire Weinan — Google LLC
- Joseph White — Dell Technologies
- Linda Wu — NVIDIA Corporation, Super Micro Computer, Inc.
- Justin York — Hewlett Packard Enterprise

# Introduction

---

Redfish is a standard that uses RESTful interface semantics to access a schema-based data model to conduct management operations. It is suitable for a wide range of devices, from stand-alone servers to composable infrastructures and large-scale cloud environments.

The initial Redfish scope targeted servers. DMTF and its alliance partners expanded that scope to cover most data center IT equipment and other solutions, and both in- and out-of-band access methods.

Additionally, DMTF and other organizations that use Redfish as part of their industry standard or solution have added educational material.

This document defines the RESTful interface protocol and the various concepts and services necessary to implement a Redfish interface. The definition of the schema based data model and standard messages for the Redfish interface are covered separately in the following documents:

- DMTF DSP8010, *Redfish Schema Bundle*, <https://www.dmtf.org/dsp/DSP8010> contains the individual schema definition files in multiple schema description languages.
- DMTF DSP0268, *Redfish Data Model Specification*, <https://www.dmtf.org/dsp/DSP0268> contains the normative descriptions and example payloads for all standard Redfish schema in a single reference guide.
- DMTF DSP8011, *Redfish Standard Registries Bundle*, <https://www.dmtf.org/dsp/DSP8011> contains the message registries used for error reporting and event messages.

# 1 Scope

This specification defines the required protocols, data model, behaviors, and other architectural components for an interoperable, multivendor, remote, and out-of-band capable interface. This interface meets the cloud-based and web-based IT professionals' expectations for scalable platform management. While large and hyperscale environments are the primary focus, clients can use the specification for individual system management.

The specification defines the required elements for all Redfish implementations, and the optional elements that system vendors and manufacturers can choose. This specification also defines at which points an implementation can provide OEM-specific extensions.

The specification sets normative requirements for *Redfish services* and associated materials, such as Redfish schema files. In general, the specification does not set requirements for Redfish clients but indicates how a client can successfully and effectively access and use a Redfish service.

The specification does not require that implementations of the Redfish interfaces and functions require particular hardware or firmware.

## 2 Normative references

The following documents are referred to in the text in such a way that some or all of their content constitutes requirements of this document. For dated references, only the edition cited applies. For undated references, the latest edition of the referenced document (including any amendments) applies.

- DMTF DSP0270, *Redfish Host Interface Specification*, <https://www.dmtf.org/dsp/DSP0270>
- Redfish Schema: RedfishExtensions v1.0.0, [https://redfish.dmtf.org/schemas/v1/RedfishExtensions\\_v1.xml](https://redfish.dmtf.org/schemas/v1/RedfishExtensions_v1.xml)
- *Transport Layer Security (TLS) Parameters*, <https://www.iana.org/assignments/tls-parameters/tls-parameters.xhtml>
- *JSON Schema: A Media Type for Describing JSON Documents* *draft-handrews-json-schema-01*, <https://tools.ietf.org/html/draft-handrews-json-schema-01>
- *JSON Schema Validation: A Vocabulary for Structural Validation of JSON* *draft-handrews-json-schema-validation-01*, <https://tools.ietf.org/html/draft-handrews-json-schema-validation-01>
- IETF RFC1738, T. Berners-Lee et al., *Uniform Resource Locators (URL)*, <https://tools.ietf.org/html/rfc1738>
- IETF RFC3986, T. Berners-Lee et al., *Uniform Resource Identifier (URI): Generic Syntax*, <https://tools.ietf.org/html/rfc3986>
- IETF RFC4122, P. Leach et al., *A Universally Unique IDentifier (UUID) URN Namespace*, <https://tools.ietf.org/html/rfc4122>
- IETF RFC5280, D. Cooper et al., *Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile*,
- IETF RFC6455, I. Fette et al., *The WebSocket Protocol*, <https://tools.ietf.org/html/rfc6455>
- IETF RFC6585, M. Nottingham et al., *Additional HTTP Status Codes*, <https://tools.ietf.org/html/rfc6585>
- IETF RFC6749, D. Hardt, Ed., *The OAuth 2.0 Authorization Framework*, <https://tools.ietf.org/html/rfc6749>
- IETF RFC6901, P. Bryan, Ed. et al., *JavaScript Object Notation (JSON) Pointer*, <https://tools.ietf.org/html/rfc6901>
- IETF RFC7230, R. Fielding et al., *Hypertext Transfer Protocol (HTTP/1.1): Message Syntax and Routing*, <https://tools.ietf.org/html/rfc7230>
- IETF RFC7231, R. Fielding et al., *Hypertext Transfer Protocol (HTTP/1.1): Semantics and Content*, <https://tools.ietf.org/html/rfc7231>
- IETF RFC7232, R. Fielding et al., *Hypertext Transfer Protocol (HTTP/1.1): Conditional Requests*, <https://tools.ietf.org/html/rfc7232>
- IETF RFC7234, R. Fielding et al., *Hypertext Transfer Protocol (HTTP/1.1): Caching*, <https://tools.ietf.org/html/rfc7234>
- IETF RFC7540, M. Belshe et al., *Hypertext Transfer Protocol Version 2 (HTTP/2)*, <https://tools.ietf.org/html/rfc7540>
- IETF RFC7519, M. Jones et al., *JSON Web Token (JWT)*, <https://tools.ietf.org/html/rfc7519>
- IETF RFC7525, Y. Sheffer et al., *Recommendations for Secure Use of Transport Layer Security (TLS) and*

*Datagram Transport Layer Security (DTLS)*, <https://tools.ietf.org/html/rfc7525>

- IETF RFC7578, L. Masinter et al., *Returning Values from Forms: multipart/form-data*, <https://tools.ietf.org/html/rfc7578>
- IETF RFC7617, J. Reschke et al., *The 'Basic' HTTP Authentication Scheme*, <https://tools.ietf.org/html/rfc7617>
- IETF RFC8259, T. Bray, Ed., *The JavaScript Object Notation (JSON) Data Interchange Format*, <https://tools.ietf.org/html/rfc7617>
- IETF RFC8288, M. Nottingham, *Web Linking*, <https://tools.ietf.org/html/rfc8288>
- ISO 639-1:2002, *Codes for the representation of names of languages - Part 1: Alpha-2 code*, <https://www.iso.org/standard/22109.html>
- 24 February 2014, *OData Version 4.0 Part 1: Protocol*, <https://docs.oasis-open.org/odata/odata/v4.0/os/part1-protocol/odata-v4.0-os-part1-protocol.html>
- 24 February 2014, *OData Version 4.0 Part 3: Common Schema Definition Language (CSDL)*, <https://docs.oasis-open.org/odata/odata/v4.0/os/part3-csdl/odata-v4.0-os-part3-csdl.html>
- 10 March 2016, *OData Version 4.0 Plus Errata 03 OASIS Standard incorporating Draft 01 of Errata 03*, <https://docs.oasis-open.org/odata/odata/v4.0/errata03/csd01/complete/vocabularies/Org.OData.Measures.V1.xml>
- 20 November 2014, SNIA *TLS Specification for Storage Systems*, [https://www.snia.org/tech\\_activities/standards/curr\\_standards/tls](https://www.snia.org/tech_activities/standards/curr_standards/tls)
- *The OpenAPI Specification*, <https://swagger.io/specification/>
- *The Unified Code for Units of Measure*, <https://ucum.org/ucum.html>
- 9 September 2021, *Fetch Living Standard*, <https://fetch.spec.whatwg.org/>
- 17 September 2021, *9.2 Server-sent events in the HTML Living Standard*, <https://html.spec.whatwg.org/multipage/server-sent-events.html>

## 3 Terms, definitions, symbols, and abbreviated terms

Some terms and phrases in this document have specific meanings beyond their typical English meanings. This clause defines those terms and phrases.

The terms "shall" ("required"), "shall not", "should" ("recommended"), "should not" ("not recommended"), "may", "need not" ("not required"), "can" and "cannot" in this document are to be interpreted as described in ISO/IEC Directives, Part 2, Clause 7. The terms in parenthesis are alternatives for the preceding term, for use in exceptional cases when the preceding term cannot be used for linguistic reasons. Note that ISO/IEC Directives, Part 2, Clause 7 specifies additional alternatives. Occurrences of such additional alternatives shall be interpreted in their normal English meaning.

The terms "clause", "subclause", "paragraph", and "annex" in this document are to be interpreted as described in ISO/IEC Directives, Part 2, Clause 6.

The terms "normative" and "informative" in this document are to be interpreted as described in ISO/IEC Directives, Part 2, Clause 3. In this document, clauses, subclauses, or annexes labeled "(informative)" do not contain normative content. Notes and examples are always informative elements.

The term "deprecated" in this document is to be interpreted as material that is not recommended for use in new development efforts. Existing and new implementations may use this material, but they should move to the favored approach. Deprecated material may be implemented in order to achieve backwards compatibility. Deprecated material should contain references to the last published version that included the deprecated material as normative material and to a description of the favored approach. Deprecated material may be removed from the next major version of the specification.

This document defines these additional terms:

### 3.1 Hardware terms

#### 3.1.1 baseboard management controller (BMC)

embedded device or service

Note 1 to entry: Typically, an independent microprocessor or system-on-chip with associated firmware in a computer system that completes out-of-band systems monitoring and management-related tasks.

#### 3.1.2 IPMI

Intelligent Platform Management Interface

### **3.1.3 KVM-IP**

keyboard, video, mouse redirection over IP

### **3.1.4 NIC**

network interface controller

### **3.1.5 PCI**

Peripheral Component Interconnect

### **3.1.6 PCIe**

Peripheral Component Interconnect Express

## **3.2 Web development terms**

---

### **3.2.1 CORS**

cross-origin resource sharing

### **3.2.2 CRUD**

basic **C**reate, **R**ead, **U**pdate, and **D**elete operations that any interface can support

### **3.2.3 CSRF**

cross-site request forgery

### **3.2.4 HTTP**

Hypertext Transfer Protocol

### **3.2.5 HTTPS**

Hypertext Transfer Protocol Secure

Note 1 to entry: TLS secures HTTP.

### 3.2.6 hypermedia API

API that enables you to navigate through URIs that a service returns

### 3.2.7 IP

Internet Protocol

### 3.2.8 JSON

JavaScript Object Notation

### 3.2.9 message

complete HTTP-formatted or HTTPS-formatted request or response

Note 1 to entry: In the REST-based Redfish protocol, every request results in a response.

### 3.2.10 OData

Open Data Protocol (OData), as defined in [OData Version 4.0 Part 1: Protocol](#)

### 3.2.11 OData service document

resource that provides information about the [service root](#) for generic OData clients

### 3.2.12 operation

HTTP `POST` , `GET` , `PUT` , `PATCH` , `HEAD` , and `DELETE` request methods that map to generic [CRUD](#) operations

### 3.2.13 parent resource

parent to another resource if the initial segment of the resource URI is the same as the URI of the other resource, but is at least one level higher

Note 1 to entry: For example, `/redfish/v1/Chassis/A88` is a parent resource of `/redfish/v1/Chassis/A88/Assembly` .

### 3.2.14 property

name-value pair in a Redfish-defined request or response

Note 1 to entry: A property can be any valid JSON data type.

### **3.2.15 request**

message from a client to a service

### **3.2.16 response**

message from a service to a client in response to a request message

### **3.2.17 subscription**

registration of a destination to receive events

### **3.2.18 TCP**

Transmission Control Protocol

### **3.2.19 TLS**

Transport Layer Security

### **3.2.20 XSS**

cross-site scripting

## **3.3 Redfish terms**

---

### **3.3.1 action**

a custom operation specific to one or more resources other than one of the standard CRUD operations

Note 1 to entry: For more information, see the [Actions](#) clause.

### **3.3.2 collection**

see [resource collection](#)

### **3.3.3 event**

data structure that corresponds to one or more alerts that have occurred

---

Note 1 to entry: For more information, see the [Eventing](#) clause.

### 3.3.4 excerpt

subset of data that is copied from one resource and presented in another resource

Note 1 to entry: An excerpt provides data in convenient locations without duplication of entire resources.

### 3.3.5 member

single [resource](#) instance in a [resource collection](#)

### 3.3.6 Redfish client

communicates with a [Redfish service](#) and accesses one or more of the service's resources or functions

### 3.3.7 Redfish protocol

discovers, connects to, and inter-communicates with a [Redfish service](#)

### 3.3.8 Redfish schema

a set of human and machine-readable documents that define Redfish resources using one or more of the supported [schema definition languages](#)

### 3.3.9 Redfish service

implementation of the protocols, resources, and functions that deliver the interface that this specification defines and its associated behaviors for one or more managed systems

Note 1 to entry: Also known as the *service*.

### 3.3.10 resource

URI-addressable Redfish data structure

### 3.3.11 resource collection

set of similar resources where the number of instances can shrink or grow

### 3.3.12 resource tree

tree structure of resources accessible through a well-known starting URI

Note 1 to entry: A client can discover the available resources on a Redfish service by following the resource hyperlinks from the base of the tree.

### 3.3.13 resource type

set of definitions for [properties](#) and [actions](#) contained within a [resource](#) and documented in the [Redfish schema](#) files

### 3.3.14 service root

starting-point resource for locating and accessing the other resources and associated metadata that make up an instance of a Redfish service

### 3.3.15 subordinate resource

is subordinate to another resource if the initial segment of the resource URI is the same as the URI of the other resource, but is at least one level deeper

Note 1 to entry: For example, `/redfish/v1/Chassis/A88/Assembly` is a subordinate resource of the `Chassis` resource named `A88`.

### 3.3.16 task

representation of a [long-running operation](#)

### 3.3.17 task monitor

opaque service-generated URI that the client that initiates the request can use to monitor an [asynchronous operation](#)

## 4 Typographical conventions

---

The following typographical convention indicates deprecated material:

---

**DEPRECATED**

Deprecated material appears here.

**END DEPRECATED**

---

In places where this typographical convention cannot be used, such as tables or figures, the "DEPRECATED" label is used alone.

## 5 Overview

Redfish is a management standard that uses a data model representation with a RESTful interface.

Being RESTful, Redfish is easier to use and implement.

Being model-oriented, it can express the relationships between components and the semantics of the Redfish services and components within them. The model is also easy to extend.

By requiring JSON representation, Redfish enables easy integration with programming environments. It is also easy to interpret by humans.

An interoperable Redfish schema defines this model, which is freely available and published in OpenAPI YAML, OData CSDL, and JSON Schema formats.

### 5.1 Goals

As an architecture, data model, and set of protocols that enable a client to access Redfish services, Redfish has these goals.

[Table 1](#) describes these goals:

**Table 1 — Redfish goals**

Goal	Purpose
Scalable	Can scale on stand-alone machines or racks of equipment.
Flexible	Can implement through existing hardware or entirely as a software service.
Extensible	Can easily add new and vendor-specific capabilities to the data model.
Backward-compatible	Can add capabilities while preserving investments in earlier implementations.
Interoperable	Provides consistent functionality across multiple vendor implementations.
Standards-based	Built on ubiquitous and secure protocols. Leverages other standards where applicable.
Simple	Easy-to-use without the need for highly specialized programming skills or systems knowledge.
Lightweight	Designed to reduce complexity and implementation costs. Minimizes the required footprint for implementations.

## 5.2 Design tenets

To deliver these goals, Redfish:

- Provides a RESTful interface by using a JSON payload and a data model.
- Separates the protocol from the data model, which enables the independent revision and use of each.
- Specifies versioning rules for protocols and schema.
- Leverages strength of ubiquitous standards where it meets architectural requirements, such as JSON, HTTP, OData, OpenAPI, and the RFCs that this document references.
- Organizes the data model so that it provides clearly demarcated and value-add features in the same payload as standardized items.
- Makes data in payloads as obvious in context as possible.
- Maintains implementation flexibility. Does not tie the interface to any particular underlying implementation or architecture.
- Focuses on widely used capabilities. To avoid complexity, does not add functions that only a small percentage of users value.

## 5.3 Limitations

Redfish minimizes the need for clients to complete upgrades by using strict versioning and forward-compatibility rules, and separation of the protocols from the data model. However, Redfish does not guarantee that clients never need to update their software. For example, clients might need to upgrade to manage new system or component types, or update the data model.

Interoperable does not mean identical. Many elements of Redfish are optional. Clients should be prepared to discover the optional elements by using the built-in discovery methods.

The *resource tree* reflects the topology of the system and its devices. Consequently, different hardware or device types result in different resource trees, even for identical systems from the same manufacturer. References between resources may result in a graph instead of a tree. Clients that traverse the resource tree should provide logic to avoid infinite loops.

Additionally, not all Redfish resources use simple REST read-and-write semantics. Different use cases may follow other types of client logic. For example, clients cannot simply read user credentials or certificates from one service and write them to another service.

Finally, the hyperlink values between resources and other elements can vary across implementations. Clients should not assume that they can reuse hyperlinks across different Redfish service instances.

## 5.4 Additional design background and rationale

---

### 5.4.1 REST-based interface

Redfish exposes many service applications as RESTful interfaces. This document defines a RESTful interface.

Redfish defines a RESTful interface because it:

- Enables a lightweight implementation, using fewer layers than previous standards.
- Is a prevalent access method in the industry.
- Is easy to learn, document, and implement in modern programming languages.
- Has a number of development environments and a healthy tooling ecosystem.
- Fits with the design goal of simplicity.
- Equally applies to software application space as it does to embedded environments, which enables convergence and sharing of code within the management ecosystem.
- Adapts well to any data modeling language.
- Has industry-provided security and discovery mechanisms.

### 5.4.2 Data-oriented

The Redfish data model is developed by focusing on the contents of the payload. By concentrating on the contents of the payload first, Redfish payloads are easily mapped to schema definition languages and encoding types. The data model is defined in various schema languages, including OpenAPI YAML, OData CSDL, and JSON Schema.

### 5.4.3 Separation of protocol from data model

Redfish separates the protocol operations from the data model and versions the protocol independently from the data model. This enables clients to extend and change the data model as needed without requiring the protocol version to change.

### 5.4.4 Hypermedia API service root

Redfish has a single service root URI and clients can discover all resources through referenced URIs. The [hypermedia API](#) enables the discovery of resources through hyperlinks.

### 5.4.5 OpenAPI v3.0 support

The OpenAPI v3.0 provides a rich ecosystem of tools for using RESTful interfaces that meet the design requirements of that specification. Starting with *Redfish Specification v1.6.0*, the Redfish schemas support the OpenAPI YAML file format and URI patterns that conform to the OpenAPI Specification were defined. Conforming Redfish services that

support the Redfish protocol version v1.6.0 or later implement those URI patterns to enable use of the OpenAPI ecosystem.

For details, see [OpenAPI Specification v3.0](#).

#### 5.4.6 OData conventions

With the popularity of RESTful APIs, there are nearly as many RESTful interfaces as there are applications. While following REST patterns helps promote good practices, due to design differences between the many RESTful APIs there few common conventions between them.

To provide for interoperability between APIs, [OData](#) defines a set of common RESTful conventions and annotations. Redfish follows OData conventions for describing schema, [URL conventions](#), and definitions for typical properties in a JSON payload.

### 5.5 Service elements

---

#### 5.5.1 Synchronous and asynchronous operation support

Some operations can take more time than a client typically wants to wait. For this reason, some operations can be asynchronous at the discretion of the service. The request portion of an asynchronous operation is no different from the request portion of a synchronous operation.

To determine whether an operation was completed synchronously or asynchronously, clients can review the [HTTP status codes](#). For more information, see the [Asynchronous operations](#) clause.

#### 5.5.2 Eventing mechanism

Redfish enables clients to receive messages outside the normal request and response paradigm. The service uses these messages, or *events*, to asynchronously notify the client of a state change or error condition, usually of a time critical nature.

This specification defines two styles of eventing:

- Push-style eventing.

When the service detects the need to send an event, it calls `HTTP POST` to push the event message to the client. Clients can enable reception of events by creating a [subscription](#) entry in the event service, or an administrator can create subscriptions as part of the Redfish service configuration.

- [Server-sent events \(SSE\)](#)-style eventing.

The client opens an SSE connection to the service through a `GET` on the `ServerSentEventUri`-specified URI in the event service.

For information, see the [Eventing](#) clause.

### 5.5.3 Actions

Actions are Redfish operations that do not easily map to RESTful interface semantics. These types of operations may not directly affect properties in the Redfish resources. The Redfish schema defines certain standard actions for common Redfish resources. For these standard actions, the Redfish schema contains the normative language on the behavior of the action.

### 5.5.4 Service discovery

While the service itself is at a well-known URI, clients need to discover the network address of the service. Like UPnP, Redfish uses SSDP for discovery. A wide variety of devices, such as printers and client operating systems, support SSDP. It is simple, lightweight, IPv6 capable, and suitable for implementation in embedded environments.

For more information, see the [Discovery](#) clause.

### 5.5.5 Remote access support

Remote management functionality typically includes access mechanisms for redirecting operator interfaces such as serial console, keyboard video and mouse (KVM-IP), command shell, or command-line interface, and virtual media. While these mechanisms are critical functionality, they cannot be reasonably implemented as a RESTful interface.

Therefore, this standard does not define the protocols or access mechanisms for those services but encourages implementations that leverage existing standards. However, the Redfish schema includes resources and properties that enable client discovery of these capabilities and access mechanisms to enable interoperability.

## 5.6 Security

---

The challenge of remote interface security is to protect both the interface and exchanged data. To accomplish this, Redfish provides authentication and encryption. As part of this security, Redfish defines and requires minimum levels of encryption.

For more information, see the [Security details](#) clause.

## 6 Protocol details

In this document, the Redfish protocol refers to the RESTful mapping to HTTP, TCP/IP, and other protocol, transport, and messaging layer aspects. HTTP is the application protocol that transports the messages and TCP/IP is the transport protocol. The RESTful interface is a mapping to the message protocol.

The Redfish protocol is designed around a web service-based interface model, which provides network and interaction efficiency for both user interface (UI) and automation usage. Specifically, the protocol can leverage existing tool chains.

[Table 2](#) describes the items that the Redfish protocol uses:

**Table 2 — Redfish protocol**

Item	Description
<a href="#">HTTP methods</a>	Maps to common CRUD operations.
<a href="#">Actions</a>	Expands operations beyond CRUD-type operations.
<a href="#">Media types</a>	Negotiates the type of data sent in the message body.
<a href="#">HTTP status codes</a>	Indicates the success or failure of the server's request.
<a href="#">Error responses</a>	Returns more information than HTTP status codes.
<a href="#">TLS</a>	Secures messages. See <a href="#">Security details</a> .
<a href="#">Asynchronous semantics</a>	Manages long-running operations.

A Redfish interface shall be exposed through a web service endpoint implemented by using HTTP version 1.1. See [RFC7230](#), [RFC7231](#), and [RFC7232](#).

A Redfish interface may additionally be exposed through a web service endpoint implemented by using HTTP version 2.0. See [RFC7540](#).

The subsequent clauses describe how the Redfish interface uses and adds constraints to HTTP to ensure interoperability of Redfish implementations.

### 6.1 Universal Resource Identifiers

A Universal Resource Identifier (URI) identifies a resource, including the service root and all Redfish resources.

- A URI shall identify each unique instance of a resource.

- URIs shall not include any [RFC1738](#)-defined unsafe characters.
  - For example, the {, }, , |, ^, ~, [ , ], ` , and \ characters are unsafe because gateways and other transport agents can sometimes modify these characters.
  - Do not use the # character for anything other than the start of a fragment.
- URIs shall not include any percent-encoding of characters. This restriction does not apply to the [query parameters](#) portion of the URI.

A `GET` operation on a URI returns a representation of the resource with properties and hyperlinks to associated resources. The service root URI is well known and is based on the protocol version.

To discover the URIs to additional resources, extract the associated resource hyperlinks from earlier responses. The [hypermedia API](#) enables the discovery of resources through hyperlinks.

Redfish considers the [RFC3986](#)-defined scheme, authority, service root, and version, and unique resource path component parts of the URI.

For example, this URI:

```
https://mgmt.vendor.com/redfish/v1/Systems/1
```

Contains these component parts:

- https: is the scheme.
- mgmt.vendor.com is the authority to which to delegate the URI.
- redfish/v1 is the service root and version.
- Systems/1 is the unique resource path.

In a URI:

- The scheme and authority component parts are not part of the unique resource path because redirection capabilities and local operations may cause the connection portion to vary.
- The service root and resource path component parts *uniquely identify* the resource in a Redfish service.

In an implementation:

- The resource path component part shall be unique.
- A [relative reference](#) in the body and HTTP headers payload can identify a resource in that same implementation.
- An absolute URI in the body and HTTP headers payload can identify a resource in a different implementation.

For the absolute URI definition, see [RFC3986](#).

For example, a `POST` operation may return the `/redfish/v1/Systems/2` URI in the `Location` header of the response, which points to the `POST`-created resource.

Assuming that the client connects through the `mgmt.vendor.com` appliance, the client accesses the resource through the `https://mgmt.vendor.com/redfish/v1/Systems/2` absolute URI.

URIs that conform to [RFC3986](#) may also contain the `?query`, `#frag`, and `#frag`, components. For information about queries, see [Query parameters](#). When a URI includes a fragment (`frag`) to submit an operation, the server ignores the fragment.

If a property in a response references another property within a resource, use the [RFC6901](#)-defined URI fragment identifier representation format. If the property is a [reference property](#) in the schema, the fragment shall reference a valid [resource identifier](#). For example, the following fragment identifies a property at index 0 of the `Fans` array in the `/redfish/v1/Chassis/MultiBladeEncl/Thermal` resource:

```
{
  "@odata.id": "/redfish/v1/Chassis/MultiBladeEncl/Thermal#/Fans/0"
}
```

For requirements on constructing Redfish URIs, see the [resource URI patterns annotation](#) clause.

## 6.2 HTTP methods

[Table 3](#) describes the mapping of HTTP methods to the Redfish-supported operations. If the **Required** column contains **Yes**, a Redfish interface shall support the HTTP method. If the **Required** column contains **No**, a Redfish interface may support the HTTP method.

**Table 3 — Mapping of HTTP methods to Redfish-supported operations**

HTTP method	Interface semantic	Required
POST	Create resource Resource action Eventing	Yes
GET	Retrieve resource	Yes
PUT	Replace resource	No
PATCH	Update resource	Yes
DELETE	Delete resource	Yes
HEAD	Retrieve resource header	No

HTTP method	Interface semantic	Required
OPTIONS	Retrieve header Cross-origin resource sharing (CORS) pre-flight	No

For HTTP methods that the Redfish service does not support or that [Table 3](#) omits, the Redfish service shall return the HTTP [405 Method Not Allowed](#) status code or the HTTP [501 Not Implemented](#) status code.

## 6.3 HTTP redirect

---

HTTP redirect enables a service to redirect a request to another URL. Among other things, HTTP redirect enables Redfish resources to alias areas of the data model.

All Redfish clients shall correctly handle HTTP redirect.

The service for the redirected resource shall enforce the [authentication](#) and [authorization](#) requirements for the redirected resource.

## 6.4 Media types

---

Some resources may be available in more than one type of representation. The media type indicates the representation type.

In HTTP messages, the media type is specified in the `Content-Type` header. To tell a service to return the response through certain media types, the client sets the HTTP `Accept` header to a list of the media types.

- All resources shall be available through the JSON `application/json` media type.
- Redfish services shall make every resource available in a JSON-based representation as a JSON object, as specified in [RFC8259](#). Receivers shall not reject a JSON-encoded message, and shall offer at least one JSON-based response representation. An implementation may offer additional non-JSON media type representations.

To request compression in the response body, clients specify an `Accept-Encoding` request header.

## 6.5 ETags

---

To reduce unnecessary RESTful accesses to resources, the Redfish service should support the association of a separate entity tag (ETag) with each resource.

- Implementations should support the return of [ETag properties](#) for each resource.
- Implementations should support the return of ETag headers for each single-resource response.

- Implementations shall support the return of ETag headers for `GET` requests of `ManagerAccount` resources.

Because the service knows whether the new version of the object is substantially different, the service generates and provides the ETag as part of the resource payload.

The ETag mechanism supports both **strong** and **weak** validation. If a resource supports an ETag, it shall use the [RFC7232](#)-defined ETag.

This specification does not mandate a particular algorithm for ETag creation, but ETags should be highly collision-free.

An ETag can be:

- A hash
- A generation ID
- A time stamp
- Some other value that changes when the underlying object changes

If a client performs a `PUT` operation or `PATCH` operation to update a resource, it should include an ETag from a previous `GET` in the HTTP `If-Match` or `If-None-Match` header. Both **strong** and **weak** ETags are allowed in these headers. If a service supports the return of the ETag header on a resource, it may respond with the HTTP [428 Precondition Required](#) status code if the `If-Match` or `If-None-Match` header is missing from the `PUT` or `PATCH` request for the same resource, as specified in [RFC6585](#). Services should perform **weak** ETag comparison when verifying the ETag provided by the client in `PUT` or `PATCH` operations. Clients should treat ETags received from services as opaque values and not modify them when providing them in `PUT` or `PATCH` operations.

In addition to the return of the ETag property on each resource, a Redfish service should return the ETag header on:

- A client `PUT`, `POST`, or `PATCH` operation
- A `GET` operation for an individual resource

The format of the ETag header is:

```
ETag: <string>
```

For responses to `$expand` requests:

- The `@odata.etag` property of each resource in the response shall contain the ETag of the resource as if it were not expanded.
- The ETag header should contain the ETag of the entire response body.

Services should omit properties named `DateTime` from ETag calculations. Properties named `DateTime` represent

clock settings that increment monotonically over time. Services should implement methods to reduce the frequency of ETag updates for other types of fast changing properties, such the `Reading` property in the `Sensor` resource, for ETag calculations.

Special handling of ETag calculations is important for clients that use the `If-Match` header to prevent collisions with other clients. For example, if a client is attempting to modify a `Manager` resource and the `DateTime` property updates the ETag every second, a client performing a `PATCH` operation on the `Manager` resource with the `If-Match` header will likely contain an old ETag if `DateTime` is not omitted from ETag calculations, causing the service to respond with the HTTP `412 Precondition Failed` status code. However, omitting a property from ETag calculations might be problematic for clients using the `If-Match-None` header to reduce the amount of data transmitted over the network. For example, if the `Reading` property is omitted from the ETag calculation in a `Sensor` resource, a client using the `If-Match-None` header for `GET` operations could continually get the HTTP `304 Not Modified` status code regardless of how much the sensor reading changed. In this case, other methods to reduce the frequency of ETag changes are preferred.

## 6.6 Protocol version

---

The protocol version is separate from the resources' version or the Redfish schema version that the resources support.

Each Redfish protocol version is strongly typed by using the URI of the Redfish service in combination with the resource obtained at that URI, called the `ServiceRoot` resource.

The root URI for this version of the Redfish protocol shall be `/redfish/v1/`.

The URI defines the major version of the protocol.

The `RedfishVersion` property of the `ServiceRoot` resource defines the protocol version, which includes the major version, minor version, and errata version of the protocol, as defined in the Redfish schema for that resource.

The protocol version is a string in the format:

`<MajorVersion>.<MinorVersion>.<ErrataVersion>`

where

- `<MajorVersion>` is an integer that represents the major version. Indicates a backward-incompatible change.
- `<MinorVersion>` is an integer that represents the minor version. Indicates a minor update. Redfish introduces functionality but does not remove any functionality. The minor version preserves compatibility with earlier minor versions.
- `<ErrataVersion>` is an integer that represents the errata version. Indicates a fix to the earlier version.

Any resource that a client discovers through hyperlinks that the service root or any service root-referenced service or resource returns shall conform to the same protocol version that the service root supports.

A `GET` operation on the `/redfish` resource shall return this response body:

```
{
  "v1": "/redfish/v1/"
}
```

## 6.7 Redfish-defined URIs and relative reference rules

**Table 4** describes the Redfish-defined URIs. Redfish services shall support URIs in the table if the **Required** column contains **Yes**. Redfish services should support URIs in the table if the **Required** column contains **No**.

**Table 4 — Redfish-defined URIs**

URI	Returns	Required
<code>/redfish</code>	Version. A major update that does not preserve compatibility with earlier minor versions.	Yes
<code>/redfish/v1/</code>	Redfish service root.	Yes
<code>/redfish/v1/odata</code>	Redfish OData service document.	Yes
<code>/redfish/v1/\$metadata</code>	Redfish metadata document.	Yes
<code>/redfish/v1/openapi.yaml</code>	Redfish OpenAPI YAML document.	No
<code>/redfish/v1/Schemas/&lt;SchemaFile&gt;</code>	Local copy of a Redfish schema file, where <code>&lt;SchemaFile&gt;</code> is the file name of the local schema file.	No
<code>/redfish/v1/Registries/&lt;RegistryFile&gt;</code>	Local copy of a Redfish registry file, where <code>&lt;RegistryFile&gt;</code> is the file name of the local registry file.	No
<code>/redfish/v1/TaskService/TaskMonitors/&lt;TaskMonitorId&gt;</code>	Redfish task monitor, where <code>&lt;TaskMonitorId&gt;</code> is the identifier of the task monitor.	No

In addition, **Table 5** describes the URIs that services shall process without a trailing slash in one of these ways:

- Redirect it to the associated Redfish-defined URI.
- Treat it as the equivalent URI to the associated Redfish-defined URI.

**Table 5 — Redfish-defined URIs without trailing slashes**

URI	Associated Redfish-defined URI
/redfish/v1	/redfish/v1/
/redfish/	/redfish

All other Redfish service-supported URIs shall match the [resource URI patterns definitions](#), except the supplemental resources that the `@Redfish.Settings`, `@Redfish.ActionInfo`, and `@Redfish.CollectionCapabilities payload annotations` reference. The client shall treat the URIs for these supplemental resources as opaque.

All Redfish-defined URIs and URIs starting with `/redfish` are reserved for future standardization by DMTF and DMTF alliance partners, except OEM extension URIs, which shall conform to the requirements of the [OEM URIs](#) clause.

All relative references that the service uses shall start with either:

- A double forward slash (`//`) and include the authority (network-path), such as `//mgmt.vendor.com/redfish/v1/Systems`.
- A single forward slash (`/`) and include the absolute-path, such as `/redfish/v1/Systems`.

For details, see [RFC3986](#).

## 7 Service requests

---

This clause describes the requests that clients can send to Redfish services.

### 7.1 Request headers

---

[Table 6](#) lists the HTTP-defined request headers and their requirements for Redfish services and clients.

For Redfish services:

- Redfish services shall process the HTTP-defined headers in [Table 6](#) if the **Service requirement** column contains **Yes** or **Conditional**.
- Redfish services should process the HTTP-defined headers in [Table 6](#) and [Table 7](#) if the **Service requirement** column contains **No**.

For Redfish clients (sending the HTTP requests):

- Redfish clients shall include the HTTP-defined headers in [Table 6](#) if the **Client requirement** column contains **Yes** or **Conditional**.
- Redfish clients should transmit the HTTP-defined headers in [Table 6](#) and [Table 7](#) if the **Client requirement** column contains **No**.

Table 6 — Request headers

Header	Service requirement	Client requirement	Supported values	Description
Accept	Yes	No	RFC7231	<p>Communicates to the server the media type or types that this client is prepared to accept.</p> <p>Services shall support resource requests with <code>Accept</code> header values of <code>application/json</code> OR <code>application/json; charset=utf-8</code>.</p> <p>Services shall support XML metadata requests with <code>Accept</code> header values of <code>application/xml</code> OR <code>application/xml; charset=utf-8</code>.</p> <p>Services shall support OpenAPI YAML schema requests with <code>Accept</code> header values of <code>application/yaml</code> OR <code>application/yaml; charset=utf-8</code> OR <code>application/vnd.oai.openapi</code> OR <code>application/vnd.oai.openapi; charset=utf-8</code>.</p> <p>Services shall support SSE requests with <code>Accept</code> header values of <code>text/event-stream</code> OR <code>text/event-stream; charset=utf-8</code>.</p> <p>Services shall support any request with <code>Accept</code> header values of <code>application/*</code>, <code>application/*; charset=utf-8</code>, <code>*/*</code>, or <code>*/*; charset=utf-8</code>.</p>
Accept-Encoding	No	No	RFC7231	Indicates whether the client can handle gzip-encoded responses. If a service cannot return an acceptable response to a request with this header, it shall respond with the HTTP <a href="#">406 Not Acceptable</a> status code. If the request omits this header, the service should not return gzip-encoded responses.
Accept-Language	No	No	RFC7231	The languages that the client accepts in the response. If the request omits this header, uses the service's default language for the response.
Authorization	Conditional	Conditional	RFC7617, RFC6749	<p>Required for <a href="#">HTTP Basic authentication</a> and <a href="#">OAuth 2.0</a>.</p> <p>A client can access unsecured resources without this header on systems that support Basic authentication.</p>

Header	Service requirement	Client requirement	Supported values	Description
Content-Length	No	No	RFC7231	<p>The size of the message body.</p> <p>To indicate the size of the body, a client can use the <code>Transfer-Encoding: chunked</code> header.</p> <p>If a service needs to use <code>Content-Length</code> and does not support <code>Transfer-Encoding</code>, it responds with the HTTP <a href="#">406 Not Acceptable</a> status code.</p>
Content-Type	Conditional	Conditional	RFC7231	<p>The request format. Required for operations with a request body.</p> <p>Services shall accept the <code>Content-Type</code> header set to either <code>application/json</code> or <code>application/json; charset=utf-8</code>.</p> <p>It is recommended that clients use these values in requests because other values can cause an error.</p>
Host	Yes	No	RFC7230	Enables support of multiple origin hosts at a single IP address.
If-Match	Conditional	No	RFC7232	<p>To ensure that clients update the resource from a known state, <code>PUT</code> and <code>PATCH</code> requests for resources for which a service returns ETags shall support <code>If-Match</code>.</p> <p>While not required for clients, it is highly recommended for <code>PUT</code> and <code>PATCH</code> operations.</p>
If-None-Match	No	No	RFC7232	<p>A service only returns the resource if the current ETag of that resource does not match the ETag sent in this header.</p> <p>If the ETag in this header matches the resource's current ETag, the <code>GET</code> operation returns the HTTP <a href="#">304 Not Modified</a> status code.</p>
Last-Event-ID	No	No	HTML5 SSE	<p>The event source's last <code>id</code> field from the SSE stream.</p> <p>Requests history event data.</p> <p>See <a href="#">Server-sent events</a>.</p>
Max-Forwards	No	No	RFC7231	<p>Limits gateway and proxy hops.</p> <p>Prevents messages from remaining in the network indefinitely.</p>

Header	Service requirement	Client requirement	Supported values	Description
OData-MaxVersion	No	No	4.0	The maximum OData version that an OData-aware client understands.
OData-Version	Yes	No	4.0	<p>The OData version.</p> <p>Services shall reject requests that specify an unsupported OData version.</p> <p>If a service encounters an unsupported OData version, it should reject the request with the HTTP <a href="#">412 Precondition Failed</a> status code.</p>
Origin	Yes	No	Fetch Living Standard, 3.1. <a href="#">origin header</a>	Enables web applications to consume a Redfish service while preventing CSRF attacks.
User-Agent	Yes	No	RFC7231	<p>Traces product tokens and their versions.</p> <p>The header can list multiple product tokens.</p>
Via	No	No	RFC7230	<p>Defines the network hierarchy and recognizes message loops.</p> <p>Each pass inserts its own <code>via</code> header.</p>

Redfish services shall understand and be able to process the headers in [Table 7](#) as defined by this specification if the **Service requirement** column contains **Yes**.

**Table 7 — Request headers part 2**

Header	Service requirement	Client requirement	Supported values	Description
X-Auth-Token	Yes	Conditional	Opaque encoded octet strings	<p>Authenticates user sessions.</p> <p>The token value shall be indistinguishable from random.</p> <p>While services shall support this header, a client can access unsecured resources without establishing a session.</p>

## 7.2 GET (read requests)

### 7.2.1 GET (read requests) overview

The `GET` operation retrieves resources from a Redfish service. Clients make a `GET` request to the individual resource URI. Clients may obtain the resource URI from published sources, such as the OpenAPI document, or from a [resource identifier property](#) in a previously retrieved resource response, such as the [links property](#).

The service shall return the resource representation using one of the media types listed in the `Accept` header, subject to the requirements of the [media types](#). If the `Accept` header is absent, the service shall return the resource's representation as `application/json`. Services may but are not required to support the convention of retrieving individual properties within a resource by appending a segment containing the property name to the URI of the resource.

- The HTTP `GET` operation shall retrieve a resource without causing any side effects.
- The service shall ignore the content of the body on a `GET`.
- The `GET` operation shall be idempotent in the absence of outside changes to the resource.

If supported by the service, clients can perform a conditional `GET` operation by specifying an [If-None-Match request header](#) that contains the ETag of the resource.

### 7.2.2 Resource collection requests

Clients retrieve a resource collection by making a `GET` request to the resource collection URI. The response includes the resource collection's properties and an array of its [members](#).

No requirements are placed on implementations to return a consistent set of members when a series of requests that use paging query parameters are made over time to obtain the entire set of members. These calls can result in missed or duplicate elements if multiple `GET` requests use paging to retrieve the `Members` array instances.

- Clients shall not make assumptions about the URIs for the members of a resource collection.
- Retrieved resource collections shall always include the [count \( Members@odata.count \) property](#) to specify the total number of entries in its `Members` array.
- Regardless of the [next link \( Members@odata.nextLink \) property](#) or paging, the [count \( Members@odata.count \) property](#) shall return the total number of resources that the `Members` array references.

A subset of the members can be retrieved using client paging [query parameters](#).

A service might not be able to return all of the contents of a resource collection request in a single response body. In this case, the response can be paged by the service. If a service pages a response to a resource collection request, the following rules shall apply:

- Responses can contain a subset of the full resource collection's members.
- Individual members shall not be split across response bodies.
- A [next link \(`Members@odata.nextLink`\) property](#) annotation shall be supplied in the response body with the URI to the next set of members in the collection.
- The [next link \(`Members@odata.nextLink`\) property](#) shall adhere to the rules in the [Next link property](#) clause.
- `GET` operations on the [next link \(`Members@odata.nextLink`\) property](#) shall return the subsequent section of the resource collection response.

### 7.2.3 Service root request

The root URL for Redfish version 1.x services shall be `/redfish/v1/`.

The service returns the `ServiceRoot` resource, as defined by this specification, as a response for the root URL.

Services shall not require authentication to retrieve the service root and `/redfish` resources.

### 7.2.4 OData service and metadata document requests

Redfish services expose two OData-defined documents at specific URIs to enable generic OData clients to navigate the Redfish service.

- Service shall expose an [OData metadata document](#) at the `/redfish/v1/$metadata` URI.
- Service shall expose an [OData service document](#) at the `/redfish/v1/odata` URI.
- Service shall not require authentication to retrieve the OData metadata document or the OData service document.

## 7.3 Query parameters

---

### 7.3.1 Query parameter overview

To paginate, retrieve subsets of resources, or expand the results in a single response, clients can include the query parameters. Some query parameters apply only to resource collections.

Services:

- Shall only support query parameters on `GET` operations.
- Should support the `$top`, `$skip`, `only`, and `excerpt` query parameters.
- May support the `$expand`, `$filter`, and `$select` query parameters.
- Shall include the `ProtocolFeaturesSupported` object in the service root, if the service supports query parameters.
  - This object indicates which parameters and options have been implemented.
- Shall ignore unknown or unsupported query parameters that do not begin with `$`.

- Shall use the `&` operator to separate multiple query parameters in a single request.
- Should ignore the `=` character if provided as the last character for the `only`, `excerpt`, `includeoriginofcondition` query parameters.
- Should treat the names of query parameters as case-insensitive strings.

Services shall return:

- The HTTP `501 Not Implemented` status code for any unsupported query parameters that start with `$`.
- An [extended error](#) that indicates the unsupported query parameters for this resource.
- The HTTP `400 Bad Request` status code for any query parameters that contain values that are invalid, or values applied to query parameters without defined values, such as `excerpt` or `only`.

Services should return:

- The HTTP `400 Bad Request` status code and an [error response](#) with the `QueryNotSupportedOnResource` message from the Base Message Registry for any implemented query parameters that are not supported on a resource in the request.
- The HTTP `400 Bad Request` status code and an [error response](#) with the `QueryNotSupportedOnResource` message from the Base Message Registry for any supported query parameters that apply only to resource collections but are used on singular resources. This includes query parameters such as `$filter`, `$top`, `$skip`, and `only`.
- The HTTP `400 Bad Request` status code and an [error response](#) with the `QueryNotSupportedOnOperation` message from the Base Message Registry for any supported query parameters on operations other than `GET`.

Services shall process query parameters in this order:

- `$filter`
- `$skip`
- `$top`
- Apply server-side pagination
- `$expand`
- `excerpt`
- `$select`

[Table 8](#) describes the query parameters:

**Table 8 — Query parameters**

Query parameter	Description and example
<code>excerpt</code>	<p>Returns a subset of the resource's properties that match the defined <code>Excerpt</code> schema annotation.</p> <p>If no <code>Excerpt</code> schema annotation is defined for the resource, the entire resource is returned.</p> <p><b>Example:</b></p> <pre>https://resource?excerpt</pre>
<code>\$expand=&lt;string&gt;</code>	<p>Returns a hyperlink and its contents in-line with retrieved resources, as if a <code>GET</code> call response was included in-line with that hyperlink.</p> <p>See <a href="#">The <code>\$expand</code> query parameter</a>.</p> <p><b>Example:</b></p> <pre>https://resource?\$expand=*(levels=3)</pre> <pre>https://resourcecollection?\$expand=.(levels=1)</pre>
<code>\$filter=&lt;string&gt;</code>	<p>Applies to resource collections. Returns a subset of collection members that match the <code>\$filter</code> expression.</p> <p>See <a href="#">The <code>\$filter</code> query parameter</a>.</p> <p><b>Example:</b></p> <pre>https://resourcecollection?\$filter=SystemType eq 'Physical'</pre>
<code>only</code>	<p>Applies to resource collections. If the target resource collection contains exactly one member, clients can use this query parameter to return that member's resource.</p> <p>If the collection contains either zero members or more than one member, the response returns the resource collection, as expected.</p> <p>Services should return the HTTP <code>400 Bad Request</code> and an <a href="#">error response</a> with the <code>QueryCombinationInvalid</code> message from the Base Message Registry if <code>only</code> is being combined with other query parameters.</p> <p><b>Example:</b></p> <pre>https://resourcecollection?only</pre>

Query parameter	Description and example
\$select=<string>	<p>Returns a subset of the resource's properties that match the <code>\$select</code> expression.</p> <p>See <a href="#">The \$select query parameter</a>.</p> <p><b>Example:</b></p> <pre>https://resource?\$select=SystemType,Status</pre>
\$skip=<integer>	<p>Applies to resource collections. Returns a subset of the members in a resource collection, or an empty set of members if the <code>\$skip</code> value is greater than or equal to the member count. This paging query parameter defines the number of <b>members</b> in the <b>resource collection</b> to skip.</p> <p><b>Example:</b></p> <pre>https://resourcecollection?\$skip=5</pre>
\$top=<integer>	<p>Applies to resource collections. Defines the number of members to show in the response.</p> <p>Minimum value is <code>0</code>, though a value of <code>0</code> returns an empty set of members.</p> <p><b>Example:</b></p> <pre>https://resourcecollection?\$top=30</pre>
includeoriginofcondition	<p>Returns the resource referenced by any <code>OriginOfCondition</code> properties in-line with the requested resource if not inside a <code>Links</code> property. Can be considered a specialized type of <code>\$expand=</code>, specifically for <code>OriginOfCondition</code> properties.</p> <p><b>Example:</b></p> <pre>https://resource?includeoriginofcondition</pre>

Services may support OEM-defined query parameters. OEM-defined query parameter names shall not contain characters that conflict with syntax for query parameter parsing, such as `&`. OEM-defined query parameters shall be in the form:

`OEM-<OemIdentifier>-<ParameterName>`

where

- `<OemIdentifier>` is the unique identifier of the OEM, including possible subdivisioning, that follows the same naming as defined in the [OEM-specified object naming](#) clause. Separator underscores (`_`) may be excluded for improved readability.
- `<ParameterName>` is the parameter name.

For example, if Contoso defined a `StatusOnly` parameter, the query parameter would be `OEM-Contoso-StatusOnly`.

### 7.3.2 The \$expand query parameter

The `$expand` query parameter enables a client to request a response that includes not only the requested resource, but also includes the contents of the [subordinate](#) or hyperlinked resources. The definition of this query parameter follows the [OData Protocol Specification](#).

The `$expand` query parameter has a set of possible options that determine which hyperlinks in a resource are included in the expanded response.

**Table 9** describes the Redfish-supported options for the `$expand` query parameter. The service may implement some of these options but not others. Any other supported syntax for `$expand` is outside the scope of this specification.

**Table 9 — The \$expand query parameter options**

Option	Description	Example
asterisk ( <code>*</code> )	Shall expand all hyperlinks, including those in <a href="#">payload annotations</a> , such as <code>@Redfish.Settings</code> , <code>@Redfish.ActionInfo</code> , and <code>@Redfish.CollectionCapabilities</code> .	<code>https://resource?\$expand=*</code>
<code>\$levels</code>	<p>Number of levels the service should cascade the <code>\$expand</code> operation.</p> <p>The default level shall be 1.</p> <p>For example, <code>\$levels=2</code> expands both the hyperlinks in the current resource (level 1), and the hyperlinks in the resulting expanded resources (level 2).</p>	<code>https://resourcecollection?\$expand=.(\$levels=2)</code>
period ( <code>.</code> )	Shall expand all hyperlinks <b>not</b> in any <a href="#">links property</a> instances of the resource, including those in payload annotations, such as <code>@Redfish.Settings</code> , <code>@Redfish.ActionInfo</code> , and <code>@Redfish.CollectionCapabilities</code> .	<code>https://resourcecollection?\$expand=.</code>
tilde ( <code>~</code> )	Shall expand all hyperlinks found in all <a href="#">links property</a> instances of the resource.	<code>https://resourcecollection?\$expand=~</code>

Examples of `$expand` usage include:

- Get all members of the software inventory collection.

With `$expand`, the client can request multiple `SoftwareInventory` collection member resources in one request rather than fetching them one at a time.

Example request: GET /redfish/v1/UpdateService/FirmwareInventory?\$expand=.

- Get a computer system and subordinate resources.

With `$levels`, a single `GET` request can include the subordinate resource collections, such as `Processors` and `Memory`.

Example request: GET /redfish/v1/Systems/1?\$expand=.(\$levels=3)

- Get all UUIDs of the members of the computer system collection.

To accomplish this result, include both `$select` and `$expand` on the URI. When combining `$select` and `$expand`, the full property path from the expansion source is required in the `$select` query parameter.

Example request: GET /redfish/v1/Systems?\$select=Members/UUID&\$expand=.(\$levels=1)

When services execute `$expand`, they should include all supported properties in the requested resources.

When clients use `$expand`, they should be aware that the payload may increase beyond what can be sent in a single response.

If a client does not have the privileges required to access referenced resources in an expansion request, the service shall not expand the referenced resources. Services should not fail the request if the client has the privileges to access the URI of the requested resource.

If a service cannot return the payload due to its size, it shall return the HTTP `507 Insufficient Storage` status code.

If a service cannot return the payload corresponding to an individual `member` of a resource collection, but can return other members, it should:

- Return the HTTP `200 OK` status code.
- Return the `@odata.id` property for the member that it cannot expand in the `Members` array.
- Return `extended information` with the member indicating the reason that member was not returned, such as when a provider internal to the service returns an error or times out.

Responses may automatically include expanded resources due to the `expanded resource annotation` on `reference properties`, such as the `Members` property in the `LogEntryCollection` resource. `$expand` operates in parallel to the usage of this schema annotation. For example, since `Members` is automatically expanded, the following three requests produce identical responses:

- /redfish/v1/Systems/1/LogServices/FaultLog/Entries
- /redfish/v1/Systems/1/LogServices/FaultLog/Entries?\$expand=.
- /redfish/v1/Systems/1/LogServices/FaultLog/Entries?\$expand=.(\$levels=1)

The following example expands the `RoleCollection` resource with the level set to 1:

```
{  
    "@odata.id": "/redfish/v1/AccountService/Roles",  
    "@odata.type": "#RoleCollection.RoleCollection",  
    "Name": "Roles Collection",  
    "Members@odata.count": 3,  
    "Members": [  
        {  
            "@odata.id": "/redfish/v1/AccountService/Roles/Administrator",  
            "@odata.type": "#Role.v1_1_0.Role",  
            "Id": "Administrator",  
            "Name": "User Role",  
            "Description": "Admin User Role",  
            "IsPredefined": true,  
            "AssignedPrivileges": [ "Login", "ConfigureManager", "ConfigureUsers",  
                "ConfigureSelf", "ConfigureComponents" ]  
        },  
        {  
            "@odata.id": "/redfish/v1/AccountService/Roles/Operator",  
            "@odata.type": "#Role.v1_1_0.Role",  
            "Id": "Operator",  
            "Name": "User Role",  
            "Description": "Operator User Role",  
            "IsPredefined": true,  
            "AssignedPrivileges": [ "Login", "ConfigureSelf", "ConfigureComponents" ]  
        },  
        {  
            "@odata.id": "/redfish/v1/AccountService/Roles/ReadOnly",  
            "@odata.type": "#Role.v1_1_0.Role",  
            "Id": "ReadOnly",  
            "Name": "User Role",  
            "Description": "ReadOnly User Role",  
            "IsPredefined": true,  
            "AssignedPrivileges": [ "Login", "ConfigureSelf" ]  
        }  
    ]  
}
```

### 7.3.3 The `$select` query parameter

The `$select` query parameter indicates that the implementation should return a subset of the resource's properties that match the `$select` expression. If a request omits the `$select` query parameter, the response returns all properties by default. The definition of this query parameter follows the [OData Protocol](#) Specification.

The `$select` expression shall not affect the resource itself.

The `$select` expression defines a comma-separated list of properties to return in the response body.

The syntax for properties in objects or properties in arrays of objects shall be the object and property names concatenated with a slash ( / ). For arrays, the expression shall not contain the array index.

An example of `$select` usage is:

```
GET /redfish/v1/Systems/1?$select=Name,SystemType>Status/State
```

For successful responses, when services execute `$select`, they shall return all requested properties of the referenced resource that are supported. Services shall omit unsupported properties requested in the `$select` query parameter from the response. If a requested property is an object, the service shall return the entire object. The `@odata.id` and `@odata.type` properties shall be in the response payload and contain the same values as if `$select` was omitted. If the `@odata.context` property is supported, it shall be in the response payload and should be in the `context property` recommended format. If the `@odata.etag` property is supported, it shall be in the response payload and contain the same values as if `$select` was omitted. Services may include [object-level messages](#) or [messages for requested properties](#) in the response payload.

For unsuccessful responses, the response body shall contain an [error response](#) regardless of the properties requested in the `$select` query parameter.

Any other supported syntax for `$select` is outside the scope of this specification.

### 7.3.4 The `$filter` query parameter

The `$filter` parameter enables a client to request a subset of the resource collection's members based on the `$filter` expression. The definition of this query parameter follows the [OData Protocol](#) Specification.

The `$filter` query parameter defines a set of properties and literals with an operator.

A literal value can be:

- A string enclosed in single quotes.
- A number.
- A boolean value.

If the literal value does not match the data type for the specified property, the service should reject `$filter` requests with the HTTP [400 Bad Request](#) status code.

The `$filter` section of the OData ABNF Components Specification contains the grammar for the allowable syntax of the `$filter` query parameter, with the additional restriction that only built-in filter operations are supported.

[Table 10](#) lists the Redfish-supported values for the `$filter` query parameter. Any other supported syntax for `$filter` is outside the scope of this specification.

**Table 10 — The \$filter query parameter options**

Value	Description	Example
<code>()</code>	Precedence grouping operator.	<code>(Status/State eq 'Enabled' and Status/Health eq 'OK') or SystemType eq 'Physical'</code>
<code>and</code>	Logical and operator.	<code>ProcessorSummary/Count eq 2 and MemorySummary/TotalSystemMemoryGiB gt 64</code>
<code>eq</code>	Equal comparison operator.	<code>ProcessorSummary/Count eq 2</code>
<code>ge</code>	Greater than or equal to comparison operator.	<code>ProcessorSummary/Count ge 2</code>
<code>gt</code>	Great than comparison operator.	<code>ProcessorSummary/Count gt 2</code>
<code>le</code>	Less than or equal to comparison operator.	<code>MemorySummary/TotalSystemMemoryGiB le 64</code>
<code>lt</code>	Less than comparison operator.	<code>MemorySummary/TotalSystemMemoryGiB lt 64</code>
<code>ne</code>	Not equal comparison operator.	<code>SystemType ne 'Physical'</code>
<code>not</code>	Logical negation operator.	<code>not (ProcessorSummary/Count eq 2)</code>
<code>or</code>	Logical or operator.	<code>ProcessorSummary/Count eq 2 or ProcessorSummary/Count eq 4</code>

When evaluating expressions, services shall use the following operator precedence:

- Grouping
- Logical negation
- Relational comparison. `gt`, `ge`, `lt`, and `le` all have equal precedence.
- Equality comparison. `eq` and `ne` both have equal precedence.
- Logical `and`
- Logical `or`

If the service receives an unsupported `$filter` query parameter, it shall reject the request and return the HTTP [501 Not Implemented](#) status code.

## 7.4 HEAD

The `HEAD` method differs from the `GET` method in that it shall not return message body information.

However, the `HEAD` method completes the same authorization checks and returns all the same meta information and status codes in the HTTP headers as a `GET` method.

Services may support the `HEAD` method to:

- Return meta information in the form of HTTP response headers.
- Verify hyperlink validity.

Services may support the `HEAD` method to verify resource accessibility.

Services shall not support any other use of the `HEAD` method.

The `HEAD` method shall be idempotent in the absence of outside changes to the resource.

Services shall reject `HEAD` requests that contain [query parameters](#). Services should return the HTTP [400 Bad Request](#) status code if provided with a query parameter in a `HEAD` request.

## 7.5 Data modification requests

---

### 7.5.1 Data modification requests overview

To create, modify, and delete resources, clients issue the following operations:

- [POST \(create\)](#)
- [PATCH \(update\)](#)
- [PUT \(replace\)](#)
- [DELETE \(delete\)](#)
- [POST \(action\)](#) on the resource

The following clauses describe the success and error response requirements common to all data modification requests.

### 7.5.2 Modification success responses

For [POST \(create\)](#) operations, the response from the service, after the create request succeeds, should be one of these responses:

- The HTTP [201 Created](#) status code. If a response body is provided, it contains the JSON representation of the newly created resource after the request has been applied.
- The HTTP [202 Accepted](#) status code with a `Location` header set to the URI of a [task monitor](#) when the processing of the request requires additional time to be completed.
  - After processing of the [task](#) is complete, the created resource may be returned in response to a request to

the task monitor URI with the HTTP `201 Created` status code.

- The HTTP `204 No Content` status code with no response body.

For **PATCH (update)**, **PUT (replace)**, and **DELETE (delete)** operations, the response from the service, after successful modification, should be one of the following responses:

- The HTTP `200 OK` status code with a body that contains the JSON representation of the targeted resource after the modification has been applied, or, for the delete operation, a representation of the deleted resource.
- The HTTP `202 Accepted` status code with a `Location` header set to the URI of a task monitor when the processing of the modification requires additional time.
  - After processing of the task is complete, the modified resource may be returned in response to a request to the task monitor URI with the HTTP `200 OK` status code.
- The HTTP `204 No Content` status code with no response body.

For details on successful responses to action requests, see [POST \(action\)](#).

### 7.5.3 Modification error responses

If the resource exists but does not support the requested operation, services shall return the HTTP `405 Method Not Allowed` status code.

Otherwise, if the service returns a client `4xx` or service `5xx` [status code](#), the service encountered an error and the resource shall not have been modified or created as a result of the operation.

## 7.6 PATCH (update)

To update a resource's properties, the service shall support the `PATCH` method.

The request body defines the changes to make to one or more properties in the resource that the request URI references. The `PATCH` request does not change any properties that are not in the request body. Services may accept a `PATCH` method with an empty JSON object, which indicates that the service should make no changes to the resource.

For resources that allow for properties to not be updated immediately, clients can perform `PATCH` requests to a designated settings resource. For more information, see the [Settings resource](#) clause.

See the [Modification success responses](#) clause for behavior when the `PATCH` operation is successful.

If supported by the service, clients can perform a conditional `PATCH` operation by specifying an `If-Match OR If-None-Match` [request header](#) that contains the ETag of the resource.

Services may reject the update on certain properties based on their own policies and, in this case, not make the

requested update. Services should not require clients to provide properties other than those the client is attempting to modify, including the `Oem` property.

A partial success of a `PATCH` operation occurs when a modification request for multiple properties results in at least one property updated successfully, but one or more properties could not be updated. In these cases, the service shall return the HTTP `200 OK` status code and a resource representation with **extended information** that lists the properties that could not be updated. Examples include:

- A property is read-only, unknown, or unsupported.
- A service-side error occurred, such as a write failure for an EEPROM.

If all properties in the update request are read-only, unknown, or unsupported, but the resource can be updated, the service shall return the HTTP `400 Bad Request` status code and an **error response** with messages that show the non-updatable properties.

The service shall ignore OData annotations in the request body, such as the **resource identifier**, **type**, and **ETag** properties, except for the conditions listed below. If the update request only contains OData annotations, the service shall return the HTTP `400 Bad Request` status code and an **error response** with the `Nooperation` message from the Base Message Registry (preferred) or a modification success response, except for the conditions listed below.

- Writable **reference properties**.
- In **deep operations** when specifying subordinate resources.

In the absence of outside changes to the resource, the `PATCH` operation should be idempotent, although the original `ETag` value may no longer match.

## 7.7 PATCH on array properties

---

The **Array properties** clause describes the three styles of array properties in a resource.

Within a `PATCH` request, the service shall accept `null` to remove an element, and accept an empty object `{}` to leave an element unchanged. Array properties that use the fixed or variable length style remove those elements, while array properties that use the rigid style replace removed elements with `null` elements. A service may indicate the maximum size of an array by padding `null` elements at the end of the array sequence.

When processing a `PATCH` request, the order of operations shall be:

- Modifications
- Deletions
- Additions

A `PATCH` request with fewer elements than in the current array shall remove the remaining elements of the array.

For example, a fixed length-style `Flavors` array indicates that the service supports a maximum of six elements, by padding the array with `null` elements, with four populated.

```
{  
  "Flavors": ["Chocolate", "Vanilla", "Mango", "Strawberry", null, null]  
}
```

A client could issue the following `PATCH` request to remove `Vanilla`, replace `Strawberry` with `Cherry`, and add `Coffee` and `Banana` to the array, while leaving the other elements unchanged.

```
{  
  "Flavors": [{}, null, {}, "Cherry", "Coffee", "Banana"]  
}
```

After the `PATCH` operation, the resulting array is:

```
{  
  "Flavors": ["Chocolate", "Mango", "Cherry", "Coffee", "Banana", null]  
}
```

## 7.8 PUT (replace)

To completely replace a resource, services may support the `PUT` method. The service may add properties to the response resource that the client omits from the request body, the resource definition requires, or the service normally supplies.

The `PUT` operation should be idempotent in the absence of outside changes to the resource, with the possible exception that the operation might change ETag values.

See the [Modification success responses](#) clause for behavior when the `PUT` operation is successful.

If supported by the service, clients can perform a conditional `PUT` operation by specifying an [`If-Match`](#) OR [`If-None-Match`](#) request header that contains the ETag of the resource.

Services may reject requests that do not include properties that the resource definition (schema) requires.

## 7.9 POST (create)

To create a resource, services shall support the `POST` method on resource collections.

The `POST` request is submitted to the resource collection to which the new resource will belong. See the [Modification success responses](#) clause for behavior when the `POST` operation is successful.

The body of the create request contains a representation of the object to create. The service may ignore any service-controlled properties, such as `Id`, which would force the service to overwrite those properties. Additionally, the service shall set the `Location` header in the response to the URI of the new resource.

- Submitting a `POST` request to a resource collection is equivalent to submitting the same request to the `Members` property of that resource collection. Services that support the addition of `Members` to a resource collection shall support both forms.
  - For example, if a client adds a member to the resource collection at `/redfish/v1/EventService/Subscriptions`, it can perform a `POST` request to either `/redfish/v1/EventService/Subscriptions` or `/redfish/v1/EventService/Subscriptions/Members`.
- The `POST` operation shall not be idempotent.
- Services may allow the inclusion of the `@Redfish.OperationApplyTime` property in the request body. See [Operation apply time](#).
- Services should return the HTTP `400 Bad Request` status code for requests containing properties with the value `null`.
- Services should not require clients to provide properties that are not annotated as [required on create](#).

## 7.10 DELETE (delete)

To remove a resource, the service shall support the `DELETE` method. Resources [subordinate](#) to the resource removed by a `DELETE` method are typically removed, as the contents of subordinate resources are dependent on the [parent resource](#). In some cases, related resources may also be relocated in the resource tree based on their definition and usage. Other resources in the resource tree may also be removed or incur side effects of a resource removal.

See the [Modification success responses](#) clause for behavior when the `DELETE` operation is successful.

- If the resource was already deleted, the service may return the HTTP `404 Not Found` status code or a [success code](#).
- The service may allow the inclusion of the `@Redfish.OperationApplyTime` property in the request body. See [Operation apply time](#).

## 7.11 POST (action)

Services shall support the `POST` method as a way for clients to send actions to resources.

- The `POST` operation may not be idempotent.
- Services may allow the inclusion of the `@Redfish.OperationApplyTime` property in the request body. See [Operation apply time](#).

### Operation apply time.

- Services should not require clients to provide parameters that are not annotated as [required](#).

To request actions on a resource, send the HTTP `POST` method to the URI of the action. The `target` property in the resource's [Actions](#) property shall contain the URI of the action. The URI of the action shall be in the format:

```
<ResourceUri>/Actions/<QualifiedActionName>
```

where

- `<ResourceUri>` is the URI of the resource that supports the action.
- `Actions` is the name of the property that contains the actions for a resource, as defined by this specification.
- `<QualifiedActionName>` is the qualified name of the action. Includes the [resource type](#).

To determine the available [actions](#) and the [valid parameter values](#) for those actions, clients can query a resource directly.

Clients provide parameters for the action as a JSON object within the request body of the `POST` operation. For information about the structure of the request and required parameters, see the [Actions property](#) clause. Some parameter information may require that the client examine the [Redfish schema](#) that corresponds to the resource.

If the action request does not contain all required parameters, the service shall return the HTTP [400 Bad Request](#) status code. If the action request contains unsupported parameters, the service shall ignore the unsupported parameters or return the HTTP [400 Bad Request](#) status code. If an action does not have any required parameters, the service shall accept an empty JSON object in the HTTP body for the action request.

**Note:** Previous versions of the specification allowed a service to reject action requests with an empty JSON object in the HTTP body if the client is not providing parameters for the action request. Services implementing version v1.20.0 or higher are required to accept an empty JSON object in an action request when the client is not providing parameters for the action request.

[Table 11](#) describes the HTTP status codes and additional information that the service shall return a response to a successful `POST` (action) request:

**Table 11 — POST (action) status codes**

To indicate	HTTP status code	Additional information
Success, and the action's schema definition does not contain an <a href="#">action response</a> .	<a href="#">200 OK</a>	An <a href="#">error response</a> , with a message that indicates success or any additional relevant messages. If the action was successfully processed and completed without errors, warnings, or other notifications for the client, the service should return the <code>Success</code> message from the Base Message Registry in the <code>code</code> property in the response body.

To indicate	HTTP status code	Additional information
Success, and the action's schema definition contains an <a href="#">action response</a> .	<a href="#">200 OK</a>	The response body conforms to the <a href="#">action response</a> defined in the schema.
A new resource was created, and the action's schema definition does not contain an <a href="#">action response</a> .	<a href="#">201 Created</a>	A <a href="#">Location</a> response header set to the URI of the created resource. An <a href="#">error response</a> , with a message that indicates success or any additional relevant messages. If the action was successfully processed and completed without errors, warnings, or other notifications for the client, the service should return the <a href="#">Success</a> message or <a href="#">Created</a> message from the Base Message Registry in the <code>code</code> property in the response body.
A new resource was created, and the action's schema definition contains an <a href="#">action response</a> .	<a href="#">201 Created</a>	A <a href="#">Location</a> response header set to the URI of the created resource. The response body conforms to the <a href="#">action response</a> defined in the schema.
Additional time is required to process.	<a href="#">202 Accepted</a>	A <a href="#">Location</a> response header set to the URI of a <a href="#">task monitor</a> .
Success, and the action's schema definition does not contain an <a href="#">action response</a> .	<a href="#">204 No Content</a>	No response body.

If an action requested by the client has no effect, such as a reset of a `ComputerSystem` where the `ResetType` parameter is set to `On` and the `ComputerSystem` is already `On`, the service should respond with the HTTP [200 OK](#) status code and return the `NoOperation` message from the Base Message Registry.

If an error was detected and the action request was not processed, the service shall return an HTTP `4xx` or HTTP `5xx` [status code](#). The response body, if provided, shall contain an [error response](#) that describes the error or errors.

Example successful action response:

```
{
  "error": {
    "code": "Base.1.8.Success",
    "message": "Successfully Completed Request",
    "@Message.ExtendedInfo": [
      {
        "@odata.type": "#Message.v1_1_1.Message",
        "MessageId": "Base.1.8.Success",
        "Message": "Successfully Completed Request",
        "Severity": "OK",
        "MessageSeverity": "OK",
        "Resolution": "None"
      }
    ]
  }
}
```

## 7.12 Operation apply time

Services may accept the `@Redfish.OperationApplyTime` annotation in the following request bodies:

- `POST` (create)
- `DELETE` (delete)
- `POST` (action)
- The JSON part for multipart HTTP POST operations, such as with the [multipart HTTP push update](#) in the update service.

This annotation enables the client to control when an operation is carried out.

For example, if the client wants to delete a particular `Volume` resource, but can only safely do so when a reset occurs, the client can use this annotation to instruct the service to delete the `Volume` on the next reset.

If multiple operations are pending, the service shall process them in the order in which the service receives them.

Services that support the `@Redfish.OperationApplyTime` annotation for create operations on a resource collection and delete operations on members of a resource collection shall include the `@Redfish.OperationApplyTimeSupport` response annotation for the resource collection.

The following example is a response for a resource collection that supports the `@Redfish.OperationApplyTime` annotation in requests to create new members and delete existing members:

```
{  
    "@odata.id": "/redfish/v1/Systems/1/Storage/SATAEmbedded/Volumes",  
    "@odata.type": "#VolumeCollection.VolumeCollection",  
    "Name": "Storage Volume Collection",  
    "Description": "Storage Volume Collection",  
    "Members@odata.count": 2,  
    "Members": [  
        {"@odata.id": "/redfish/v1/Systems/1/Storage/SATAEmbedded/Volumes/1"},  
        {"@odata.id": "/redfish/v1/Systems/1/Storage/SATAEmbedded/Volumes/2"}  
    ],  
    "@Redfish.OperationApplyTimeSupport": {  
        "@odata.type": "#Settings.v1_3_3.OperationApplyTimeSupport",  
        "SupportedValues": ["Immediate", "OnReset"]  
    }  
}
```

In the previous example, a client can annotate their create request body on the `VolumeCollection` itself, or a delete operation on the `Volumes` within the `VolumeCollection`.

The following sample request deletes a `Volume` on the next reset:

```
DELETE /redfish/v1/Systems/1/Storage/SATAEmbedded/Volumes/2 HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
    "@Redfish.OperationApplyTime": "OnReset"
}
```

Services that support the `@Redfish.OperationApplyTime` annotation for an action shall include the `@Redfish.OperationApplyTimeSupport` response annotation for the action.

The following example is a response for a `ComputerSystem` resource that supports the `@Redfish.OperationApplyTime` annotation in the reset action request:

```
{
    "@odata.id": "/redfish/v1/Systems/1",
    "@odata.type": "#ComputerSystem.v1_5_0.ComputerSystem",
    "Actions": {
        "#ComputerSystem.Reset": {
            "target": "/redfish/v1/Systems/1/Actions/ComputerSystem.Reset",
            "ResetType@Redfish.AllowableValues": ["On", "ForceOff", "ForceRestart",
                "Nmi", "ForceOn", "PushPowerButton"],
            "@Redfish.OperationApplyTimeSupport": {
                "@odata.type": "#Settings.v1_3_3.OperationApplyTimeSupport",
                "SupportedValues": ["Immediate", "AtMaintenanceWindowStart"],
                "MaintenanceWindowStartTime": "2017-05-03T23:12:37-05:00",
                "MaintenanceWindowDurationInSeconds": 600,
                "MaintenanceWindowResource": {
                    "@odata.id": "/redfish/v1/Systems/1"
                }
            }
        }
    },
    ...
}
```

In the previous example, a client can annotate their reset action request body on the `ComputerSystem` in the payload.

The following sample request completes a reset at the start of the next maintenance window:

```
POST /redfish/v1/Systems/1/Actions/ComputerSystem.Reset HTTP/1.1
Content-Type: application/json; charset=utf-8
```

```
Content-Length: <computed length>
OData-Version: 4.0

{
    "ResetType": "ForceRestart",
    "@Redfish.OperationApplyTime": "AtMaintenanceWindowStart"
}
```

Services that support the `@Redfish.OperationApplyTime` annotation for the JSON part of a multipart HTTP POST operation shall include the `@Redfish.OperationApplyTimeSupport` response annotation for the property that specifies the URI of the multipart POST operation.

The following example is a response for an `UpdateService` resource that supports the `@Redfish.OperationApplyTime` annotation in the JSON part of a multipart request:

```
{
    "@odata.id": "/redfish/v1/UpdateService",
    "@odata.type": "#UpdateService.v1_8_0.UpdateService",
    "MultipartHttpPushUri": "/redfish/v1/UpdateService/update-multipart",
    "MultipartHttpPushUri@Redfish.OperationApplyTimeSupport": {
        "@odata.type": "#Settings.v1_3_3.OperationApplyTimeSupport",
        "SupportedValues": ["Immediate", "AtMaintenanceWindowStart"],
        "MaintenanceWindowStartTime": "2017-05-03T23:12:37-05:00",
        "MaintenanceWindowDurationInSeconds": 600,
        "MaintenanceWindowResource": {
            "@odata.id": "/redfish/v1/UpdateService"
        }
    },
    ...
}
```

Services that support the `@Redfish.OperationApplyTime` annotation for a resource collection, action, or multipart HTTP POST operation shall create a `task`, and respond with the HTTP `202 Accepted` status code with a `Location` header set to the URI of a `task monitor`, if the client's request body contains `@Redfish.OperationApplyTime` in the request.

The `Settings` Redfish schema defines the structure of the `@Redfish.OperationApplyTimeSupport` object and the `@Redfish.OperationApplyTime` annotation value.

## 7.13 Deep operations

Implementations may support operations that modify the current resource as well as subordinate resources. These operations are known as deep operations. They give the client the ability to modify more than one resource with a single operation.

[Table 12](#) describes the types of deep operations that this specification defines:

**Table 12 — Deep operations**

Operation	Description	Example
Deep PATCH (update)	Modify a resource and one or more subordinate resources.	Modify a ComputerSystem resource as well as subordinate Storage and NetworkInterface resources.
Deep POST (create)	Create one or more resources in a resource collection and optionally subordinate resources.	Create multiple ManagerAccount resources.

- Services that support deep PATCH for updating resources shall set the value of the DeepPATCH property in the DeepOperations property in the ProtocolFeaturesSupported property within the service root to true .
- Services that support deep POST for creating resources shall set the value of the DeepPOST property in the DeepOperations property in the ProtocolFeaturesSupported property within the service root to true .
- The Members property in resource collections shall not be removed when using a deep PATCH .
- Action URLs shall not support deep POST operations.
- If the service supports deep operations, the MaxLevels property in the DeepOperations property in the ProtocolFeaturesSupported property in the service root shall indicate the maximum number of levels that the service supports for deep operations.
- To request deep operations on a resource, send the HTTP method to the deep operation URI of the resource. The URI for deep operations on any resource shall be in the format: <ResourceUri>.Deep .
- The schema used for validating the root level of the request body shall be the schema of the resource in the resource URI.
  - The subordinate resources included in the request body shall be validated against their corresponding schema.

The body of deep operations contains the resource being modified as well as the subordinate resources being modified. This resource can be a collection or a single instance. These resources could be subordinate resources, subordinate resource collections, or subordinate members of resource collections. The client can omit properties from the request such as those it does not want to modify or that the service controls. Requests that include references to multiple instances, such as members of a collection, shall include the Members property as part of the request body.

To determine which members of subordinate resource collections are to be modified by a deep PATCH , services shall use the @odata.id property provided by the client to identify the member of the resource collection to be modified.

Clients may provide the @odata.etag property in subordinate resources being modified by a deep PATCH . If the request contains the If-Match or If-None-Match header, the service shall compare the ETag in the request header with the ETag of the resource specified by the URI. If this check passes, the operation can proceed using the @odata.etag values contained in the body of the subordinate resources. In this case, the operation on each subordinate resource shall be completed independently, where some subordinate values that pass the condition

check proceed and the resources that fail do not proceed. In this case, annotated [extended information](#) shall be included in the subordinate resource representation of the response.

Failure semantics for deep operations are similar to that of other [operations of similar type](#). If any properties in a deep `PATCH` operation succeeded, the result is a `200 OK` with the results returned in the response, and the service should include extended information indicating warnings or errors. For a deep `POST` operation, if any member of the collection was created then a `201 Created` shall be returned, and any members that were not created should have extended information in their place holders with sufficient identifying information, such as returning all of the properties provided in the `POST` request body for that member, as well as extended information indicating why the creation was not successful. When sending a deep `POST` request, the value of the `Location header` shall be that of one of the URLs created and should be that of one of the least subordinate URLs, such as that of a `ComputerSystem` resource and not one of the devices subordinate to the `ComputerSystem` resource.

If the request body for a deep operation contains resources that are not modifiable, but no modifications are requested for those resources, services shall not treat this as a modification request for those resources. For example, if the service root is not modifiable, meaning `PATCH` is not accepted on the resource, a client is allowed to provide the service root in the deep operation request body if there are no modifications to the service root.

Deep `POST` shall not be allowed on the `SessionCollection` resource.

The following deep `PATCH` example modifies two members of the `RoleCollection` resource:

```
PATCH /redfish/v1/AccountService/Roles.Deep HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
  "Members": [
    {
      "@odata.id": "/redfish/v1/AccountService/Roles/OperatorRestricted",
      "AssignedPrivileges": [ "Login", "ConfigureComponents" ]
    },
    {
      "@odata.id": "/redfish/v1/AccountService/Roles/ReadOnlyRestricted",
      "AssignedPrivileges": [ "Login" ]
    }
  ]
}
```

The following deep `POST` example creates two members in the `RoleCollection` resource:

```
POST /redfish/v1/AccountService/Roles.Deep HTTP/1.1
Content-Type: application/json; charset=utf-8
```

```

Content-Length: <computed length>
OData-Version: 4.0

{
  "Members": [
    {
      "RoleId": "OperatorRestricted",
      "AssignedPrivileges": [ "Login", "ConfigureComponents" ]
    },
    {
      "RoleId": "ReadOnlyRestricted",
      "AssignedPrivileges": [ "Login" ]
    }
  ]
}

```

The following deep `POST` example creates one member in the `EventDestinationCollection` resource and creates a `Certificate` resource subordinate to the new `EventDestination` resource:

```

POST /redfish/v1/EventService/Subscriptions.Deep HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
  "Members": [
    {
      "Destination": "https://listener1.contoso.com",
      "Protocol": "Redfish",
      "Certificates": {
        "Members": [
          {
            "CertificateType": "PEM",
            "CertificateString": "----- BEGIN CERTIFICATE ...."
          }
        ]
      }
    }
  ]
}

```

The following deep `PATCH` example modifies the asset tag and BIOS settings of a `ComputerSystem` resource:

```

PATCH /redfish/v1/Systems/47832.Deep HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>

```

```
OData-Version: 4.0

{
    "AssetTag": "Inventory Tag 12394783431",
    "Bios": {
        "@odata.id": "/redfish/v1/Systems/47832/Bios",
        "@Redfish.Settings": {
            "@odata.type": "#Settings.v1_3_3.Settings",
            "SettingsObject": {
                "@odata.id": "/redfish/v1/Systems/47832/Bios/SD",
                "Attributes": {
                    "AdminPhone": "(123) 456-789",
                    "BootMode": "Uefi"
                }
            }
        }
    }
}
```

The following example shows a deep `PATCH` with ETags in the request:

```
PATCH /redfish/v1/AccountService/Roles.Deep HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
If-Match: <Collection ETag>
OData-Version: 4.0

{
    "Members": [
        {
            "@odata.id": "/redfish/v1/AccountService/Roles/OperatorRestricted",
            "@odata.etag": "W/\"ABCDEFG\"",
            "AssignedPrivileges": [ "Login", "ConfigureComponents" ]
        },
        {
            "@odata.id": "/redfish/v1/AccountService/Roles/ReadOnlyRestricted",
            "@odata.etag": "W/\"ABCDEFG\"",
            "AssignedPrivileges": [ "Login" ]
        }
    ]
}
```

The following example response shows a partial failure of a deep `PATCH` where the ETag provided in the request for the `Role` resource named `ReadOnlyRestricted` was incorrect:

```
HTTP/1.1 200 OK
```

```
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
ETag: <Resource collection ETag>
OData-Version: 4.0

{
  "Members": [
    {
      "@odata.id": "/redfish/v1/AccountService/Roles/OperatorRestricted",
      "@odata.etag": "W/\"ABCDEFG\"",
      "AssignedPrivileges": [ "Login", "ConfigureComponents" ]
    },
    {
      "@odata.id": "/redfish/v1/AccountService/Roles/ReadOnlyRestricted",
      "@Message.ExtendedInfo": [
        {
          "@odata.type": "#Message.v1_1_1.Message",
          "MessageId": "Base.1.8.PreconditionFailed",
          "RelatedProperties": [ "/AssignedPrivileges" ]
        }
      ]
    }
  ]
}
```

# 8 Service responses

This clause describes the responses that Redfish services can return to clients.

## 8.1 Response headers

HTTP defines headers for use in response messages. [Table 13](#) defines those headers and their requirements for Redfish services:

- Redfish services shall return the HTTP-defined headers if the **Required** column contains **Yes**.
- Redfish services should return the HTTP-defined headers if the **Required** column contains **No**.
- Redfish clients shall be able to both understand and process all the HTTP-defined headers.

**Table 13 — Response headers**

Header	Required	Supported values	Description
<code>Access-Control-Allow-Origin</code>	No	<a href="#">Fetch Living Standard, 3.2.3. HTTP responses</a>	Prevents or allows requests based on originating domain. Prevents CSRF attacks.
<code>Allow</code>	Yes	<code>POST</code> , <code>PUT</code> , <code>PATCH</code> , <code>DELETE</code> , <code>GET</code> , <code>HEAD</code>	Shall be returned with the HTTP <a href="#">405</a> (Method Not Allowed) status code to indicate the valid methods for the request URI. Shall be returned with any <code>GET</code> or <code>HEAD</code> operation to indicate the other allowable operations for this resource.
<code>Cache-Control</code>	Yes	<a href="#">RFC7234</a>	Shall be supported and indicates whether a response can or cannot be cached.
<code>Content-Encoding</code>	No	<a href="#">RFC7231</a>	Encoding used to compress the message body.
<code>Content-Length</code>	No	<a href="#">RFC7231</a>	Size of the message body. An optional means of indicating size of the body uses <code>Transfer-Encoding: chunked</code> , that does not use the <code>Content-Length</code> header. If a service does not support <code>Transfer-Encoding</code> and needs <code>Content-Length</code> instead, the service shall respond with the HTTP <a href="#">411 Length Required</a> status code.

Header	Required	Supported values	Description
<code>Content-Type</code>	Yes	<a href="#">RFC7231</a>	<p>The message body's representation type.</p> <p>Services shall specify a <code>Content-Type</code> of <code>application/json</code> when returning resources as JSON.</p> <p>Services shall specify a <code>Content-Type</code> of <code>application/xml</code> when returning metadata as XML.</p> <p>Services shall specify a <code>Content-Type</code> of <code>application/yaml</code> or <code>application/vnd.oai.openapi</code> when returning OpenAPI schema as YAML.</p> <p>Services shall specify a <code>Content-Type</code> of <code>text/event-stream</code> when returning an <a href="#">SSE</a> stream.</p> <p><code>;charset=utf-8</code> shall be appended to the <code>Content-Type</code> if specified in the chosen media-type in the <code>Accept</code> header for the request.</p>
<code>ETag</code>	Conditional	<a href="#">RFC7232</a>	An identifier for a specific version of a resource, often a message digest. The <code>ETag</code> header shall be included on responses to <code>GET</code> s of <code>ManagerAccount</code> resources.
<code>Link</code>	Yes	<a href="#">RFC8288</a>	<code>Link</code> headers shall be returned, as described in the <a href="#">Link headers</a> clause.
<code>Location</code>	Conditional	<a href="#">RFC7231</a>	URI of a newly created resource. Shall be returned upon creation of a resource. <code>Location</code> and <code>X-Auth-Token</code> shall be included on responses that create user sessions.
<code>Max-Forwards</code>	No	<a href="#">RFC7231</a>	Limits gateway and proxy hops. Prevents messages from remaining in the network indefinitely.
<code>OData-Version</code>	Yes	4.0	OData version of the payload to which the response conforms.
<code>Retry-After</code>	No	<a href="#">RFC7231</a> , Section 7.1.3	Informs a client how long to wait before requesting the <code>task</code> information again.
<code>Server</code>	No	<a href="#">RFC7231</a>	<p>A product token and its version. Multiple product tokens may be listed.</p> <p><b>Note:</b> Previous versions of the Specification marked this header as required. This has been changed because no use cases for requiring it have been identified.</p>
<code>Via</code>	No	<a href="#">RFC7230</a>	Defines the network hierarchy and recognizes message loops. Each pass inserts its own <code>via</code> header.

Header	Required	Supported values	Description
WWW-Authenticate	Conditional	RFC7617	Provides a challenge to clients, such as directing a client to perform <a href="#">HTTP Basic authentication</a> , when authentication headers in the request are missing or invalid. Services shall return this header if the <code>HTTPBasicAuth</code> property in the <code>AccountService</code> resource contains <code>Enabled</code> or is not implemented.
X-Auth-Token	Yes	Opaque encoded octet strings	Contains the authentication token for user sessions. The token value shall be indistinguishable from random.

## 8.2 Link header

The `Link` header provides metadata information on the accessed resource in response to a `HEAD` or `GET` request. The metadata information can include hyperlinks from the resource and JSON Schemas that describe the resource.

The following example shows the `Link` headers for a `ManagerAccount` with an `Administrator` role, in addition to a `Settings` annotation:

```
Link: </redfish/v1/AccountService/Roles/Administrator>; path=/Links/Role
Link: <http://redfish.dmtf.org/schemas/Settings.json>
Link: </redfish/v1/JsonSchemas/ManagerAccount.v1_0_2.json>; rel=describedby
```

- The first `Link` header is an example of a hyperlink that comes from the resource. It describes hyperlinks within the resource. This type of header is outside the scope of this specification.
- The second `Link` header is an example of an annotation `Link` header as it references the JSON Schema that describes the annotation and does not have `rel=describedby`. This example references the public copy of the annotation on DMTF's Redfish schema repository.
- The third `Link` header is an example for the JSON Schema that describes the actual resource.
  - Note that the URL can reference an unversioned JSON Schema because the `@odata.type` in the resource indicates the appropriate version, or reference the versioned JSON Schema, which according to previous normative statements need to match the version in the `@odata.type` property of the resource.

A `Link` header containing `rel=describedby` shall be returned on `GET` and `HEAD` requests for Redfish resources. If the referenced JSON Schema is a versioned schema, it shall match the version contained in the value of the `@odata.type` property returned in this resource.

A `Link` header satisfying annotations should be returned on `GET` and `HEAD` requests for Redfish resources.

## 8.3 Status codes

HTTP defines status codes that appear in responses. The status codes themselves provide general information about how the request was processed, such as whether the request was successful, if the client provided bad information, or the service encountered an error when processing the request.

- When the service returns a status code in the `4xx` or `5xx` range, services should return an [extended error response](#) in the response body to provide the client more meaningful and deterministic error semantics.
- When the service returns a status code in the `2xx` range and the response contains a representation of a resource, services may use [extended information](#) to convey additional information about the resource.
- Extended error messages shall not provide privileged information when authentication failures occur.

**Note:** For security implications of extended errors, See [Security details](#).

**Table 14** lists HTTP status codes that have meaning or usage defined for a Redfish service, or are otherwise referenced by this specification. Other codes may be returned by the service as appropriate, and their usage is implementation-specific. For usage and additional requirements imposed by this specification, see the **Description** column.

- Clients shall understand and be able to process the HTTP-defined status codes in **Table 14** and constrained by additional requirements defined by this specification.
- Services shall respond with the HTTP-defined status codes in **Table 14** and constrained by additional requirements in the **Description** column.
- If no other status code in the `4xx` range is appropriate for client-side errors, the default status code should be the HTTP [400 Bad Request](#) status code.
- If no other status code in the `5xx` range is appropriate for service-side errors, the default status code should be the HTTP [500 Internal Server Error](#) status code.

**Table 14 — HTTP status codes**

HTTP status code	Description
<code>200 OK</code>	Request completed successfully and includes a representation in its body.
<code>201 Created</code>	Request to create a resource completed successfully. The <a href="#">Location header</a> shall be set to the canonical URI for the newly created resource. For <a href="#">POST (create) requests</a> , the response body may include a representation of the newly created resource. For <a href="#">POST (action) requests</a> , the response body shall include the action response.
<code>202 Accepted</code>	Request has been accepted for processing but the processing has not been completed. The <a href="#">Location header</a> shall be set to the URI of a <a href="#">task monitor</a> that can later be queried to determine the status of the operation. If a response body is provided, it shall contain a representation of the <a href="#">Task</a> resource.

HTTP status code	Description
204 No Content	Request succeeded, but no response body is provided.
301 Moved Permanently	Requested resource resides under a different URI.
302 Found	Requested resource resides temporarily under a different URI.
304 Not Modified	Service has made a conditional <code>GET</code> request where access is allowed but the resource content has not changed. Certain request headers, such as <code>If-None-Match</code> , initiate conditional requests to save network bandwidth if no change has occurred. See HTTP 1.1, sections 14.25 and 14.26.
400 Bad Request	Request could not be processed because it contains invalid information, such as an invalid input field, or is missing a required value. The response body shall return an extended error as defined in the <a href="#">Error responses</a> clause.
401 Unauthorized	Authentication credentials included with this request are missing or invalid. Additional details are described in the <a href="#">Sensitive data</a> clause.
403 Forbidden	Service recognized the credentials in the request but those credentials do not possess authorization to complete this request. This code is also returned when the user credentials provided need to be changed before access to the service can be granted. For details, see the <a href="#">Security details</a> clause.
404 Not Found	Request specified a URI of a resource that does not exist. Additional details are described in the <a href="#">Sensitive data</a> clause.
405 Method Not Allowed	HTTP verb in the request, such as <code>DELETE</code> , <code>GET</code> , <code>HEAD</code> , <code>POST</code> , <code>PUT</code> , or <code>PATCH</code> , is not supported for this request URI. The response shall include an <code>Allow</code> header that provides a list of methods that the resource identified by the URI in the client request supports. Additional details are described in the <a href="#">Sensitive data</a> clause.
406 Not Acceptable	Accept header was specified in the request and the resource identified by this request cannot generate a representation that corresponds to one of the media types in the Accept header.
409 Conflict	Creation or update request could not be completed because it would cause a conflict in the current state of the resources that the platform supports. For example, a conflict occurred due to an attempt to set multiple properties that work in a linked manner by using incompatible values.
410 Gone	Requested resource is no longer available at the service and no forwarding address is known. This condition is expected to be considered permanent. Clients with hyperlink editing capabilities should delete references to the URI in the client request after user approval. If the service does not know or cannot determine whether the condition is permanent, client should use the HTTP <code>404 Not Found</code> status code. This response is cacheable unless otherwise indicated.
411 Length Required	Request did not use the <code>Content-Length</code> header to specify the length of its content but perhaps used the <code>Transfer-Encoding: chunked</code> header instead. The addressed resource requires the <code>Content-Length</code> header.
412 Precondition Failed	Precondition check, such as check of the <code>OData-Version</code> , <code>If-Match</code> , or <code>If-None-Match</code> header, failed.

HTTP status code	Description
413 Payload Too Large	Request payload, or a part in a multipart request, is larger than the maximum size the service supports.
415 Unsupported Media Type	Request specifies a <code>Content-Type</code> for the body that is not supported.
428 Precondition Required	Request did not provide the required precondition, such as an <code>If-Match</code> OR <code>If-None-Match</code> header.
431 Request Header Field Too Large	Service is unwilling to process the request because either an individual header field or the collection of all header fields are too large.
500 Internal Server Error	Service encountered an unexpected condition that prevented it from fulfilling the request. The response body shall return an extended error as defined in the <a href="#">Error responses</a> clause.
501 Not Implemented	Service does not currently support the functionality required to fulfill the request. This response is appropriate when the service does not recognize the request method and cannot support the method for any resource.
503 Service Unavailable	Service currently cannot handle the request due to temporary overloading or maintenance of the service. A service may use this response to indicate that the request URI is valid but the service is performing initialization or other maintenance on the resource. A service may also use this response to indicate that the service itself is undergoing maintenance, such as finishing initialization steps after reboot of the service.
507 Insufficient Storage	Service cannot build the response for the client due to the size of the response.

## 8.4 OData metadata responses

### 8.4.1 OData metadata responses overview

OData metadata describes resources, resource collections, capabilities, and service-dependent behavior to generic OData consumers with no specific understanding of this specification. Clients are not required to request metadata if they already have sufficient understanding of the target service. For example, clients are not required to request metadata to request and interpret a JSON representation of a resource that this specification defines.

A client can access the OData metadata at the `/redfish/v1/$metadata` URI.

A client can access the OData service document at the `/redfish/v1/odata` URI.

### 8.4.2 OData \$metadata

The OData metadata describes top-level service resources and resource types according to [OData Common Schema Definition Language](#). The OData metadata is represented as an XML document with an `Edmx` root element in the `http://docs.oasis-open.org/odata/ns/edmx` namespace with an OData version attribute set to `4.0`.

The service shall use the `application/xml` or `application/xml; charset=utf-8` MIME types to return the OData metadata document as an XML document.

```
<edmx:Edmx xmlns:edmx="http://docs.oasis-open.org/odata/ns/edmx" Version="4.0">
  <!-- edmx:Reference and edmx:Schema elements go here -->
</edmx:Edmx>
```

#### 8.4.2.1 Referencing other schemas

The OData metadata should include the namespaces for each of the Redfish resource types, along with the `RedfishExtensions.v1_0_0` namespace. Dynamic clients that reference the OData metadata document leverage schema definitions that are referenced to understand the definitions of the resources in the service. However, there are cases where it might not be practical to maintain an accurate document, such as when resources are dynamically discovered by the service through devices that support Redfish Device Enablement.

These references shall use either:

- An absolute URI for the Redfish schema definitions, such as on `http://redfish.dmtf.org/schemas` or `http://developers.contoso.org/schemas`.
- A relative URI to a local copy of the Redfish schema. See the [Redfish-defined URIs and relative reference rules](#) clause for recommended URI patterns.

```
<edmx:Reference Uri="http://redfish.dmtf.org/schemas/v1/ServiceRoot_v1.xml">
  <edmx:Include Namespace="ServiceRoot"/>
  <edmx:Include Namespace="ServiceRoot.v1_0_0"/>
</edmx:Reference>

...

<edmx:Reference Uri="http://redfish.dmtf.org/schemas/v1/VirtualMedia_v1.xml">
  <edmx:Include Namespace="VirtualMedia"/>
  <edmx:Include Namespace="VirtualMedia.v1_0_0"/>
</edmx:Reference>
<edmx:Reference Uri="http://redfish.dmtf.org/schemas/v1/RedfishExtensions_v1.xml">
  <edmx:Include Namespace="RedfishExtensions.v1_0_0" Alias="Redfish"/>
</edmx:Reference>
```

The service's [OData metadata document](#) shall include an `EntityContainer` that defines the top-level resources and resource collections.

#### 8.4.2.2 Referencing OEM extensions

The OData metadata document may reference additional schema documents that describe OEM-specific extensions that the service uses.

For example, the OData metadata document may reference custom types for additional resource collections.

```
<edmx:Reference Uri="http://contoso.org/Schema/CustomTypes">
  <edmx:Include Namespace="CustomTypes"/>
</edmx:Reference>
```

#### 8.4.3 OData service document

The OData service document serves as a top-level entry point for generic OData clients. More information about the OData service document can be found in the [OData JSON Format Specification](#).

```
{
  "@odata.context": "/redfish/v1/$metadata",
  "value": [
    {
      "name": "Service",
      "kind": "Singleton",
      "url": "/redfish/v1/"
    },
    {
      "name": "Systems",
      "kind": "Singleton",
      "url": "/redfish/v1/Systems"
    },
    ...
  ]
}
```

The service shall use the `application/json` MIME type to return the OData service document as a JSON object.

The JSON object shall contain the `@odata.context` context property set to `/redfish/v1/$metadata`.

The JSON object shall include a `value` property set to a JSON array that contains an entry for the `service root` and each resource that is a direct child of the service root.

[Table 15](#) describes the properties that each JSON object entry includes:

**Table 15 — JSON object properties**

Property	Description
name	User-friendly resource name of the resource.
kind	Type of resource. Value is <code>Singleton</code> for all cases defined by Redfish.
url	Relative URL for the top-level resource.

## 8.5 Resource responses

Services use the `application/json` MIME type to return resources and resource collections as JSON payloads. A service shall not break responses for a single resource into multiple results.

The format of these payloads is defined by the Redfish schema. For rules about the Redfish schema and how it maps to JSON payloads, see the [Data model](#) and [Schema definition languages](#) clauses.

## 8.6 Error responses

HTTP status codes often do not provide enough information to enable deterministic error semantics. For example, if a client makes a `PATCH` call and some properties do not match while others are not supported, the HTTP [400 Bad Request](#) status code does not tell the client which values are in error. Error responses provide the client more meaningful and deterministic error semantics.

To provide the client with as much information about the error as possible, a Redfish service may provide multiple error responses in the HTTP response. Additionally, the service may provide Redfish standardized errors, OEM-defined errors, or both, depending on the implementation's ability to convey the most useful information about the underlying error.

[Table 16](#) describes the properties in the extended error response, which is a single JSON object:

**Table 16 — Error properties**

Property	Description
code	String. Defines a <code>MessageId</code> from the message registry. See the <a href="#">MessageId format clause</a> for the format of <code>MessageId</code> .
message	Displays a human-readable error message that corresponds to the message in the message registry.
<code>@Message.ExtendedInfo</code>	Displays an array of <a href="#">message objects</a> . Describes one or more error messages.

See the [Schema definition languages](#) clause for references to the schema definitions of the error response payload.

The `@Message.ExtendedInfo` property should be present in all error responses. If the `@Message.ExtendedInfo` property is present, all information necessary to process the error should be provided in the `@Message.ExtendedInfo` property. Clients should look for the `@Message.ExtendedInfo` property for error processing first, and fallback on the `code` and `message` properties if `@Message.ExtendedInfo` is not present.

The following sample error response contains two messages in the `@Message.ExtendedInfo` property that describe two different errors. The message described by the `code` and `message` properties do not provide actionable information for the client.

```
{
  "error": {
    "code": "Base.1.8.GeneralError",
    "message": "A general error has occurred. See Resolution for information on how to resolve the error.",
    "@Message.ExtendedInfo": [
      {
        "@odata.type": "#Message.v1_1_1.Message",
        "MessageId": "Base.1.8.PropertyValueNotInList",
        "RelatedProperties": [
          "/IndicatorLED"
        ],
        "Message": "The value Red for the property IndicatorLED is not in the list of acceptable values.",
        "MessageArgs": ["Red",
                      "IndicatorLED"
                    ],
        "Severity": "Warning",
        "MessageSeverity": "Warning",
        "Resolution": "Choose a value from the enumeration list that the implementation can support and resubmit the request if the operation failed."
      },
      {
        "@odata.type": "#Message.v1_1_1.Message",
        "MessageId": "Base.1.8.PropertyNotWritable",
        "RelatedProperties": [
          "/SKU"
        ],
        "Message": "The property SKU is a read only property and cannot be assigned a value.",
        "MessageArgs": ["SKU"],
        "Severity": "Warning",
        "MessageSeverity": "Warning",
        "Resolution": "Remove the property from the request body and resubmit the request if the operation failed."
      }
    ]
  }
}
```

## 9 Data model

One of the key tenets of Redfish is the separation of protocol from the data model. This separation makes the data both transport and protocol agnostic. By concentrating on the data transported in the payload of the protocol (in HTTP, it is the HTTP body), Redfish can also define the payload in any encoding and the data model is intended to be schema-language agnostic. While Redfish uses the JSON data-interchange format, Redfish provides a common encoding type that ensures property naming conventions that make development easier in JavaScript, Python, and other languages. This encoding type helps the Redfish data model be more easily accessible in modern tools and programming environments.

The data model allows an OEM to extend the model by adding an [OEM resource](#) or [extending a resource](#).

This clause describes common data model, resource, and Redfish schema requirements.

### 9.1 Resources

A [resource](#) is a single entity accessed at a specific URI. Services use the `application/json` MIME type to return resources as JSON payloads.

Each resource shall be strongly typed, defined by a [resource type](#) in a [Redfish schema document](#), and identified in the response payload by the value of the [type identifier](#) property.

Responses for a single resource shall contain the following properties:

- `@odata.id`
  - [Registry resources](#) are not required to provide `@odata.id`
- `@odata.type`
- `Id`
- `Name`

Responses may also contain other properties defined within that [resource type](#). Responses shall not include any properties not defined by that resource type.

### 9.2 Resource types

A [resource type](#) defines the set of properties that may be returned in the response payload of a Redfish resource request. Each resource type is documented in a [Redfish schema document](#), and those documents are known collectively as the [Redfish schema](#). The resource type may also include definitions for [actions](#) available for that resource.

Resource types are named to match the contents and purpose of the resource that they define. For example, the `Circuit` resource type defines the properties and actions related to a single electrical circuit. Resource types provide global uniqueness for definitions across multiple schema files and allow for schema files to reference each other. Resource types may be defined by OEMs to extend the Redfish schema, and should follow the naming rules specified by the [OEM resource types](#) clause.

## 9.3 Resource collections

---

A *resource collection* is a set of resources that share the same schema definition. Services use the `application/json` MIME type to return resource collections as JSON payloads.

Resource collection responses shall contain the following properties:

- `@odata.id`
- `@odata.type`
- `Name`
- `Members`
- `Members@odata.count`

Responses for resource collections may contain the following properties:

- `@odata.context`
- `@odata.etag`
- `Description`
- `Members@odata.nextLink`
- `Oem`

Responses for resource collections shall not contain any other properties with the exception of [payload annotations](#).

## 9.4 OEM resources

---

OEMs and other third parties can extend the Redfish data model by creating additional resource types. Extending the data model is accomplished by defining an OEM resource type, and schema file, for each resource type, and creating hyperlinks to connect instances of new resources to the *resource tree*.

Companies, OEMs, and other organizations may also use the `Oem` property in resources, the [links property](#), and the [actions property](#) to define additional [properties](#), hyperlinks, and [actions](#) for standard Redfish resource types.

While the information and semantics of these extensions are outside of the standard, the schema representing the data, the resource itself, and the semantics around the protocol shall conform to the requirements in this

specification. OEMs are encouraged to follow the design tenets and naming conventions in this specification when defining OEM resources or properties.

## 9.5 Common data types

### 9.5.1 Primitive types

[Table 17](#) describes the primitive data types for properties and action parameters in the data model:

**Table 17 — Primitive data types**

Type	Description
Boolean	A variable with a value of <code>true</code> or <code>false</code> .
Number	A number with optional decimal point or exponent. Number properties may restrict the representation to an integer or a number with decimal point.
String	A sequence of characters enclosed with double quotes ( " ).
Array	A comma-separated set of the previous types enclosed with square braces ( [ and ] ). See the <a href="#">Array properties clause</a> .
Object	A set of properties enclosed with curly braces ( { and } ). See the <a href="#">Structured properties clause</a> .
Null	<code>null</code> value, which the service uses when it is unable to determine the property's value due to an error or other temporary condition, or if the schema has requirements for using <code>null</code> for other special conditions.

When receiving values from the client, services should support other valid representations of the data in the specified JSON type. In particular, services should support valid integer and decimal values in exponential notation and integer values that contain a decimal point with no non-zero trailing digits.

### 9.5.2 Enumerations

Enumerations are frequently used in Redfish to promote readability and interoperability, especially compared to the use of string values when used for similar purposes. Enumerations aren't optimal in all cases. Properties with two values that are likely to not have additional values should consider the boolean type if the `true` and `false` values can be described by the property name. The following design tenets apply to enumerations:

- Enumeration values can be added to existing properties. Client software should be prepared to receive enumeration values that are not known if the resource schema version is higher than the client's supported version.
- Enumeration properties should avoid definition of "unknown", "other", or similar generic or placeholder values as these reduce interoperability.
- Feedback is encouraged for adding enumeration values to existing properties to cover new technologies or use cases.

- Enumeration values are generally defined to support existing or newly developed products.
- Enumeration values that are obsolete or highly unlikely to appear in implementations are not included, but they can be added.
- Enumerations may include vendor-specific values when they apply to multiple products or implementations.
- Sometimes the value `OEM` is included as an enumeration value. When this is in the enumeration, client software should be aware that there is likely an `Oem` property with additional information. In some cases, standard schema contains a standard value to further describe this enumeration value when additional OEM data is unlikely.

### 9.5.3 Empty string values

String properties should return an empty string ("") for properties configured by a user or external service that have not been set to an initial value. This allows client software to identify the property as supported by the service and avoids the use of `null`, which indicates an error condition. For example, the `AssetTag` property must be set by the end user, and therefore would return an empty string ("") until assigned a value by the user, while a failure to read the stored `AssetTag` value due to a non-volatile memory error would return `null`. To improve interoperability, implementations should avoid the use of filler strings, such as `N/A` or `<Empty>`, to represent a value not set by a user.

### 9.5.4 GUID and UUID values

Globally Unique Identifier (GUID) and Universally Unique Identifier (UUID) values are unique identifier strings and shall use the [RFC4122](#)-defined format:

```
([0-9a-fA-F]{8}-[0-9a-fA-F]{4}-[0-9a-fA-F]{4}-[0-9a-fA-F]{4}-[0-9a-fA-F]{12})
```

### 9.5.5 Date-Time values

Date-Time values are strings according to the ISO 8601 extended format, including the time offset or UTC suffix.

Date-Time values shall use the format:

```
<YYYY>-<MM>-<DD>T<hh>:<mm>:<ss>[.<SSS>](Z|((+|-)<HH>:<MM>))
```

where

- `<YYYY>` is the four-digit year.
- `<MM>` is the two-digit month (1 to 12).
- `<DD>` is the two-digit day (1 to 31).
- `T` is the time separator. Shall be a capital `T`.
- `<hh>` is the two-digit hour (0 to 23).
- `<mm>` is the two-digit minute (0 to 59).

- `<ss>` is the two-digit second (0 to 59).
- `<sss>` is optional and is the decimal fraction of a second. Shall be one or more digits where the number of digits implies the precision.
- `z` is the zero offset indicator. Shall be a capital `z`.
- `<HH>` is the two-digit hour offset (0 to 23).
- `<MM>` is the two-digit minute offset (0 to 59).

For example, `2015-03-13T04:14:33+06:00` represents March 13, 2015 at 4:14:33 with a +06:00 time offset.

When the time of day is unknown or serves no purpose, the service shall report `00:00:00z` for the time of day value.

### 9.5.6 Duration values

Duration values are strings according to the ISO 8601 duration format, with the exception of not expressing a representation for years, months, weeks, or negative values. Duration values shall use the format:

`P[<d>D][T[<h>H][<m>M][<s>[.<f>]S]]`

where

- `<d>` is the number of days.
- `<h>` is the number of hours.
- `<m>` is the number of minutes.
- `<s>` is the number of seconds.
- `<f>` is the fractional seconds.

Each field is optional and can contain more than one digit.

For example, [Table 18](#) describes the following durations:

**Table 18 — Durations**

Value	Duration
<code>P90D</code>	Ninety days.
<code>P3D</code>	Three days.
<code>PT6H</code>	Six hours.
<code>PT10S</code>	Ten seconds.
<code>PT0.001S</code>	0.001 seconds.
<code>PT1H30M</code>	One hour and 30 minutes.

**DEPRECATED:** Duration values shall use the format: `P[<y>Y][<m>M][<w>W][<d>D][T[<h>H][<m>M][<s>[.<f>]S]]`. This definition allows for specifying years, months, and weeks. ISO 8601 does not specify an exact value for the duration of a year or of a month, which introduces interoperability challenges.

### 9.5.7 Reference properties

Reference properties provide a reference to another resource in the data model. Reference properties are JSON objects that contain an `@odata.id` property. The `@odata.id` property value is the URI of the referenced resource.

### 9.5.8 Non-resource reference properties

Non-resource reference properties provide a URI to services or documents that are not Redfish-defined resources. These properties shall include the `Uri` or `URI` term in their property name and shall be of type string. For example, `AssemblyBinaryDataUri` in the Assembly schema. The access protocol, request headers, response headers, and data format of the referenced URI may be defined in schema for that property. Non-resource reference properties that refer to local HTTP/S targets shall follow the Redfish protocol, including use of Redfish sessions and access control, unless otherwise specified by the property definition in schema.

### 9.5.9 Array properties

Array properties contain a set of values or objects, and appear as JSON arrays within a response body. Array elements shall all contain values of the same data type.

Table 19 describes the array types, regardless of the data type of the elements:

**Table 19 — Array types**

Array type	Description
Fixed length	Contains a static number of elements. The property definition sets or the implementation chooses the size of the array.
Variable length	Contains a variable number of elements. The array size is not specified and the size varies among instances. The array size may change. This array style is the most common style.
Rigid	<p>The array index is meaningful. When elements are added to or removed from the array, the elements do not change their position, or index, in the array. An element that is removed from a rigid array shall be replaced by a <code>null</code> element and all other elements shall remain at their current index.</p> <p>Empty elements in a rigid array property shall be represented by <code>null</code> elements. Any array property that uses this style shall indicate the rigid style in the <a href="#">long description</a> of its schema definition.</p>

Services may pad an array property with `null` elements at the end of the sequence to indicate the array size to clients. This practice is useful for small fixed-length arrays, and for variable or rigid arrays with a restrictive maximum size. Services should not pad array properties if the maximum array size is not restrictive. For example, an array

property typically populated with two elements, that a service limits to a maximum of 16 elements, should not pad the array with 14 `null` elements.

### 9.5.10 Structured properties

Structured properties are JSON objects within a response body.

Some structured properties inherit from the `Resource.v1_0_0.ReferenceableMember` definition. Structured properties that follow this definition shall contain the `MemberId` and `resource identifier` properties.

Because the definition of structured properties can evolve over time, clients need to be aware of the inheritance model that the different structured property definitions use.

For example, the `Location` property definition in the `Resource` schema has gone through several iterations since the `Resource.v1_1_0` type was introduced, and each iteration inherits from the earlier version so that existing references in other schemas can leverage the additions.

Structured property references need to be resolved for both local and external references.

A local reference is a resource that has a structured property in its own schema, such as `ProcessorSummary` in the `ComputerSystem` resource. In these cases, the `type` property for the resource is the starting point for resolving the structured property definition.

To find the latest applicable version, clients can step the [version of the resource](#) backwards.

For example, if a service returns `#ComputerSystem.v1_4_0.ComputerSystem` as the resource type, a client can step backwards from `ComputerSystem.v1_4_0`, to `ComputerSystem.v1_3_0`, to `ComputerSystem.v1_2_0`, and so on, until it finds the `ProcessorSummary` structured property definition.

An external reference is a resource that has a property that references a definition found in a different schema, such as the `Location` property in the `Chassis` resource.

In these cases, clients can use the latest version of the external schema file as a starting point to resolve the structured property definition.

For example, if the latest version of the `Resource` schema is `1.6.0`, a client can go backward from `Resource.v1_6_0`, to `Resource.v1_5_0`, to `Resource.v1_4_0`, and so on, until it finds the `Location` structured property definition.

### 9.5.11 Message object

#### 9.5.11.1 Overview

A message object provides additional information about an [object](#), [property](#), or [error response](#).

[Table 20](#) describes the properties of the message object, which is a JSON object:

**Table 20 — Message object properties**

Property	Type	Required	Defines
MessageId	String	Yes	Error or message. Do not confuse this value with the HTTP status code. Clients can use this code to access a detailed message from a message registry.
Message	String	No	Human-readable error message that indicates the semantics associated with the error. This shall be the complete message, and not rely on substitution variables.
RelatedProperties	An array of JSON pointers	No	Properties in a JSON payload that the message describes.
MessageArgs	An array of strings	No	Substitution parameter values for the message. If the parameterized message defines a <code>MessageId</code> , the service shall include the <code>MessageArgs</code> in the response.
MessageSeverity	String (enumeration)	No	Severity of the error. Services can replace the value of the <code>MessageSeverity</code> property defined in the message registry with a value more applicable to the implementation.
Severity	String	No	Severity of the error. Services can replace the value of the <code>Severity</code> property defined in the message registry with a value more applicable to the implementation.  <b>DEPRECATED:</b> This property has been deprecated in favor of <code>MessageSeverity</code> .
Resolution	String	No	Recommended actions to take to resolve the error. Services can replace the value of the <code>Resolution</code> property defined in the message registry with a service-defined resolution.

Each instance of a message object shall contain at least a `MessageId`, together with any applicable `MessageArgs`, or a `Message` property that defines the complete human-readable error message.

A `MessageId` identifies a specific message that a [message registry](#) defines.

### 9.5.11.2 MessageId format

The `MessageId` property value shall be in the format:

```
<MessageRegistryPrefix>.<MajorVersion>.<MinorVersion>.<MessageKey>
```

where

- `<MessageRegistryPrefix>` is the name of the message registry. The message registry name shall be Pascal-

cased, except for any prepended unique OEM identifier which may include underscore ( `_` ) characters. The message registry name shall be exposed in the `RegistryPrefix` property in the message registry.

- `<MajorVersion>` is a non-negative integer that represents the major version of the message registry.
- `<MinorVersion>` is a non-negative integer that represents the minor version of the message registry.
- `<MessageKey>` is a human-readable key into the message registry. The message key shall be Pascal-cased and shall not include spaces, periods, or special characters.

To search the message registry for a message, the client can use the `MessageId`.

The message registry approach has advantages for internationalization because the message registry can be translated easily and is lightweight for implementations because large strings need not be included with the implementation.

The use of `GeneralError` from the Base Message Registry as a `MessageId` in `ExtendedInfo` is discouraged. If no better message exists or the `ExtendedInfo` array contains multiple messages, use `GeneralError` from the Base Message Registry only in the `code` property of the `error` object.

When an implementation uses `GeneralError` from the Base Message Registry in `ExtendedInfo`, the implementation should include a service-defined value for the `Resolution` property with this error to indicate how to resolve the problem.

## 9.6 Properties

---

### 9.6.1 Properties overview

Every property included in a Redfish response payload shall be defined in the schema for that [resource](#). The following attributes apply to all property definitions:

- Property names in the request and response payload shall match the casing of the `Name` attribute value in the defining schema.
- Required properties shall always be returned in a response.
- Properties not returned from a `GET` operation indicate that the property is not supported by the implementation, or by that particular resource instance. Differences in underlying product support or configuration varies among resource instances, and therefore the properties returned by each instance vary accordingly.
- If an implementation supports a property, it shall always provide a value for that property. If a value is unknown at the time of the operation due to an internal error, or inaccessibility of the data, the value of `null` is an acceptable value if supported by the schema definition.
- Resource instances should omit properties if the underlying product, service, or current configuration does not provide the function described by the property. For example, a chassis resource instance might not provide a serial number, and therefore should omit the `SerialNumber` property, while other chassis resource instances that have a serial number provide this property. See the [Special resource situations](#) clause for handling special resource situations.

- A service may implement a writable property as read-only.
- All property definitions in schema shall specify a well-known data type, with exceptions in the following standard schemas:
  - `AttributeRegistry` and `Bios` : To support arbitrary BIOS settings of different data types.

This clause also contains a set of common properties across all Redfish resources. The property names in this clause shall not be used for any other purpose.

### 9.6.2 Resource identifier (`@odata.id`) property

Registry resources in a response may include an `@odata.id` property. All other resources and resource collections in a response shall include an `@odata.id` property. The value of the identifier property shall be the resource [URI](#).

### 9.6.3 Resource type (`@odata.type`) property

All resources and resource collections in a response shall include an `@odata.type` type property. To support generic OData clients, all structured properties in a response should include an `@odata.type` type property.

The value of the type property for resources and structured properties shall be in the format:

```
#<ResourceType>.<Version>.<TermName>
```

where

- `<ResourceType>` is the resource type in the Redfish schema that defines the resource.
- `<Version>` is the resource type version, in the format: `v<MajorVersion>_<MinorVersion>_<ErrataVersion>`.
- `<TermName>` is the specific type defined within the resource type definition. For most Redfish resources, the specific type name is the same as the resource type name.

An example of a resource type value is `#ComputerSystem.v1_0_0.ComputerSystem`, where `ComputerSystem.v1_0_0` denotes the version 1.0.0 of the `ComputerSystem` resource type, and the specific type is `ComputerSystem`.

The value of the type property for resource collections shall be in the format:

```
#<ResourceType>.<ResourceType>
```

where

- `<ResourceType>` is the resource type in the Redfish schema that defines the resource collection.

An example of a resource collection type value is `#ComputerSystemCollection.ComputerSystemCollection` for the `ComputerSystemCollection` resource collection.

## 9.6.4 Resource ETag (@odata.etag) property

ETags enable clients to conditionally retrieve or update a [resource](#). Resources should include an `@odata.etag` property. For a resource, the value shall be the [ETag](#).

## 9.6.5 Resource context (@odata.context) property

Responses for [resources](#) and [resource collections](#) may contain an `@odata.context` property that describes the source of the payload.

If the `@odata.context` property is present, it shall be the context URL that describes the resource, according to [OData Protocol](#).

The context URL for a resource should be in the format:

```
/redfish/v1/$metadata#<ResourceType>.<ResourceType>
```

where

- `<ResourceType>` is the resource type of the resource or resource collection.

For example, the following context URL specifies that the results show a single `ComputerSystem` resource:

```
{  
  "@odata.context": "/redfish/v1/$metadata#ComputerSystem.ComputerSystem",  
  ...  
}
```

The context URL for a resource may be in one of the other formats that [OData Protocol](#) specifies.

## 9.6.6 Id

The `Id` property of a [resource](#) uniquely identifies the resource within the resource collection that contains it. The value of `Id` shall be a string that is unique across a [resource collection](#). Since [URIs are constructed from the value of the Id property](#), the value shall not contain any [RFC1738](#)-defined unsafe characters.

## 9.6.7 Name

The `Name` property conveys a human-readable moniker for a [resource](#). The data type of the `Name` property shall be string. The value of `Name` is NOT required to be unique across resource instances within a [resource collection](#).

## 9.6.8 Description

The `Description` property conveys a human-readable description of the `resource`. The data type of the `Description` property shall be string.

## 9.6.9 MemberId

The `MemberId` property uniquely identifies an element within an array, where a `reference property` can reference the element. The value of `MemberId` shall be a string that is unique across the array.

## 9.6.10 Count (`Members@odata.count`) property

The count property defines the total number of `resource`, or `members`, that are available in a `resource collection`. The count property shall be named `Members@odata.count` and its value shall be the total number of members available in the resource collection. The `$top` or `$skip` `query parameters` shall not affect this count. If the number of members available in the resource collection is reduced due to filtering, such as in response to the `$filter` query parameter, the count should be the total number of members available in the resource collection after the filter is applied.

## 9.6.11 Members

The `Members` property of a `resource collection` identifies the `members` of the collection. The `Members` property is required and shall be returned in the response for any resource collection. The `Members` property shall be an array of JSON objects named `Members`. The `Members` property shall not be `null`. Empty collections shall be an empty JSON array.

## 9.6.12 Next link (`Members@odata.nextLink`) property

The next link (`Members@odata.nextLink`) property value shall be an opaque URL to a resource, with the same `@odata.type`, which contains the next set of partial `members` from the original operation. The next link property shall only be present if the number of members in the resource collection is greater than the number of members returned, and if the payload does not represent the end of the requested resource collection.

The next link property may contain the `$skiptoken` query parameter. The `$skiptoken` query parameter shall contain an opaque value that allows a client to reference the next page of the collection. Clients shall not interpret or construct `$skiptoken` values.

The next link property value should contain control information provided by the client in `query parameters` that affect the size of the collection and properties returned in the response, such as the `$filter` and `$expand` query parameters. For example, if a client performs a request to `/redfish/v1/Chassis/1/Sensors?$filter=Status/Health eq 'OK'&$expand=..`, some appropriate values include:

- `/redfish/v1/Chassis/1/Sensors?$filter=Status/Health eq 'OK'&$expand=.&$skip=50`
- `/redfish/v1/Chassis/1/Sensors?$filter=Status/Health eq 'OK'&$expand=.&$skiptoken=1`
- `/redfish/v1/Chassis/1/Sensors?$skiptoken=ffb559f3-7ee9-42d6-ac8b-888c0ad24db3`
  - In this case, the skip token internally maintains the client's filter and expand request.

The `Members@odata.count` property value is the total number of resources available if the client enumerates all pages of the resource collection.

## 9.6.13 Links

The `Links` property represents the hyperlinks associated with the [resource](#), as defined by that resource's schema definition. All associated [reference properties](#) defined for a resource shall be nested under the `Links` property. All directly ([subordinate](#)) referenced properties defined for a resource shall be in the root of the resource. There are some exceptions to these rules, such as where ease of [expansion](#) or [deep operations](#) is beneficial to the user.

The `Links` property shall be named `Links` and contain a property for each related resource.

To navigate vendor-specific hyperlinks, the `Links` property shall also include an `Oem` property.

### 9.6.13.1 Reference to a related resource

A reference to a single [resource](#) is a JSON object that contains a single [resource identifier property](#). The name of this reference is the name of the relationship. The value of this reference is the URI of the referenced resource.

```
{
  "Links": {
    "ManagedBy": {
      "@odata.id": "/redfish/v1/Chassis/Encl1"
    }
  }
}
```

### 9.6.13.2 References to multiple related resources

A reference to a set of zero or more related [resources](#) is an array of JSON objects. The name of this reference is the name of the relationship. Each element of the array is a JSON object that contains a [resource identifier property](#) with the value of the URI of the referenced resource.

```
{
  "Links": {
    "Contains": [
      "@odata.id": "/redfish/v1/Chassis/1"
    ]
  }
}
```

```

    },
    {
        "@odata.id": "/redfish/v1/Chassis/Encl1"
    }
}
}

```

## 9.6.14 Actions property

The `Actions` property contains the `actions` supported by a [resource](#).

### 9.6.14.1 Action representation

Each supported action is represented as a property nested under `Actions`. The unique name that identifies the action is used to construct the property name.

This property name shall be in the format:

```
#<ResourceType>.<ActionName>
```

where

- `<ResourceType>` is the resource where the action is defined.
- `<ActionName>` is the name of the action.

The client may use this fragment to identify the action definition in the [referenced](#) Redfish schema document.

The property for the action is a JSON object and contains the following properties:

- The `target` property shall be present and defines the relative or absolute URL to invoke the action.
- The `title` property may be present and defines the action's name.

The [OData JSON Format](#) Specification defines the `target` and `title` properties.

To specify the list of supported values for a parameter, the service may include the `@Redfish.AllowableValues` annotation.

For example, the following property defines the `Reset` action for a `ComputerSystem`:

```
{
    "#ComputerSystem.Reset": {
        "target": "/redfish/v1/Systems/1/Actions/ComputerSystem.Reset",
        "title": "Computer System Reset",
    }
}
```

```
  "ResetType@Redfish.AllowableValues": ["On", "ForceOff", "GracefulRestart",
    "GracefulShutdown", "ForceRestart", "Nmi", "ForceOn",
    "PushPowerButton"]
},
...
}
```

Given this, the client could invoke a `POST` request to `/redfish/v1/Systems/1/Actions/ComputerSystem.Reset` with the following body:

```
POST /redfish/v1/Systems/1/Actions/ComputerSystem.Reset HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
  "ResetType": "On"
}
```

The resource may provide a separate `@Redfish.ActionInfo` resource to describe the parameters and values that a particular instance or implementation supports. Use the `@Redfish.ActionInfo` annotation to specify the `ActionInfo` resource, which contains a URI to the `@Redfish.ActionInfo` resource for the action. For details, see the [Action info annotation](#) clause.

#### 9.6.14.2 Action responses

Response payloads for actions may contain a JSON body that is described by the schema definition for the action. See the [Schema definition languages](#) clause for the representation of these definitions. Actions that do not define a response body may provide an [error response](#) in the response payload.

Clients can discover the schema definition of the action response based on the [property name of the action](#) found in `GET` responses for resources containing the action. For example, if an action is named `#ContosoComputerSystem.Reset`, a client can find the action response definition, if there is one available, by locating the `Reset` action definition found in the `ContosoComputerSystem` schema.

#### 9.6.15 Oem

The `Oem` property is used for [extending standard resources](#) with OEM extensions.

### 9.6.16 Status

The `Status` property represents the status of a [resource](#). The `Status` property shall follow the definition for `Status` in the `Resource` schema.

By having a common representation of status, clients can depend on consistent semantics. The `status` property is capable of indicating the current state, health of the resource, and the health of subordinate resources.

## 9.7 Naming conventions

---

The Redfish interface is intended to be easily readable and intuitive. Thus, consistency helps the consumer understand the use of a newly discovered property. While consistency is no substitute for the normative information in the *Redfish Specification* and Redfish schema, the following [naming rules](#) help with readability and client usage. In general, names in Redfish are designed and intended to be human-readable and convey the meaning of the name, in context, without the need to consult schema definitions or other documentation.

### 9.7.1 Naming rules

Standard Redfish schema and registries defined and published by DMTF as well as those created by others and republished by DMTF shall follow a set of naming conventions. These conventions are intended to ensure consistent naming and eliminate naming collisions. For schema files, the [resource type](#) is used to construct the [type property](#) and the schema file name.

Standard Redfish properties follow similar naming conventions, and should use a common definition when defined in multiple schemas across the Redfish data model. This consistency enables code re-use across resources and increases interoperability. New resource definitions should leverage existing property definitions whenever possible.

The general Redfish naming rules for resource types, registries, properties, and action parameters are as follows:

- Names shall be Pascal-cased. The first letter of each word in a name shall be uppercase and spaces between words shall be removed. For example, `ComputerSystem`, `PowerState`, and `SerialNumber`.
- Names shall differ from other defined names by more than case sensitivity. For example, `Snmpv1` cannot be used as a name if `SNMPv1` is also defined.
- Names of array properties or reference properties for resource collections should use a plural form of the name. All other names should use the singular form of the name.
- Reference properties for resource collections should omit the term "collection" in the name.
- Names shall start with an alphabetic character ( `a` to `z` and `A` to `Z` ) and consist only of alphanumeric characters ( `a` to `z`, `A` to `Z`, and `0` to `9` ).
- Both characters should be capitalized for two-character acronyms. For example, `IPAddress` OR `RemoteIP`.
- Names constructed from a single acronym or mixed-case name, such as `LDAP`, `PCIe`, or `SNMP`, should use the typical capitalization for that name.

- Names incorporating acronyms with three or more characters should follow the capitalization used in related names for consistency. For example, `EnableSNMPv1` and `EnableSNMPv2` follow the pattern used for `SNMP`.
- Pascal-casing may be used for acronyms longer than two characters to improve readability, especially when two or more acronyms appear together in a name, which should be avoided.
- The names `Settings` and `SD` are reserved for use for `settings resources` and shall not be used for schema names.
- Names that represent a string-typed HTTP URI, and is not a `reference property`, should end in `URI`. For example, `TargetComponentURI`.
- The name `ReturnType` is reserved for RDE usage and shall not be used for action parameter names.

Exceptions are allowed for the following cases:

- Well-known technology abbreviations, acronyms, or product names should follow their defined capitalization. Examples include `iSCSI`, `iSCSITarget`, and `iLO`.
- OEM appears as `oem` in schema and property names either alone or as a portion of a name, but should be `OEM` when used alone as an enumeration value.
- Underscore characters are allowed in the construction of OEM-specified object property names when required, and in OEM-defined resource types or OEM-defined registry names.

`Enumeration values` follow the above rules and exceptions. Additionally, enumeration values:

- Should start with a letter and be followed by letters or numbers to conform to schema description language requirements. Underscore characters may be used to replace other special characters, or to significantly improve readability, but this usage is discouraged.
- Should prioritize readability as they may appear unmodified on user interfaces, whereas property or schema names should follow conventions and strive for consistency.

For properties that have units or other special meaning, append a unit identifier to the name. Examples include:

- Bandwidth (Mbps). For example, `PortSpeedMbps`.
- CPU speed (MHz). For example, `ProcessorSpeedMHz`.
- Memory size (MB). For example, `MemoryMB`.
- Counts of items (Count). For example, `ProcessorCount` OR `FanCount`.
- The state of a resource (State). For example, `PowerState`.
- State values where work is in process. For example, `Applying` OR `ClearingLogic`.

## 9.7.2 URI naming rules

The following rules apply to Redfish schema-defined URLs:

- URI segments should generally follow the naming rules, and for each segment, follow the name of the property that provides the hyperlink.

- URI segments for resource collections should use the plural form of the resource collection schema name, with the `Collection` term omitted. For example, `Processors` for a `ProcessorCollection`.
- URI segments for resource collections shall not be named `Members`, as this value will conflict with POST operations on the required `Members` property. See the [POST \(create\)](#) clause for more information.
- If a hyperlink to a subordinate resource is not found at the root of the resource, the URI segments should contain the property path. For example, for the `Certificates` hyperlink found in `ManagerNetworkProtocol` within the `HTTPS` object, `HTTPS` should be one of the URI segments resulting in the URI pattern `/redfish/v1/Managers/{ManagerId}/NetworkProtocol/HTTPS/Certificates`.

Starting with *Redfish Specification v1.18.0*, in cases where the URI segment does not follow one or more of the previous rules, the schema definition for the [reference property](#) for the URI shall specify the [URI segment annotation](#) to define the URI segment to append.

## 9.8 Extending standard resources

---

### 9.8.1 Extending standard resources overview

In the context of this clause, the OEM term refers to any company, manufacturer, or organization that provides or defines an extension to DMTF-published schema and functionality for Redfish. All Redfish-specified [resources](#) include an empty structured `oem` property. The value of this predefined placeholder can encapsulate one or more OEM-specified object properties, which can contain OEM-specific property definitions.

### 9.8.2 OEM property format and content

Each property contained within the `oem` property shall be an OEM-specified JSON object. The name of each object property shall uniquely identify the OEM or organization that defines the properties contained by that object. The [OEM-specified object naming](#) clause describes this naming convention.

The OEM-specified object shall include a [type property](#) if the object:

- Is not contained in an array of objects.
- Is contained in the first object within an array of objects.
- Is contained in subsequent objects within an array of objects, whose type is different from the first array member.

To avoid naming collisions with other OEM definitions, the type property should start with the defining organization's unique OEM identifier, including possible subdivisioning. Separator underscores (`_`) may be excluded from the type property and schema file name for improved readability. The [Schema definition languages](#) clause contains file name construction rules, such as handling version information. See [OEM property examples](#) for examples.

The `oem` property can simultaneously hold multiple OEM-specified objects, including objects for more than one company or organization.

The definition of any other properties that are contained within the OEM-specified object, along with the functional specifications, validation, or other requirements for that content is OEM-specific and outside the scope of this specification. While there are no Redfish-specified limits on the size or complexity of the elements within an OEM-specified object, it is intended it is typically used for only a small number of simple properties that augment the Redfish [resource](#). If a large number of objects or a large quantity of data compared to the size of the Redfish resource is to be supported, the OEM should consider creating a [subordinate resource](#) for their extensions.

### 9.8.3 OEM-specified object naming

The OEM-specified object properties within the `Oem` property are named by using a unique OEM identifier. There are two specified forms for the identifier. The identifier shall be either an ICANN-recognized domain name (including the top-level domain suffix), with all dot (.) separators replaced with underscores (\_), or an IANA-assigned Enterprise Number prefixed with "EID\_".

**DEPRECATED:** The identifier shall be either an ICANN-recognized domain name including the top-level domain suffix, or an IANA-assigned Enterprise Number prefixed with `EID_`.

Organizations that use `.com` domain names may omit the `.com` suffix. For example, Contoso.com would use `Contoso` instead of `Contoso_com`, but Contoso.org would use `Contoso_org`. The domain name portion of an OEM identifier shall be considered to be case-insensitive. That is, the text `Contoso_biz`, `contoso_BIZ`, `conToso_biz`, and so on all identify the same OEM.

The OEM identifier portion of the object name may be followed by an underscore (\_) and any additional string to enable further subdivisions of OEM-specified objects as desired. For example, `Contoso_xxxx` or `EID_412_xxxx`. The form and meaning of any text that follows the trailing underscore is completely OEM-specific. OEM-specified extension suffixes may be case-sensitive, depending on the OEM. Generic client software should treat such extensions, if present, as opaque and not try to parse nor interpret the content.

There are organizations with which DMTF has a working relationship and have registered their OEM namespace directly in the specification to allow extensions of the ICANN domain name requirements above. The following organization OEM namespaces shall be considered reserved:

- OpenBMC
- OpenCompute

This suffix could be used in many ways, depending on OEM need. For example, the Contoso company may have a *Research* sub-organization, in which case the OEM-specified property name might be extended to `Contoso_Research`. Alternatively, it can identify a unique resource type for a functional area, geography, subsidiary, and so on.

The OEM identifier portion of the name typically identifies the company or organization that created and maintains the schema for the property. However, this practice is not a requirement. The identifier is only required to uniquely identify the party that is the top-level manager of a resource type to prevent collisions between OEM property definitions from different vendors or organizations. Consequently, the organization for the top of the resource type

may be different from the organization that provides the definition of the OEM-specified property. For example, Contoso may allow one of their customers, such as `CustomerA`, to extend a Contoso product with certain CustomerA proprietary properties. In this case, although Contoso allocated the name `Contoso_CustomerA`, it could be CustomerA that defines the content and functionality within that resource type. In all cases, OEM identifiers should not be used except with permission or as specified by the identified company or organization.

#### 9.8.4 OEM resource types

Companies, OEMs, and other organizations can define additional [resources](#) and link to them from an `Oem` property in a standard Redfish resource, preferably from the `Oem` property within the `Links` property. To avoid naming collisions with current or future standard Redfish resource types or schema files, the defining organization's unique OEM identifier, including possible subdivisioning, should be prepended to the OEM resource type name with an optional underscore ( `_` ) as separator. This unique OEM identifier should follow the same naming as defined in the [OEM-specified object naming](#) clause. The name of the OEM resource type, including the unique OEM identifier, should also be prepended to the file name of OEM schema file that specifies the OEM resource type. Separator underscores ( `_` ) may be excluded from the OEM resource type name and schema file name for improved readability. The [Schema definition languages](#) clause contains other file name construction rules, such as handling version information.

For example, OEM resource type `ContosoDrive` OR `Contoso_CustomerA_Drive` would not conflict with the standard Redfish `Drive` resource type, or conflict with another OEM's drive-related definition.

#### 9.8.5 OEM registries

Companies, OEMs, and other organizations can define additional [registries](#) and utilize them in [message objects](#), privileges or for BIOS attributes. To avoid naming collisions with current or future standard Redfish message registries, the defining organization's unique OEM identifier, including possible subdivisioning, should be prepended to the registry name with an optional underscore ( `_` ) as separator. This unique OEM identifier should follow the same naming as defined in the [OEM-specified object naming](#) clause. Separator underscores ( `_` ) may be excluded from the OEM registry name for improved readability. The OEM registry name, including the unique OEM identifier, should also be used to construct the registry file name as defined in the [Registry file naming](#) clause.

For example, OEM registry name `ContosoDriveEvent` OR `Contoso_CustomerB_DriveEvent` would not conflict with a possible future standard Redfish `driveEvent` message registry name, or conflict with another OEM's drive-related registry name.

#### 9.8.6 OEM URIs

To avoid URI collisions with other OEM resources and future Redfish standard resources, the URIs for OEM resources within the Redfish [resource tree](#) shall be in the form:

```
<BaseUri>/Oem/<OemIdentifier>/<ResourcePath>
```

where

- `<BaseUri>` is the URI segment of the standard Redfish resource starting with `/redfish/` where the `Oem` property is used. For example, `/redfish/v1/Systems/3AZ38944T523`.
- `<OemIdentifier>` is the unique identifier of the OEM, including possible subdivisioning, that follows the same naming as defined in the [OEM-specified object naming](#) clause. Separator underscores (`_`) may be excluded for improved readability.
- `<ResourcePath>` is the path to the OEM-defined resource. This path might contain multiple segments for cases where OEM-defined resources are subordinate to an OEM-defined resource. Each segment in the path contains the name of an OEM-defined resource.

For example, if Contoso defined a new `ContosoAccountServiceMetrics` OEM resource type to be linked through the `Oem` property at the `/redfish/v1/AccountService` URI, the OEM resource has the `/redfish/v1/AccountService/Oem/Contoso/AccountServiceMetrics` URI. If Contoso uses a subdivision of their OEM identifier such as `Contoso_CustomerA` the OEM resource has the URI `/redfish/v1/AccountService/Oem/Contoso_CustomerA/AccountServiceMetrics` OR `/redfish/v1/AccountService/Oem/ContosoCustomerA/AccountServiceMetrics`.

### 9.8.7 OEM property examples

The following fragment shows examples of naming and the `Oem` property as it might appear when accessing a [resource](#). The example shows that the OEM identifiers can be of different forms, that OEM-specified content can be simple or complex, and that the format and usage of extensions of the OEM identifier is OEM-specific.

```
{  
    "Oem": {  
        "Contoso": {  
            "@odata.type": "#ContosoAnvil.v1_2_1.AnvilTypes1",  
            "Slogan": "Contoso anvils never fail",  
            "Disclaimer": "* Most of the time"  
        },  
        "Contoso_biz": {  
            "@odata.type": "#ContosoBizEngine.v1_1_0.RelatedSpeed",  
            "Speed": "Plaid"  
        },  
        "EID_412": {  
            "@odata.type": "#AdatumPowerExtensions.v1_0_1.PowerInfoExt",  
            "ReadingInfo": {  
                "Accuracy": "5",  
                "IntervalSeconds": "20"  
            }  
        },  
        "Contoso_CustomerA": {  
            "@odata.type": "#ContosoCustomerASling.v1_0_0.SlingPower",  
            "AvailableTargets": [ "Rabbit", "Duck", "Runner" ],  
            "LaunchPowerOptions": [ "Low", "Medium", "Eliminate" ],  
            "LaunchPower": "Eliminate",  
        }  
    }  
}
```

```

        "Target": "Rabbit"
    },
},
...
}
```

### 9.8.8 OEM actions

OEM-specific actions appear in the JSON payload as properties of the `Oem` object, nested under an [Actions property](#).

The name of the property that represents the action, which shall follow the form:

```
#<OEMSchemaName>.<Action>
```

where

- `<OEMSchemaName>` is the name of the schema containing the OEM extension. To avoid naming collisions with other OEM definitions, the schema name should start with the defining organization's unique OEM identifier, including possible subdivisioning. Separator underscores (`_`) may be excluded for improved readability.
- `<Action>` is the action name.

```
{
  "Actions": {
    "Oem": {
      "#Contoso_ABC_ComputerSystem.Ping": {
        "target": "/redfish/v1/Systems/1/Actions/Oem/Contoso_ABC_ComputerSystem.Ping",
        "@Redfish.ActionInfo": "/redfish/v1/Systems/1/Oem/Contoso_ABC/PingActionInfo"
      },
      "#ContosoCustomerAComputerSystem.CustomPing": {
        "target": "/redfish/v1/Systems/1/Actions/Oem/ContosoCustomerAComputerSystem.CustomPing",
        "@Redfish.ActionInfo": "/redfish/v1/Systems/1/Oem/ContosoCustomerA/CustomPingActionInfo"
      }
    },
    ...
  }
}
```

The URI of the OEM action in the `target` property shall be in the form:

```
<ResourceUri>/Actions/Oem/<OEMSchemaName>.<Action>
```

where

- `<ResourceUri>` is the URI of the resource that supports invoking the action. For example, `/redfish/v1/Systems/1/`.
- `Actions` is the name of the property containing the actions for a resource.
- `Oem` is the name of the OEM property within the `Actions` property.
- `<OEMSchemaName>. <Action>` is the name of the schema containing the OEM extension followed by the action name. For example, `Contoso_ABC_ComputerSystem.Ping`.

The URI to the `ActionInfo` resource should be in the form:

```
<ResourceUri>/Oem/<OemIdentifier>/<Action>ActionInfo
```

where

- `<ResourceUri>` is the URI of the resource that supports invoking the action. For example, `/redfish/v1/Systems/1/`.
- `<OemIdentifier>` is the unique identifier of the OEM, including possible subdivisioning, that follows the same naming as defined in the [OEM-specified object naming](#) clause. Separator underscores (`_`) may be excluded for improved readability.
- `<Action>` is the action name.

## 9.9 Payload annotations

### 9.9.1 Payload annotations overview

[Resources](#), [objects within a resource](#), and [properties](#) may include additional annotations as properties with the name, in the format:

```
[<PropertyName>]@<Namespace>. <TermName>
```

where

- `<PropertyName>` is the name of the property to annotate. If absent, the annotation applies to the entire JSON object, which may be an entire resource.
- `<Namespace>` is the namespace that defines the annotation term.
- `<TermName>` is the annotation term to apply to the resource or property of the resource.

Services shall limit the annotation usage to the `odata`, `Redfish`, and `Message` namespaces. The [OData JSON Format](#) Specification defines the `odata` namespace. The `Redfish` namespace is an alias for the `RedfishExtensions.v1_0_0` namespace.

The client can get the definition of the annotation from the [OData metadata document](#), the [HTTP Link header](#), or

may ignore the annotation entirely, but should not fail reading the resource due to unrecognized annotations, including new annotations that the `Redfish` namespace defines.

### 9.9.2 Allowable values for strings

Services may use the `@Redfish.AllowableValues` annotation to specify the list of allowable values for a string **property** in **modification requests** or **action** parameter. The values of the annotation should only include those values that are both supported by the service and currently available as valid values for the particular instance of the property or action parameter. The annotation shall contain a JSON array of strings that define the allowable values for the property or action parameter.

The following example shows that the `FavoriteFruit` property supports four values.

```
{  
    "FavoriteFruit": "Kiwi",  
    "FavoriteFruit@Redfish.AllowableValues": [ "Orange", "Pineapple", "Kiwi", "Starfruit" ]  
}
```

### 9.9.3 Allowable patterns for string values

Services may use the `@Redfish.AllowablePattern` annotation to specify a pattern that describes the valid values for a string **property** in **modification requests** or **action** parameter. The annotation shall contain a regular expression that describes the supported pattern for the property or action parameter. If a pattern is specified in the schema definition for this property, this annotation may further restrict the allowable values, but shall not allow values which would violate the schema-defined pattern. Services shall not use the `@Redfish.AllowablePattern` for **enumerations**.

The following example shows the `AssetTag` property allows alphanumeric characters, colons, and dashes, and allows a string length up to 31 characters.

```
{  
    "AssetTag": "22HOU-34566",  
    "AssetTag@Redfish.AllowablePattern": "^[\\w:-]{0,31}$"  
}
```

### 9.9.4 Allowable values for numbers and durations

Services may use the `@Redfish.AllowableNumbers` annotation to specify one or more ranges of allowable values and an optional incremental step value between valid values for a numeric **property** in **modification requests** or **action** parameter or **duration** property or action parameter. The annotation shall contain an array of strings, each specifying a range of values and an optional step value. Each element in the array shall contain a number or a duration in the format:

```
<min>:<max>:<step>
```

where

- `<min>` is the supported value or the lowest value in an inclusive range.
- `<max>` is the highest value in an inclusive range.
- `<step>` is the incremental step value added to the `<min>` value in series within the inclusive range.

If the value does not contain `:` characters, the value specifies a single supported value.

If the value specifies a range and `<min>` is omitted, the minimum supported value shall be assumed to be a value of zero.

If the value specifies a range and `:<step>` is omitted, no step is defined for the supported range.

The following example shows the usage of the `@Redfish.AllowableNumbers` annotation for different properties.

`PacketSizeBytes` supports a range of 1024 to 65536 in increments of 256. `TemperatureThresholdCelsius` supports a range of 0 to 50 with no step restrictions. `ClockSpeedMHz` supports 800, 1150, or 1600 to 5000 in increments of 100. `TimeoutDuration` supports a duration between 5 minutes and 24 hours, in 5 minute increments.

```
{
  "PacketSizeBytes": 2048,
  "PacketSizeBytes@Redfish.AllowableNumbers": [ "1024:65536:256" ],
  "TemperatureThresholdCelsius": 37,
  "TemperatureThresholdCelsius@Redfish.AllowableNumbers": [ ":50" ],
  "ClockSpeedMHz": 2200,
  "ClockSpeedMHz@Redfish.AllowableNumbers": [ "800", "1150", "1600:5000:100" ],
  "TimeoutDuration": "PT2H",
  "TimeoutDuration@Redfish.AllowableNumbers": [ "PT5M:PT24H:PT5M" ]
}
```

## 9.9.5 Extended information

The following clauses describe the methods of providing extended information:

- [Extended object information](#)
- [Extended property information](#)

### 9.9.5.1 Extended object information

To specify object-level status information, services may annotate a JSON object with the `@Message.ExtendedInfo` annotation.

```
{
  "@odata.id": "/redfish/v1/Managers/1/SerialInterfaces/1",
  "@odata.type": "#SerialInterface.v1_0_0.SerialInterface",
  "Name": "Managed Serial Interface 1",
  "Description": "Management for Serial Interface",
  "Status": {
    "State": "Enabled",
    "Health": "OK"
  },
  "InterfaceEnabled": true,
  "SignalType": "Rs232",
  "BitRate": "115200",
  "Parity": "None",
  "DataBits": "8",
  "StopBits": "1",
  "FlowControl": "None",
  "ConnectorType": "RJ45",
  "PinOut": "Cyclades",
  "@Message.ExtendedInfo": [
    {
      "MessageId": "Base.1.8.PropertyDuplicate",
      "Message": "Indicates that a duplicate property was included in the request body.",
      "RelatedProperties": [
        "/InterfaceEnabled"
      ],
      "Severity": "Warning",
      "MessageSeverity": "Warning",
      "Resolution": "Remove the duplicate property from the request body and resubmit the request if the operation failed."
    }
  ]
}
```

The property contains an array of [message objects](#).

### 9.9.5.2 Extended property information

Services may use `@Message.ExtendedInfo`, prepended with the name of the property to annotate an individual property in a JSON object with extended information:

```
{
  "@odata.id": "/redfish/v1/Managers/1/SerialInterfaces/1",
  "@odata.type": "#SerialInterface.v1_0_0.SerialInterface",
  "Name": "Managed Serial Interface 1",
  "Description": "Management for Serial Interface",
  "Status": {
    "State": "Enabled",
    "Health": "OK"
  },
}
```

```
    "InterfaceEnabled": true,
    "SignalType": "Rs232",
    "BitRate": 115200,
    "Parity": "None",
    "DataBits": 8,
    "StopBits": 1,
    "FlowControl": "None",
    "ConnectorType": "RJ45",
    "PinOut": "Cyclades",
    "PinOut@Message.ExtendedInfo": [
        {
            "MessageId": "Base.1.8.PropertyValueNotInList",
            "Message": "The value Contoso for the property PinOut is not in the list of acceptable values.",
            "Severity": "Warning",
            "MessageSeverity": "Warning",
            "Resolution": "Choose an enumeration list value that the implementation supports. Resubmit the request if the operation failed."
        }
    ]
}
```

### 9.9.5.3 Extended information implementation notes

Services should provide the extended object or property payload annotations in responses to `PATCH` or `PUT` requests for the following conditions:

- If a client attempts to modify a property that is read-only, include the `PropertyNotWritable` message from the Base Message Registry.
- If an internal error prevents an update to a property, include the `PropertyNotUpdated` message from the Base Message Registry.
- If properties are not updated and the service does not have sufficient space to report all messages, include the `MaximumErrorsExceeded` message from the Base Message Registry.

Services should provide the extended object or property payload annotations in responses to `PATCH` or `PUT` requests for other non-successful conditions not listed previously.

### 9.9.6 Action info annotation

The `@Redfish.ActionInfo` term within the [action representation](#) conveys the parameter requirements and allowable values on parameters for [actions](#). This term contains a URI to the `ActionInfo` resource.

The URI to the `ActionInfo` resource should be in the form:

```
<ResourceUri>/<Action>ActionInfo
```

where

- `<ResourceUri>` is the URI of the resource that supports invoking the action. For example, `/redfish/v1/Systems/1/`.
- `<Action>` is the action name.

Example `#ComputerSystem.Reset` action with the `@Redfish.ActionInfo` annotation and resource:

```
{
  "Actions": {
    "#ComputerSystem.Reset": {
      "target": "/redfish/v1/Systems/1/Actions/ComputerSystem.Reset",
      "@Redfish.ActionInfo": "/redfish/v1/Systems/1/ResetActionInfo"
    }
  },
  ...
}
```

The `ResetActionInfo` resource contains a more detailed description of the parameters and the supported values. This resource follows the `ActionInfo` schema definition.

```
{
  "@odata.id": "/redfish/v1/Systems/1/ResetActionInfo",
  "@odata.type": "#ActionInfo.v1_0_0.ActionInfo",
  "Id": "ResetActionInfo",
  "Name": "Reset Action Info",
  "Parameters": [
    {
      "Name": "ResetType",
      "Required": true,
      "DataType": "String",
      "AllowableValues": ["On", "ForceOff", "ForceRestart", "Nmi",
        "ForceOn", "PushPowerButton"]
    }
  ]
}
```

### 9.9.7 Settings and settings apply time annotations

See the [Settings resource](#) clause.

### 9.9.8 Operation apply time and operation apply time support annotations

See the [Operation apply time](#) clause.

### 9.9.9 Maintenance window annotation

The [settings apply time](#) and [operation apply time](#) annotations enable an operation to be performed during a

maintenance window. The `@Redfish.MaintenanceWindow` term at the root of a resource configures the start time and duration of a maintenance window for a resource.

The following example body for the `/redfish/v1/Systems/1` resource configures the maintenance window to start at `2017-05-03T23:12:37-05:00` and last for 600 seconds.

```
{  
    "@odata.id": "/redfish/v1/Systems/1",  
    "@odata.type": "#ComputerSystem.v1_5_0.ComputerSystem",  
    "@Redfish.MaintenanceWindow": {  
        "@odata.type": "#Settings.v1_3_3.MaintenanceWindow",  
        "MaintenanceWindowStartTime": "2017-05-03T23:12:37-05:00",  
        "MaintenanceWindowDurationInSeconds": 600  
    },  
    ...  
}
```

### 9.9.10 Collection capabilities annotation

Resource collections may contain a collection capabilities annotation. The `@Redfish.CollectionCapabilities` term at the root of a resource collection shows what properties a client is allowed to use in a [POST request](#) for creating a resource.

The following `ComputerSystemCollection` example body contains the collection capabilities annotation. The `useCase` property contains the `ComputerSystemComposition` value, and the `CapabilitiesObject` property contains the `/redfish/v1/Systems/Capabilities` value. The resource at `/redfish/v1/Systems/Capabilities` describes the `POST` request format for creating a `ComputerSystem` resource for compositions.

```
{  
    "@odata.id": "/redfish/v1/Systems",  
    "@odata.type": "#ComputerSystemCollection.ComputerSystemCollection",  
    "Name": "Computer System Collection",  
    "Members@odata.count": 0,  
    "Members": [],  
    "@Redfish.CollectionCapabilities": {  
        "@odata.type": "#CollectionCapabilities.v1_1_0.CollectionCapabilities",  
        "Capabilities": [  
            {"CapabilitiesObject": {  
                "@odata.id": "/redfish/v1/Systems/Capabilities"  
            },  
            "UseCase": "ComputerSystemComposition",  
            "Links": {  
                "TargetCollection": {  
                    "@odata.id": "/redfish/v1/Systems"  
                }  
            }  
        ]  
    }  
}
```

```

        }
    }]
}
}
```

The `CapabilitiesObject` resource follows the same schema for the resource that the resource collection contains. It contains annotations to show which properties the client can use in the `POST` request body. Services may omit properties marked as `required` for the resource referenced by the `CapabilitiesObject` property, but may provide annotations for the property to express `POST` request body requirements. [Table 21](#) describes the `CapabilitiesObject` resource annotations. These annotations describe which properties are required, optional, or if other rules are associated with the properties.

**Table 21 — CapabilitiesObject resource annotations**

Annotation	Description
<code>&lt;PropertyName&gt;@Redfish.RequiredOnCreate</code>	Required in the <code>POST</code> request body.
<code>&lt;PropertyName&gt;@Redfish.OptionalOnCreate</code>	Not required in the <code>POST</code> request body.
<code>&lt;PropertyName&gt;@Redfish.SetOnlyOnCreate</code>	Cannot be modified after the resource is created.
<code>&lt;PropertyName&gt;@Redfish.UpdatableAfterCreate</code>	Can be modified after the resource is created.
<code>&lt;PropertyName&gt;@Redfish.AllowableValues</code>	Can be set to any of the listed values.
<code>@Redfish.RequestedCountRequired</code>	Required in the <code>POST</code> request body for the corresponding object to indicate the number of requested object instances. Used for <a href="#">composition requests</a> .
<code>@Redfish.ResourceBlockLimits</code>	Indicates restrictions regarding quantities of <code>ResourceBlock</code> resources of a given type in the <code>POST</code> request body. Used for <a href="#">composition requests</a> .

Example `CapabilitiesObject` resource:

```
{
  "@odata.id": "/redfish/v1/Systems/Capabilities",
  "@odata.type": "#ComputerSystem.v1_8_0.ComputerSystem",
  "Id": "Capabilities",
  "Name": "Capabilities for the system collection",
  "Name@Redfish.RequiredOnCreate": true,
  "Name@Redfish.SetOnlyOnCreate": true,
  "Description@Redfish.OptionalOnCreate": true,
```

```

    "Description@Redfish.SetOnlyOnCreate": true,
    "HostName@Redfish.OptionalOnCreate": true,
    "HostName@Redfish.UpdatableAfterCreate": true,
    "Links@Redfish.RequiredOnCreate": true,
    "Links": {
        "ResourceBlocks@Redfish.RequiredOnCreate": true,
        "ResourceBlocks@Redfish.UpdatableAfterCreate": true
    },
    "@Redfish.ResourceBlockLimits": {
        "MinCompute": 1,
        "MaxCompute": 1,
        "MaxStorage": 8
    }
}

```

### 9.9.11 Requested count and allow over-provisioning annotations

Table 22 describes the `@Redfish.RequestedCount` and `@Redfish.AllowOverprovisioning` annotations.

Clients use these annotations in [composition requests](#) to define the number of [resource](#) to allocate and to indicate whether the Redfish service can provision more resources than the client requests:

**Table 22 — RequestCount and AllowOverprovisioning annotations**

Annotation	Description
<code>@Redfish.RequestedCount</code>	Number of requested resources.
<code>@Redfish.AllowOverprovisioning</code>	Boolean. If <code>true</code> , the service may provision more resources than the <code>@Redfish.RequestedCount</code> annotation requests. Default is <code>false</code> .

Example client request for at least four and possibly more `Processor` resources:

```

{
    "Processors": {
        "Members": [
            {
                "@Redfish.RequestedCount": 4,
                "@Redfish.AllowOverprovisioning": true
            }
        ],
        ...
    }
}

```

### 9.9.12 Zone affinity annotation

The zone affinity annotation is used by clients in [composition requests](#) to indicate the components for the composition come from the specified resource zone. The `@Redfish.ZoneAffinity` term in the request body contains the value of the `Id` property of the requested resource zone.

Example client request for components to be allocated from the resource zone with the `Id` property containing `1`:

```
{  
    "@Redfish.ZoneAffinity": "1",  
    ...  
}
```

### 9.9.13 Supported certificates annotation

[Resource collections](#) of type `CertificateCollection` should contain a supported certificates annotation. The `@Redfish.SupportedCertificates` term at the root of a resource collection shows the different certificate formats allowed in the resource collection.

Example `CertificateCollection` that only supports PEM style certificates:

```
{  
    "@odata.id": "/redfish/v1/Managers/BMC/NetworkProtocol/HTTPS/Certificates",  
    "@odata.type": "#CertificateCollection.CertificateCollection",  
    "Name": "Certificate collection",  
    "Members@odata.count": 1,  
    "Members": [  
        {"@odata.id": "/redfish/v1/Managers/BMC/NetworkProtocol/HTTPS/Certificates/1"}  
    ],  
    "@Redfish.SupportedCertificates": ["PEM"]  
}
```

### 9.9.14 Deprecated annotation

Services may annotate [properties](#) with `@Redfish.Deprecated` if the schema definition has the property marked as deprecated.

Example deprecated property:

```
{
```

```
    "VendorID": "0xABCD",
    "VendorID@Redfish.Deprecated": "This property has been deprecated in favor of ModuleManufacturerID.",
    ...
}
```

### 9.9.15 Writable properties annotation

Services may annotate [resources](#) or [objects](#) with `@Redfish.WriteableProperties` to list the [properties](#) in the resource or object whose value may be changed by a [modification request](#). This annotation shall only list those properties that are defined as read-write in schema, are implemented as read-write by the service, and can be modified given the current configuration of the resource. When present at the root of a resource, the properties listed shall include only writable properties at the root level. Writable properties within objects shall be shown in additional annotation instances within the object. For arrays of objects, if the annotation appears in only the first element of the array, the values shall apply to all elements of the array.

**Note:** While "writable" is the preferred spelling of the word, the name of the annotation in payloads contains "writeable".

Example writable properties annotations that show writable properties at the root level and an array of objects within each element:

```
{
    "UserName": "John Smith",
    "Alias": "Major Goofball",
    "FavoriteFruit": "Pineapple",
    "@Redfish.WriteableProperties": [ "Alias", "FavoriteFruit" ],
    "Pets": [
        {
            "PetName": "Waffles",
            "PetType": "Donkey",
            "Friendly": true,
            "@Redfish.WriteableProperties": [ "Friendly" ]
        },
        {
            "PetName": "Fluffy",
            "PetType": "Snake",
            "Friendly": false
        }
    ]
}
```

## 9.10 Settings resource

A settings resource represents the future intended state of a [resource](#). Some resources have properties that can be updated and the updates take place immediately. However, some properties need to be updated at a future point in time, such as after a system reset. While the *active resource* represents the current state, the settings resource represents the future intended state.

For resources that support a future intended state, the response shall contain a property with the [@Redfish.Settings payload annotation](#). When a settings annotation is used, the following conditions shall apply:

- The settings resource shall be of the same schema definition as the active resource.
- The settings resource should contain a subset of updatable properties from the active resource. Additionally, it shall contain [required properties](#), which are always mandatory.
- The settings resource shall not contain the `@Redfish.Settings` annotation.
- The settings resource may contain the `@Redfish.SettingsApplyTime` annotation.
- The URI for the settings resource shall reflect that it is subordinate to the active resource. The URI should be in the form `<BaseUri>/Settings` or `<BaseUri>/SD` where `<BaseUri>` is the URI of the active resource.

The settings resource shall contain the properties that are updated at a future point in time. For resources that support a future intended state, [Table 23](#) describes the behavior of supported properties in the resource and settings resource that a service should support.

**Table 23 — Active resource and settings resource property behavior**

Property	Active resource behavior	Settings resource behavior
Read-only, <a href="#">required</a> .	Returned in the resource response to a <code>GET</code> request.	Returned in the settings resource response to a <code>GET</code> request.
Read-only, not required.	Returned in the resource response to a <code>GET</code> request.	Not returned in the settings resource response to a <code>GET</code> request.
Writable, updates immediately, but not at a future point in time.	Active value returned in the resource response to a <code>GET</code> request. <a href="#">Modification requests</a> change the active value immediately.	Not returned in the settings resource response to a <code>GET</code> request. Modification requests are rejected.
Writable, updates immediately or at a future point in time.	Active value returned in the resource response to a <code>GET</code> request. <a href="#">Modification requests</a> change the active value immediately.	Future value returned in the settings resource response to a <code>GET</code> request if a future value is pending, otherwise not returned. Modification requests change the future value.

Property	Active resource behavior	Settings resource behavior
Writable, updates at a future point in time, but not immediately.	Active value returned in the resource response to a <code>GET</code> request. <a href="#">Modification requests</a> are rejected.	Future value returned in the settings resource response to a <code>GET</code> request. Modification requests change the future value.

The `@Redfish.Settings` annotation includes several properties that help clients monitor when the service has consumed the active resource and determine the success or failure of applying the values.

- The `Messages` property is a collection of messages that represent the results of the last time the values of the settings resource were applied.
- The `ETag` property contains the ETag of the settings resource that was last applied. Immediate updates made directly to the active resource are not reflected in it.
- The `Time` property indicates the time when the settings resource was last applied. Immediate updates made directly to the active resource are not reflected in it.

The following active resource example body supports a settings resource. A client can use the `SettingsObject` property to locate the URI of the settings resource.

```
{
  "@Redfish.Settings": {
    "@odata.type": "#Settings.v1_3_3.Settings",
    "SettingsObject": {
      "@odata.id": "/redfish/v1/Systems/1/Bios/SD"
    },
    "Time": "2017-05-03T23:12:37-05:00",
    "ETag": "\"A89B031B62\"",
    "Messages": [
      {
        "MessageId": "Base.1.8.PropertyNotWritable",
        "RelatedProperties": ["/Attributes/ProcTurboMode"]
      }
    ],
    ...
  }
}
```

If a service enables a client to indicate when to apply settings:

- The settings resource shall contain a property with the `@Redfish.SettingsApplyTime` annotation.
  - Only settings resources shall contain the `@Redfish.SettingsApplyTime` annotation.
- The `@Redfish.Settings` annotation in the active resource shall contain the `SupportedApplyTimes` property for showing the allowable values for `ApplyTime` within `@Redfish.SettingsApplyTime`.
- Clients can [modify](#) the `@Redfish.SettingsApplyTime` annotation to indicate when to apply the settings.

In the following example request, the client indicates that the settings resource values are applied on reset during the specified maintenance window:

```
{  
  "@Redfish.SettingsApplyTime": {  
    "@odata.type": "#Settings.v1_3_3.PreferredApplyTime",  
    "ApplyTime": "InMaintenanceWindowOnReset",  
    "MaintenanceWindowStartTime": "2017-05-03T23:12:37-05:00",  
    "MaintenanceWindowDurationInSeconds": 600  
  },  
  ...  
}
```

## 9.11 Special resource situations

### 9.11.1 Overview

Resources need to exhibit common semantic behavior whenever possible. This can be difficult in some situations discussed in this clause.

### 9.11.2 Absent resources

Resources may be absent or their state unknown at the time a client requests information about that resource. For resources that represent removable or optional components, absence provides useful information to clients because it indicates a capability, such as an empty PCIe slot, DIMM socket, or drive bay, that would not be apparent if the resource simply did not exist.

This also applies to resources that represent a limited number of items or unconfigured capabilities within an implementation, but this usage should be applied sparingly and should not apply to resources limited in quantity due to arbitrary limits. For example, an implementation that limits `SoftwareInventory` to a maximum of 20 items should not populate 18 absent resources when only two items are present.

For resources that provide useful data in an absent state and where the URI is expected to remain constant, such as when a DIMM is removed from a memory socket, the resource should exist and should return the `Absent` value for the `State` property in the `Status` object.

In this circumstance, any required properties that have no known value shall contain `null`. Properties whose support is based on the configuration choice or the type of component installed, and therefore unknown while in the absent state, should not be returned. Likewise, subordinate resources for an absent resource should not be populated until their support can be determined. For example, the `Power` and `Thermal` resources under a `Chassis` resource should not exist for an absent Chassis.

Client software should be aware that when absent resources are later populated, the updated resource may

represent a different configuration or physical item, and previous data, including read-only properties, obtained from that resource may be invalid. For example, the `Memory` resource shows details about a single DIMM socket and the installed DIMM. When that DIMM is removed, the `Memory` resource remains as an absent resource to indicate the empty DIMM socket. Later, a new DIMM is installed in that socket, and the `Memory` resource represents data about this new DIMM, which could have completely different characteristics.

### 9.11.3 Transiently unavailable resources

Resources may be transiently unavailable due to the underlying configuration of a device. For example, a network adapter might not support management protocols while in a low power state.

For resources that provide useful data in a transiently unavailable state and where the URI is expected to remain constant, such as when a network adapter is powered off and is unable to respond to management requests, the resource should exist and should return the `StandbyOffline` value for the `State` property in the `Status` object.

In this circumstance, any required properties that have no known value shall contain `null`. Properties whose support is based on the configuration choice or the type of component installed, and therefore unknown while in the transiently unavailable state, should not be returned. Likewise, subordinate resources for a transiently unavailable resource should not be populated until their support can be determined.

## 9.12 Registries

Registry resources assist the client in interpreting Redfish resources beyond the Redfish schema definitions. To get more information about a resource, event, message, or other item, use an identifier to search registries. This information can include other properties, property restrictions, and the like. Registries are themselves resources.

Table 24 describes the types of registries that Redfish supports:

Table 24 — Registries

Registry	Description	See
Attribute	<p>Determines the semantics of each property within the <code>Attributes</code> object in a BIOS or BIOS <code>settings</code> resource. Because BIOS information can vary from platform to platform, Redfish cannot define a fixed schema for these values. Attribute registries should be assigned unique identifiers to allow users to match a given registry with compatible products.</p> <p>This registry contains both property descriptions and other information, such as data type, allowable values, and user menu information.</p>	
Message	<p>Constructs a message from a <code>MessageId</code> and other message information to present to an end user. The messages in these registries appear in both eventing and error responses to operations.</p> <p>This registry is the most common type of registry.</p>	<ul style="list-style-type: none"><li>• <a href="#">Error responses</a></li><li>• <a href="#">Eventing</a></li></ul>

Registry	Description	See
Privilege	<p>Maps the resources in a Redfish service to the privileges that can complete specified operations against those resources.</p> <p>A client can use this information to:</p> <ul style="list-style-type: none"> <li>Determine which roles should have specific privileges.</li> <li>Map accounts to those roles so that the accounts can complete operations on Redfish resources.</li> </ul>	Privilege model

## 9.13 Schema annotations

### 9.13.1 Schema annotations overview

The schema definitions of the data model use schema annotations to provide additional documentation for developers. This clause describes the different types of schema annotations that the Redfish data model uses. For information about how each of the annotations are implemented in their respective schema languages, see the [Schema definition languages](#) clause.

### 9.13.2 Description annotation

The description annotation can be applied to any type, property, action, or parameter to provide a description of Redfish schema elements suitable for end users or user interface help text.

All schemas that are published or republished by DMTF's Redfish Forum shall include a description annotation on the following schema definitions:

- Redfish types
- [Properties](#)
- [Reference properties](#)
- Enumeration values
- [Resources and resource collections](#)
- [Structured types](#)

### 9.13.3 Long description annotation

The long description annotation can be applied to any type, property, action, or parameter to provide a formal, normative specification of the schema element.

When the long descriptions in the Redfish schema contain normative language, the service shall be required to conform with the statement.

All schemas that are published or republished by DMTF's Redfish Forum shall include a long description annotation on the following schema definitions:

- Redfish types
- [Properties](#)
- [Reference properties](#)
- [Resources and resource collections](#)
- [Structured types](#)

#### 9.13.4 Resource capabilities annotation

The resource capabilities annotation can be applied to [resources](#) and [resource collections](#) to express the different type of HTTP operations a client can invoke on the given resource or resource collection.

- Insert capabilities indicate whether a client can perform a `POST` request on the resource to create a resource.
- Update capabilities indicate whether a client can perform a `PATCH` or `PUT` request on the resource.
- Delete capabilities indicate whether a client can perform a `DELETE` request on the resource.
- A service may implement a subset of the capabilities that are allowed on the resource or resource collection.

All schemas that are published or republished by DMTF's Redfish Forum for resources and resource collections shall include resource capabilities annotations.

#### 9.13.5 Resource URI patterns annotation

The resource URI patterns annotation expresses the valid URI patterns for a [resource](#) or [resource collection](#).

The strings for the URI patterns may use `{` and `}` characters to express parameters within a given URI pattern, where the value between the `{` and `}` characters contains the schema name followed by `Id`. Items between the `{` and `}` characters are treated as identifiers within the URI for given instances of a Redfish resource. Clients interpret this as a string to be replaced to access a given resource. A URI pattern may contain multiple identifier terms to support multiple levels of nested resource collections. The identifier term in the URI pattern shall match the `Id` string property for the corresponding resource, or the `MemberId` string property for the corresponding object within a resource. The process for forming the strings that are concatenated to form the URI pattern are in the [URI naming rules](#) clause.

The following string is an example URI pattern that describes a `ManagerAccount` resource: `/redfish/v1/AccountService/Accounts/{ManagerAccountId}`

Using the previous example, `{ManagerAccountId}` is replaced by the `Id` property of the corresponding `ManagerAccount` resource. If the `Id` property for a `ManagerAccount` resource is `John`, the full URI for that resource is `/redfish/v1/AccountService/Accounts/John`.

The URI patterns are constructed based on the formation of the [resource tree](#). When constructing the URI pattern for

a subordinate resource, the URI pattern for the current resource is used and appended. For example, the `RoleCollection` resource is subordinate to `AccountService`. Because the URI pattern for `AccountService` is `/redfish/v1/AccountService`, the URI pattern for the `RoleCollection` resource is `/redfish/v1/AccountService/Roles`.

In some cases, the subordinate resource is found inside of a [structured property](#) of a resource. In these cases, the name of the structured property appears in the URI pattern for the subordinate resource. For example, the `CertificateCollection` resource is subordinate to the `ManagerNetworkProtocol` resource from the `HTTPS` property. Because the URI pattern for `ManagerNetworkProtocol` is `/redfish/v1/Managers/{ManagerId}/NetworkProtocol`, the URI pattern for the `CertificateCollection` resource is `/redfish/v1/Managers/{ManagerId}/NetworkProtocol/HTTPS/Certificates`.

All schemas that are published or republished by DMTF's Redfish Forum for resources and resource collections shall be annotated with the resource URI patterns annotation.

All Redfish resources and Redfish resource collections implemented by a service shall match the URI pattern described by the resource URI patterns annotation for their given definition.

### 9.13.6 Additional properties annotation

The additional properties annotation specifies whether a type can contain additional [properties](#) outside of those defined in the schema. Types that do not support additional properties shall not contain properties beyond those described in the schema.

### 9.13.7 Permissions annotation

The permissions annotation specifies whether a client can modify the value of a [property](#), or if the property is read-only.

A service can implement a modifiable property as read-only.

The value of a write-only property, such as `Password`, cannot be read, and shall be `null` in responses.

All schemas that are published or republished by DMTF's Redfish Forum shall include a permissions annotation for all properties that are not [structured properties](#).

### 9.13.8 Required annotation

The required annotation specifies whether a service needs to support a [property](#). Required properties shall be annotated with the required annotation. All other properties are optional.

### 9.13.9 Required on create annotation

The required on create annotation specifies that a [property](#) is required to be provided by the client on creation of the [resource](#). Properties not annotated with the required on create annotation are not required to be provided by the client on a create operation.

### 9.13.10 Units of measure annotation

In addition to following the [naming rules](#), [properties](#) representing units of measure shall be annotated with the units of measure annotation to specify the units of measurement for the property.

The value of the annotation shall be a string that contains the case-sensitive "(c/s)" symbol of the unit of measure as listed in the [Unified Code for Units of Measure \(UCUM\)](#), unless the symbolic representation does not reflect common usage. If the unit in common usage is not available in UCUM, curly braces should wrap the value, such as `{value}`, to follow UCUM parsing rules. For example, `RPM` is commonly used to report fan speeds in revolutions-per-minute, and the preferred representation in UCUM is `{rev}/min`, but the value `{RPM}` is acceptable. For units with prefixes, the case-sensitive ( `c/s` ) symbol for the prefix as listed in UCUM should be prepended to the unit symbol. For example, the mebibyte ( $1024^2$  bytes), which has the UCUM `Mi` prefix and `By` symbol, would use `MiBy` as the value for the annotation. For values that also include rate information, such as megabits per second, the rate unit's symbol should be appended and use a slash ( `/` ) character as a separator. For example, `Mbit/s`.

**DEPRECATED:** Previous versions of this specification recommended `RPM` as a commonly used unit of measure for certain properties. New recommendations are provided that follow UCUM parsing rules.

### 9.13.11 Expanded resource annotation

The expanded resource annotation can be applied to a [reference property](#) to specify that the default behavior for the service is to include the contents of the related [resource](#) or [resource collection](#) in responses. This behavior follows the same semantics of the [expand query parameter](#) with a level of 1.

Reference properties annotated with this term shall be expanded by the service, even if not requested by the client. A service may page [resource collections](#).

### 9.13.12 Owning entity annotation

The owning entity annotation can be applied to a schema to specify the name of the entity responsible for development, publication, and maintenance of a given schema.

### 9.13.13 Deprecated annotation

The deprecated annotation specifies if a [property](#), enumeration, or other schema element has been deprecated.

Schema elements marked as deprecated contain a schema version that shows when the element was deprecated, as well as text that specifies the favored approach.

The deprecated annotation also specifies if [resource URI patterns](#) have been deprecated. Deprecated resource URI patterns shall also be included in the resource URI patterns annotation.

Existing and new implementations may use deprecated schema elements or URIs, but they should move to the favored approach. Deprecated schema elements may be implemented to achieve backwards compatibility. Deprecated schema elements may be removed from the next major version of the schema.

### 9.13.14 URI segment annotation

The URI segment annotation can be applied to a [reference property](#) to specify the segment appended to the URI of the resource when constructing the URI of a subordinate resource if the segment differs from the property name. For more information, see the [URI naming rules](#) clause.

### 9.13.15 URI annotation

The URI annotation can be applied to a string [property](#) to indicate the property value contains a URI. This is used for cases where a service provides a hyperlink to a URI outside of the Redfish model or where a [reference property](#) is not used to avoid undesirable expansion.

## 9.14 Versioning

---

As stated previously, a resource can be an individual entity or a resource collection, which acts as a container for a set of resources.

A [resource collection](#) does not contain any version information because it defines a single `Members` property, and the overall collection definition never grows over time.

A [resource](#) has both unversioned and versioned definitions.

References from other resources use the unversioned definition of a resource to ensure no version dependencies exist between the definitions. The unversioned definition of a resource contains no property information about the resource.

The versioned definition of a resource contains a set of properties, actions, and other definitions associated with the resource. The version of a resource follows the format:

`v<X>.<Y>.<Z>`

where

- `<x>` is an integer that represents the major version. Indicates a backward-incompatible change.
- `<y>` is an integer that represents the minor version. Indicates a minor update. Redfish introduces new functionality but does not remove any functionality. The minor version preserves compatibility with earlier minor versions. For example, a new property introduces a new minor version of the resource.
- `<z>` is an integer that represents the errata version. Indicates a fix in an earlier version. For example, a fix to a [schema annotation](#) on a property introduces an errata version of the resource.

## 9.15 Localization

---

The creation of separate localized copies of Redfish schemas and registries is allowed and encouraged. Localized schema and registry files may be submitted to DMTF for republication in the Redfish schema repository.

Property names, parameter names, and enumeration values in the JSON response payload are never localized but translated copies of those names may be provided as additional annotations in the localized schema for use by client applications. A separate file for each localized schema or registry shall be provided for each supported language. The English-language versions of Redfish schemas and registries shall be the normative versions, and alterations of meaning due to translation in localized versions of schemas and registries shall be forbidden.

Schemas and registries in non-English languages shall use the appropriate schema annotations to identify their language. Descriptive property, parameter, and enumeration text not translated into the specified language shall be removed from localized versions. This removal enables software and tools to combine normative and localized copies, especially for minor schema version differences.

# 10 File naming and publication

For consistency in publication and to enable programmatic access, all Redfish-related files shall follow a set of rules to construct the name of each file. The [Schema definition languages](#) clause describes the file name construction rules, while the following clauses describe the construction rules for other file types.

## 10.1 Registry file naming

Redfish message registry files, privilege registry files, and BIOS attribute registry files shall use the registry name to construct the file name, in this format:

```
<RegistryName>.v<MajorVersion>.v<MinorVersion>.v<Errata>.json
```

For example, the file name of the Base Message Registry v1.0.2 is `Base.v1.0.2.json`.

The registry name should be unique to avoid conflict with other registry files. The clause [OEM registries](#) describes registry name to use for OEM registry files.

## 10.2 Profile file naming

The document that describes a profile follows the Redfish schema file naming conventions. The file name format for profiles shall be:

```
<ProfileName>.v<MajorVersion>_v<MinorVersion>_v<Errata>.json
```

For example, the file name of the BasicServer profile v1.2.0 is `BasicServer.v1_2_0.json`. The file name shall include the profile name and version, which matches those property values within the document.

## 10.3 Dictionary file naming

The binary file describing a Redfish Device Enablement dictionary follows the Redfish schema file naming conventions for the schema definition language that the dictionary is converted from. Because a single dictionary file contains all minor revisions of the schema, only the major version appears in the file name. The file names for Dictionaries shall be formatted as:

```
<DictionaryName>_v<MajorVersion>.dict
```

For example, the file name of the Chassis dictionary v1.2.0 is `chassis_v1.dict`.

## 10.4 Localized file naming

Localized schemas and registries shall follow the same file naming conventions as the English language versions. When multiple localized copies are present in a repository and which have the same file name, files in languages other than English shall be organized into sub-folders named to match the [ISO 639-1](#) language code for those files. English language files may be duplicated in an `en` sub-folder for consistency.

## 10.5 DMTF Redfish file repository

All Redfish schemas, registries, dictionaries, and profiles published or republished by DMTF's Redfish Forum are available from the [DMTF website](#) for download. Programs may use the following durable URLs to access the repository. Programs incorporating remote repository access should implement a local cache to reduce latency, program requirements for Internet access and undue traffic burden on DMTF's website.

Organizations creating Redfish-related files such as OEM schemas, Redfish interoperability profiles, or message registries are encouraged to use the form at <https://redfish.dmtf.org/redfish/portal> to submit those files to DMTF for republication in DMTF's Redfish file repository.

[Table 25](#) describes how files are organized on the site:

**Table 25 — Redfish file repository**

URL	Folder contents
<code>redfish.dmtf.org/schemas</code>	Current (most recent minor or errata) release of each schema file in CSDL, JSON Schema, and/or OpenAPI formats.
<code>redfish.dmtf.org/schemas/v1</code>	Durable URL for programmatic access to all v1.xx schema files. Every v1.xx minor or errata release of each schema file in CSDL, JSON Schema, OpenAPI formats.
<code>redfish.dmtf.org/schemas/v1/{code}</code>	Durable URL for programmatic access to localized v1.xx schema files. Localized schemas are organized in sub-folders using the two-character ISO 639-1 language code as the {code} segment.
<code>redfish.dmtf.org/schemas/archive</code>	Sub-folders contain schema files specific to a particular version release.
<code>redfish.dmtf.org/registries</code>	Current (most recent minor or errata) release of each registry file.
<code>redfish.dmtf.org/registries/v1</code>	Durable URL for programmatic access to all v1.xx registry files. Every v1.xx minor or errata release of each registry file.
<code>redfish.dmtf.org/registries/v1/{code}</code>	Durable URL for programmatic access to localized v1.xx registry files. Localized schemas are organized in sub-folders using the two-character ISO 639-1 language code as the {code} segment.
<code>redfish.dmtf.org/registries/archive</code>	Sub-folders contain registry files specific to a particular version release.

URL	Folder contents
<a href="http://redfish.dmtf.org/profiles">redfish.dmtf.org/profiles</a>	Current release of each Redfish interoperability profile (.json) file and associated documentation.
<a href="http://redfish.dmtf.org/profiles/v1">redfish.dmtf.org/profiles/v1</a>	Durable URL for programmatic access to all v1.xx Redfish interoperability profile (.json) files.
<a href="http://redfish.dmtf.org/profiles/archive">redfish.dmtf.org/profiles/archive</a>	Sub-folders contain profile files specific to a particular profile version or release.
<a href="http://redfish.dmtf.org/dictionaries">redfish.dmtf.org/dictionaries</a>	Durable URL for programmatic access to all v1.xx Redfish Device Enablement dictionary files.
<a href="http://redfish.dmtf.org/dictionaries/v1">redfish.dmtf.org/dictionaries/v1</a>	Durable URL for programmatic access to all v1.xx Redfish Device Enablement dictionary files.
<a href="http://redfish.dmtf.org/dictionaries/archive">redfish.dmtf.org/dictionaries/archive</a>	Sub-folders contain dictionary files specific to a particular version release.

# 11 Schema definition languages

Individual resources and their dependent types and actions are defined within a Redfish schema document. This clause describes how these documents are constructed in the following formats:

- [OData Common Schema Definition Language](#)
- [JSON Schema](#)
- [OpenAPI](#)

## 11.1 OData Common Schema Definition Language

### 11.1.1 OData Common Schema Definition Language overview

OData Common Schema Definition Language (CSDL) is an XML schema format defined by the [OData CSDL](#) Specification. The following clause describes how Redfish uses CSDL to describe resources and resource collections.

### 11.1.2 File naming conventions for CSDL

Redfish CSDL schema files shall be named using the [resource type](#) name for the schema, followed by `_v` and the major version of the schema. Because a single CSDL schema file contains all minor revisions of the schema, only the major version appears in the file name. The file name shall be formatted as:

```
<ResourceType>_v<MajorVersion>.xml
```

For example, version 1.3.0 of the Chassis schema is `Chassis_v1.xml`.

### 11.1.3 Core CSDL files

[Table 26](#) describes the core CSDL files:

**Table 26 — Core CSDL files**

File	Description
<code>RedfishError_v1.xml</code>	Payload definition of the <a href="#">Redfish error response</a> .
<code>RedfishExtensions_v1.xml</code>	All definitions for Redfish types and annotations.
<code>Resource_v1.xml</code>	All base definitions for resources, resource collections, and common properties, such as <code>Status</code> .

## 11.1.4 CSDL format

The outer element of the OData schema representation document shall be the `Edmx` element, and shall have a `Version` attribute with a value of `4.0`.

```
<edmx:Edmx xmlns:edmx="http://docs.oasis-open.org/odata/ns/edmx" Version="4.0">
  <!-- edmx:Reference and edmx:DataServices elements go here -->
</edmx:Edmx>
```

The [Referencing other CSDL files](#) and [CSDL data services](#) clauses describe the items that are found within the `Edmx` element.

### 11.1.4.1 Referencing other CSDL files

CSDL files may use `Reference` tags to reference types defined in other CSDL documents.

The `Reference` element uses the `Uri` attribute to specify a CSDL file. The `Reference` element also contains one or more `Include` tags that specify the `Namespace` attribute containing the types to be referenced, along with an optional `Alias` attribute for that namespace.

Type definitions generally reference the OData and Redfish namespaces for common type annotation terms. Redfish CSDL files shall contain the `Alias` attribute on the following namespaces:

- `Org.OData.Core.V1` is aliased as `OData`.
- `Org.OData.Measures.V1` is aliased as `Measures`.
- `RedfishExtensions.v1_0_0` is aliased as `Redfish`.
- `Validation.v1_0_0` is aliased as `Validation`.

```
<edmx:Reference
  Uri="http://docs.oasis-open.org/odata/odata/v4.0/cs01/vocabularies/Org.OData.Core.V1.xml">
  <edmx:Include Namespace="Org.OData.Core.V1" Alias="OData"/>
</edmx:Reference>
<edmx:Reference
  Uri="http://docs.oasis-open.org/odata/odata/v4.0/os/vocabularies/Org.OData.Measures.V1.xml">
  <edmx:Include Namespace="Org.OData.Measures.V1" Alias="Measures"/>
</edmx:Reference>
<edmx:Reference Uri="http://redfish.dmtf.org/schemas/v1/RedfishExtensions_v1.xml">
  <edmx:Include Namespace="RedfishExtensions.v1_0_0" Alias="Redfish"/>
  <edmx:Include Namespace="Validation.v1_0_0" Alias="Validation"/>
</edmx:Reference>
<edmx:Reference Uri="http://redfish.dmtf.org/schemas/v1/Resource_v1.xml">
  <edmx:Include Namespace="Resource"/>
```

```
<edmx:Include Namespace="Resource.v1_0_0"/>
</edmx:Reference>
```

#### 11.1.4.2 CSDL data services

Define structures, enumerations, and other definitions in CSDL within a namespace. Use a `Schema` tag to define the schema and use the `Namespace` attribute to declare the name of the namespace.

Redfish uses namespaces to differentiate different versions of the schema. CSDL enables structures to inherit from other structures, which enables newer namespaces to define only the changes. The [Elements of CSDL namespaces](#) clause describes this behavior.

Namespaces containing unversioned `resource` and `resource collection` definitions shall use the `resource type` to name the namespace, in this format:

```
<ResourceType>
```

For example, the unversioned namespace of the Chassis resource is `Chassis`.

Namespaces containing versioned `resource` definitions shall use the `resource type` to name the namespace, in this format:

```
<ResourceType>.v<MajorVersion>_<MinorVersion>_<Errata>
```

For example, the version 1.3.0 namespace of the Chassis resource is `Chassis.v1_3_0`.

The `Schema` element is a child of the `DataServices` element, which is a child of the `Edmx` element:

```
<edmx:DataServices>
  <Schema xmlns="http://docs.oasis-open.org/odata/ns/edm" Namespace="MyTypes.v1_0_0">
    <!-- Type definitions for version 1.0.0 of MyTypes go here -->
  </Schema>
  <Schema xmlns="http://docs.oasis-open.org/odata/ns/edm" Namespace="MyTypes.v1_1_0">
    <!-- Type definitions for version 1.1.0 of MyTypes go here -->
  </Schema>
</edmx:DataServices>
```

#### 11.1.5 Elements of CSDL namespaces

The following clauses describe the definitions within each namespace:

- [Qualified names](#)

- Entity type and complex type elements

#### 11.1.5.1 Qualified names

Many definitions in CSDL use references to qualified names. CSDL defines this as a string in the form:

```
<Namespace>.<TypeName>
```

where

- <Namespace> is the namespace name.
- <TypeName> is the name of the element in the namespace.

For example, if a reference is made to `MyType.v1_0_0.MyDefinition`, the definition can be found in the `MyType.v1_0_0` namespace with an element named `MyDefinition`.

#### 11.1.5.2 Entity type and complex type elements

Use the `EntityType` and `ComplexType` tags to define the entity type and complex type elements, respectively. These elements define a JSON structure and their set of properties by defining `property elements` and `navigation property elements` within the `EntityType` or `ComplexType` tags.

All entity types and complex types shall contain a `Name` attribute, which specifies the name of the definition.

Entity types and complex types may contain a  `BaseType` attribute, which specifies a `qualified name`. When the  `BaseType` attribute is present, all definitions of the referenced  `BaseType` are available to the entity type or complex type being defined.

All resources and `resource collections` shall be defined with the entity type element. Resources inherit from `Resource.v1_0_0.Resource`, and resource collections inherit from `Resource.v1_0_0.ResourceCollection`.

All `structured properties` shall be defined with the complex type element or entity type element. Most structured properties are defined with the complex type element. Some use the entity type element that inherits from `Resource.v1_0_0.ReferenceableMember`. The entity type element enables references to be made by using the `Navigation Property element`, whereas the complex type element does not allow for this usage.

Example entity type and complex type element:

```
<EntityType Name="TypeA" BaseType="Resource.v1_0_0.Resource">
  <Annotation Term="OData.Description" String="Entity description."/>
  <Annotation Term="OData.LongDescription" String="Entity normative description."/>
  <!-- Property and navigation property definitions go here -->
</EntityType>
```

```
<ComplexType Name=".PropertyTypeA">
  <Annotation Term="OData.Description" String="Structured property description."/>
  <Annotation Term="OData.LongDescription" String="Structured property normative description."/>
  <!-- Property and navigation property definitions go here -->
</ComplexType>
```

### 11.1.5.3 Action element

Use the `Action` tag to define the action element. This element defines an [action](#) that can be performed on a [resource](#).

All Redfish actions shall be defined with the action element. All action elements shall contain a `Name` attribute, which specifies the name of the action. The action shall be represented in payloads as the [qualified name](#) of the action, preceded by `#`.

In Redfish, all action elements shall contain the `IsBound` attribute that is always set to `true`, which indicates that the action appears as a member of a structured type.

The action element shall contain one or more `Parameter` tags that specify the `Name` and `Type` of each parameter.

Because all action elements in Redfish use the `IsBound="true"` term, the first parameter is called the *binding parameter* and specifies the [structured type](#) to which the action belongs. All Redfish actions shall contain a binding parameter. The binding parameter shall be one of the following [complex type elements](#):

- For standard actions, the `Actions` complex type for the resource.
- For OEM actions, the `OemActions` complex type for the resource.

The remaining `Parameter` elements shall describe additional parameters to be passed to the action. The term `Nullable="false"` in a parameter shall indicate the parameter is required in the action request body.

```
<Schema xmlns="http://docs.oasis-open.org/odata/ns edm" Namespace="MyType">
  <Action Name="MyAction" IsBound="true">
    <Parameter Name="Thing" Type="MyType.Actions"/>
    <Parameter Name="Parameter1" Type="Edm.Boolean"/>
    <Parameter Name="Parameter2" Type="Edm.String" Nullable="false"/>
  </Action>

  <ComplexType Name="Actions">
    ...
  </ComplexType>

  ...
</Schema>
```

```
</Schema>
```

Some action parameters may specify a type that is defined by an entity type element. In these cases, the parameter in the request is a [reference object](#) to a resource within the service.

#### 11.1.5.4 Action element for OEM actions

OEM-specific actions shall be defined by using the action element with the binding parameter set to the `OemActions` complex type for the resource. For example, the following definition defines the OEM `#Contoso.Ping` action for a `ComputerSystem`.

```
<Schema xmlns="http://docs.oasis-open.org/odata/ns/edm" Namespace="Contoso">
  <Action Name="Ping" IsBound="true">
    <Parameter Name="ComputerSystem" Type="ComputerSystem.v1_0_0.OemActions"/>
  </Action>
</Schema>
```

#### 11.1.5.5 Action with a response body

A response body for an action shall be defined using the `ReturnType` tag within an action element. For example, the following definition defines the `GenerateTicket` action with a response that contains the definition specified by `GenerateTicketResponse`.

```
<Action Name="GenerateTicket" IsBound="true">
  <Parameter Name="ExampleResource" Type="ExampleResource.v1_0_0.Actions"/>
  ...
  <ReturnType Type="ExampleResource.v1_0_0.GenerateTicketResponse" Nullable="false"/>
</Action>

<ComplexType Name="GenerateTicketResponse">
  <Annotation Term="OData.AdditionalProperties" Bool="false"/>
  <Annotation Term="OData.Description" String="The response body for GenerateTicket."/>
  <Property Name="TicketId" Type="Edm.String" Nullable="false">
    <Annotation Term="OData.Permissions" EnumMember="OData.Permission/Read"/>
    <Annotation Term="OData.Description" String="The ticket identifier."/>
    <Annotation Term="Redfish.Required"/>
  </Property>
</ComplexType>
```

Using the above example, the following payload is an example response for the `GenerateTicket` action.

```
{
  "TicketId": "40478281bd0f6b9e7131db6c4f673438"
}
```

### 11.1.5.6 Property element

Properties of resources, resource collections, and structured properties are defined using the property element. The `Property` tag defines a property element inside `entity type` and `complex type` elements.

All property elements shall contain a `Name` attribute, which specifies the name of the property.

All property elements shall contain a `Type` attribute specifies the data type. The `Type` attribute shall be one of the following names or types:

- A `qualified name` that references an `enum type element`.
- A `qualified name` that references a `complex type element`.
- A primitive data type.
- An array of the previous names or types by using the `Collection` term.

Table 27 describes the primitive data types:

**Table 27 — Primitive data types**

Type	Meaning
<code>Edm.Boolean</code>	True or False.
<code>Edm.DateTimeOffset</code>	Date-time string.
<code>Edm.Decimal</code>	Number, optionally containing a decimal point.
<code>Edm.Double</code>	Number, optionally containing a decimal point and optionally containing an exponent.
<code>Edm.Duration</code>	Duration string.
<code>Edm.Guid</code>	GUID/UUID string.
<code>Edm.Int64</code>	Signed 64-bit integer.
<code>Edm.String</code>	UTF-8 string.

Property elements may specify a `Nullable` attribute. If the attribute is `false`, the property shall not contain `null`. If the attribute is `true` or absent, the property may contain `null`.

Example property element:

```
<Property Name="Property1" Type="Edm.String" Nullable="false">
  <Annotation Term="OData.Description" String="Property1 description."/>
  <Annotation Term="OData.LongDescription" String="Property1 normative description."/>
  <Annotation Term="OData.Permissions" EnumMember="OData.Permission/Read"/>
  <Annotation Term="Redfish.Required"/>
  <Annotation Term="Measures.Unit" String="Watts"/>
</Property>
```

#### 11.1.5.7 Navigation property element

Reference properties of resources, resource collections, and structured properties are defined using the navigation property element. The `NavigationProperty` tag defines a navigation property element inside [entity type](#) and [complex type elements](#).

All navigation property elements shall contain a `Name` attribute, which specifies the name of the property.

All navigation property elements shall contain a `Type` attribute specifies the data type. The `Type` attribute shall be a [qualified name](#) that references an [entity type element](#). The `Collection` term with the qualified name shall indicate the property is an array.

Navigation property elements may specify a `Nullable` attribute. If the attribute is `false`, the property shall not contain `null`. If the attribute is `true` or absent, the property may contain `null`.

Unless the reference property is to be [expanded](#), all navigation properties in Redfish shall contain the `OData.AutoExpandReferences` annotation element to show that the reference is always available.

Example navigation property element:

```
<NavigationProperty Name="RelatedType" Type="MyTypes.TypeB">
  <Annotation Term="OData.Description" String="RelatedType description."/>
  <Annotation Term="OData.LongDescription" String="RelatedType normative description."/>
  <Annotation Term="OData.AutoExpandReferences"/>
</NavigationProperty>
```

#### 11.1.5.8 Enum type element

Use the `EnumType` tag to define the enum type element. This element defines a set of enumeration values, which may be applied to one or more properties.

All enum type elements shall contain a `Name` attribute, which specifies the name of the set of enumeration values.

Enum type elements shall contain `Member` tags that define the members of the enumeration. The `Member` tags shall contain a `Name` attribute that specifies the string value of the member name.

```
<EnumType Name="EnumTypeA">
  <Annotation Term="OData.Description" String="EnumTypeA type description."/>
  <Annotation Term="OData.LongDescription" String="EnumTypeA type normative description."/>
  <Member Name="MemberA">
    <Annotation Term="OData.Description" String="The description of MemberA"/>
  </Member>
  <Member Name="MemberB">
    <Annotation Term="OData.Description" String="The description of MemberB"/>
  </Member>
</EnumType>
```

#### 11.1.5.9 Annotation element

Annotations in CSDL are expressed using the `Annotation` tag. Any schema element in CSDL may contain annotations.

The following examples show how each [Redfish schema annotation](#) is expressed in CSDL.

- The [OData Core Schema](#) defines terms with the `odata` prefix.
- The [OData Measures Schema](#) defines terms with the `Measures` prefix.
- The [RedfishExtensions Schema](#) defines terms with the `Redfish` prefix.

Example [description annotation](#):

```
<Annotation Term="OData.Description" String="The console color."/>
```

Example [long description annotation](#):

```
<Annotation Term="OData.LongDescription"
  String="This property shall contain the console color."/>
```

Example [additional properties annotation](#):

```
<Annotation Term="OData.AdditionalProperties"/>
```

Example [permissions annotation](#) (read-only):

```
<Annotation Term="OData.Permissions" EnumMember="OData.Permission/Read"/>
```

Example [permissions annotation](#) (read/write):

```
<Annotation Term="OData.Permissions" EnumMember="OData.Permission/ReadWrite"/>
```

Example [permissions annotation](#) (write-only, `null` in responses):

```
<Annotation Term="OData.Permissions" EnumMember="OData.Permission/Write"/>
```

Example [required annotation](#):

```
<Annotation Term="Redfish.Required"/>
```

Example [required on create annotation](#):

```
<Annotation Term="Redfish.RequiredOnCreate"/>
```

Example [units of measure annotation](#):

```
<Annotation Term="Measures.Unit" String="MiBy"/>
```

Example [expanded resource annotation](#):

```
<Annotation Term="OData.AutoExpand"/>
```

Example [insert capabilities annotation](#) (showing `POST` is not allowed):

```
<Annotation Term="Capabilities.InsertRestrictions">
  <Record>
    <PropertyValue Property="Insertable" Bool="false"/>
  </Record>
</Annotation>
```

Example [update capabilities annotation](#) (showing `PATCH` and `PUT` are allowed):

```
<Annotation Term="Capabilities.UpdateRestrictions">
  <Record>
    <PropertyValue Property="Updatable" Bool="true"/>
    <Annotation Term="OData.Description" String="The desired speed can be changed."/>
  </Record>
</Annotation>
```

Example [delete capabilities annotation](#) (showing `DELETE` is allowed):

```
<Annotation Term="Capabilities.DeleteRestrictions">
  <Record>
    <PropertyValue Property="Deletable" Bool="true"/>
    <Annotation Term="OData.Description" String="The resource can be deleted."/>
  </Record>
</Annotation>
```

Example [resource URI patterns annotation](#):

```
<Annotation Term="Redfish.Uris">
  <Collection>
    <String>/redfish/v1/AccountService/Accounts/{ManagerAccountId}</String>
  </Collection>
</Annotation>
```

Example [URI segment annotation](#):

```
<NavigationProperty Name="Tasks" Type="TaskService.TaskService" Nullable="false">
  <Annotation Term="OData.Permissions" EnumMember="OData.Permission/Read"/>
  <Annotation Term="OData.Description" String="The link to the task service."/>
  <Annotation Term="OData.AutoExpandReferences"/>
  <Annotation Term="Redfish.URISegment" String="TaskService"/>
</NavigationProperty>
```

Example [owning entity annotation](#):

```
<Annotation Term="Redfish.OwningEntity" String="DMTF"/>
```

Example [deprecated annotation](#):

```
<Annotation Term="Redfish.Revisions">
  <Collection>
    <Record>
      <PropertyValue Property="Kind" EnumMember="Redfish.RevisionKind/Deprecated"/>
      <PropertyValue Property="Version" String="v1_3_0"/>
      <PropertyValue Property="Description"
        String="This property has been deprecated in favor of ModuleManufacturerID."/>
    </Record>
  </Collection>
</Annotation>
```

Example [deprecated annotation](#) applied to [resource URI patterns](#):

```
<Annotation Term="Redfish.DeprecatedUris">
  <Collection>
    <String>/redfish/v1/Chassis/{ChassisId}/Thermal</String>
  </Collection>
</Annotation>
```

Example [URI annotation](#):

```
<Annotation Term="OData.IsURL"/>
```

## 11.2 JSON Schema

### 11.2.1 JSON Schema overview

The [JSON Schema Specification](#) defines a JSON format for describing JSON payloads. The following clause describes how Redfish uses JSON Schema to describe resources and resource collections.

### 11.2.2 File naming conventions for JSON Schema

Each Redfish JSON Schema file represents a single resource type.

Versioned Redfish JSON Schema files shall use the [resource type](#) to name the file, in this format:

```
<ResourceType>.v<MajorVersion>_<MinorVersion>_<Errata>.json
```

For example, version 1.3.0 of the Chassis schema is `Chassis.v1_3_0.json`.

Unversioned Redfish JSON Schema files shall use the [resource type](#) to name the file, in this format:

`<ResourceType>.json`

For example, the unversioned definition of the Chassis schema is `Chassis.json`.

### 11.2.3 Core JSON Schema files

[Table 28](#) describes the core JSON Schema files:

**Table 28 — Core JSON Schema files**

File	Description
<code>odata-v4.json</code>	Definitions for common OData properties.
<code>redfish-error.v1_0_0.json</code> and its subsequent versions	Payload definition of the <a href="#">Redfish error response</a> .
<code>redfish-schema-v1.json</code>	Extensions to the JSON Schema that define Redfish JSON Schema files.
<code>Resource.json</code> and its subsequent versions	All base definitions for resources, resource collections, and common properties, such as <code>Status</code> .

### 11.2.4 JSON Schema format

Each JSON Schema file shall contain a JSON object to describe [resources](#), [resource collections](#), and other definitions for the data model.

[Table 29](#) describes the JSON object, which contains the following terms:

**Table 29 — JSON Schema format**

Term	Description
<code>\$id</code>	Reference to the URI where the schema file is published.
<code>\$ref</code>	For a schema file that describes a resource or resource collection, the reference to the structural definition of the resource or resource collection.
<code>\$schema</code>	URI to the Redfish schema extensions for JSON Schema. The value should be <code>http://redfish.dmtf.org/schemas/v1/redfish-schema-v1.json</code> .
<code>copyright</code>	Copyright statement for the organization producing the JSON Schema.
<code>definitions</code>	Structures, enumerations, and other definitions defined by the schema.

Term	Description
<code>title</code>	For a schema file that describes a resource or resource collection, the matching <a href="#">type identifier</a> for the resource or resource collection.

## 11.2.5 JSON Schema definitions body

This clause describes the types of definitions found in the `definitions` term of a Redfish JSON Schema file.

### 11.2.5.1 Resource definitions in JSON Schema

To satisfy [versioning](#) requirements, the JSON Schema representation of each [resource](#) shall have one unversioned schema file and a set of versioned schema files.

The unversioned definition of a resource shall contain an `anyOf` statement. This statement shall consist of an array of `$ref` terms, which point to the following definitions:

- The JSON Schema definition for a [reference property](#).
- The versioned definitions of the resource.

The unversioned definition of a resource shall contain the `uris` term to express the [allowable URIs for the resource](#), and the `deletable`, `insertable`, and `updatable` terms to express the [capabilities of the resource](#).

The following example shows an unversioned resource definition in JSON Schema:

```
{
  "ComputerSystem": {
    "anyOf": [
      {"$ref": "http://redfish.dmtf.org/schemas/v1/odata.v4_0_3.json#/definitions/idRef"},
      {"$ref": "http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_0_0.json#/definitions/ComputerSystem"},
      {"$ref": "http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_0_1.json#/definitions/ComputerSystem"},
      {"$ref": "http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_6_0.json#/definitions/ComputerSystem"}
    ],
    "deletable": true,
    "description": "The ComputerSystem schema represents a general purpose machine or system.",
    "insertable": false,
    "longDescription": "This resource shall represent resources that represent a computing system.",
    "updatable": true,
    "uris": [
      "/redfish/v1/Systems/{ComputerSystemId}"
    ]
  }
}
```

```
    ...
}
```

The versioned definition of a resource shall contain the property definitions for the given version of the resource.

#### 11.2.5.2 Enumerations in JSON Schema

Table 30 describes the terms that constitute definitions for enumerations:

**Table 30 — JSON Schema enumerations**

Term	Description
enum	String array that contains the possible enumeration values.
enumDescriptions	Object that contains the <a href="#">descriptions</a> for each of the enumerations as name-value pairs.
enumLongDescriptions	Object that contains the <a href="#">long descriptions</a> for each of the enumerations as name-value pairs.
enumDeprecated	Object that contains the <a href="#">deprecation guidance</a> for each of the enumerations as name-value pairs.
enumVersionDeprecated	Object that contains the <a href="#">deprecation version information</a> for each of the enumerations as name-value pairs.
type	Because all enumerations in Redfish are strings, the <code>type</code> term always has the <code>string</code> value.

The following example shows an enumeration definition in JSON Schema:

```
{
  "Flavors": {
    "enum": ["Lit", "Blinking", "Off"],
    "enumDescriptions": {
      "Blinking": "The Indicator LED is blinking.",
      "Lit": "The Indicator LED is lit.",
      "Off": "The Indicator LED is off."
    },
    "enumLongDescriptions": {
      "Blinking": "This value shall represent the Indicator LED is in a blinking state where the LED is being turned on and off in repetition.",
      "Lit": "This value shall represent the Indicator LED is in a solid on state.",
      "Off": "This value shall represent the Indicator LED is in a solid off state."
    },
    "type": "string"
  },
  ...
}
```

```
}
```

### 11.2.5.3 Actions in JSON Schema

Versioned definitions of `resources` shall contain a definition called `Actions`. This definition is a container with a set of terms that point to the different `actions` supported by the resource. The names of standard actions shall be in the form:

```
#<ResourceType>.<ActionName>
```

Example `Actions` definition:

```
{
  "Actions": {
    "additionalProperties": false,
    "description": "The available actions for this resource.",
    "longDescription": "This type shall contain the available actions for this resource.",
    "properties": {
      "#ComputerSystem.Reset": {
        "$ref": "#/definitions/Reset"
      }
    },
    "type": "object"
  },
  ...
}
```

Another definition within the same schema file shall describe the action itself. This definition shall contain a term called `parameters` to describe the client request body. It also shall contain property definitions for the `target` and `title` properties shown in response payloads for the resource.

The following example shows a definition of an action:

```
{
  "Reset": {
    "additionalProperties": false,
    "description": "This action resets the system.",
    "longDescription": "This action shall perform a reset of the ComputerSystem.",
    "parameters": {
      "ResetType": {
        "$ref": "http://redfish.dmtf.org/schemas/v1/Resource.json#/definitions/ResetType",
        "description": "The type of reset to be performed.",
        "longDescription": "This parameter shall define the type of reset to be performed."
      }
    }
  }
}
```

```
        }
    },
    "properties": {
        "target": {
            "description": "Link to invoke action",
            "format": "uri",
            "type": "string"
        },
        "title": {
            "description": "Friendly action name",
            "type": "string"
        }
    },
    "type": "object"
},
...
}
```

Action parameters may specify a type that is a resource definition. In these cases, the parameter in the request shall contain a [reference object](#) to a resource within the service.

#### 11.2.5.4 OEM actions in JSON Schema

OEM-specific actions shall be defined by using an action definition in an appropriately named JSON Schema file. For example, the following definition defines the OEM `#ContosoNetworkDevice.Ping` action, assuming it's found in the versioned ContosoNetworkDevice JSON Schema file, such as `ContosoNetworkDevice.v1_0_0.json`.

```
{
    "Ping": {
        "additionalProperties": false,
        "parameters": {},
        "properties": {
            "target": {
                "description": "Link to invoke action",
                "format": "uri",
                "type": "string"
            },
            "title": {
                "description": "Friendly action name",
                "type": "string"
            }
        },
        "type": "object"
},
...
}
```

#### 11.2.5.5 Action with a response body

A response body for an action shall be defined using the `actionResponse` term within the action definition. For example, the following definition defines the `GenerateTicket` action with a response that contains the definition specified by `#/definitions/GenerateTicketResponse`.

```
{  
    "GenerateTicket": {  
        "actionResponse": {  
            "$ref": "#/definitions/GenerateTicketResponse"  
        },  
        "parameters": {}  
    },  
    "GenerateTicketResponse": {  
        "additionalProperties": false,  
        "description": "The response body for GenerateTicket.",  
        "properties": {  
            "TicketId": {  
                "description": "The ticket identifier.",  
                "readonly": true,  
                "type": "string"  
            }  
        },  
        "required": ["TicketId"],  
        "type": "object"  
    }  
}
```

In the previous example, the following payload is an example response for the `GenerateTicket` action.

```
{  
    "TicketId": "40478281bd0f6b9e7131db6c4f673438"  
}
```

#### 11.2.6 JSON Schema terms

Table 31 describes the JSON Schema terms that Redfish uses to provide [schema annotations](#) for Redfish JSON Schema:

**Table 31 — JSON Schema terms**

JSON Schema term	Related Redfish schema annotation
<code>description</code> <code>enumDescriptions</code>	<a href="#">Description</a>
<code>longDescription</code> <code>enumLongDescriptions</code>	<a href="#">Long description</a>
<code>additionalProperties</code>	<a href="#">Additional properties</a>
<code>readonly</code> <code>writeOnly</code>	<a href="#">Permissions</a>
<code>required</code>	<a href="#">Required</a>
<code>requiredOnCreate</code>	<a href="#">Required on create</a>
<code>units</code>	<a href="#">Units of measure</a>
<code>autoExpand</code>	<a href="#">Expanded resource</a>
<code>deletable</code> <code>insertable</code> <code>updatable</code>	<a href="#">Resource capabilities</a>
<code>uris</code>	<a href="#">Resource URI patterns</a>
<code>uriSegment</code>	<a href="#">URI segment</a>
<code>owningEntity</code>	<a href="#">Owning entity</a>
<code>deprecated</code> <code>versionDeprecated</code> <code>urisDeprecated</code>	<a href="#">Deprecated</a>

## 11.3 OpenAPI

---

### 11.3.1 OpenAPI overview

The [OpenAPI Specification](#) defines a format for describing JSON payloads and the set of URIs a client can access on a service. The following clause describes how Redfish uses OpenAPI to describe resources and resource collections.

### 11.3.2 File naming conventions for OpenAPI schema

Each Redfish OpenAPI file represents a single resource type.

Versioned Redfish OpenAPI files shall be named using the [resource type](#) name for the schema, following the format:

```
<ResourceType>.v<MajorVersion>_<MinorVersion>_<Errata>.yaml
```

For example, version 1.3.0 of the Chassis schema is `Chassis.v1_3_0.yaml`.

Unversioned Redfish OpenAPI files shall use the [resource type](#) name to name the file, in this format:

```
<ResourceType>.yaml
```

For example, the unversioned definition of the Chassis schema is `Chassis.yaml`.

### 11.3.3 Core OpenAPI schema files

[Table 32](#) describes the core OpenAPI schema files:

**Table 32 — Core OpenAPI schema files**

File	Description
<code>odata-v4.yaml</code>	Definitions for common OData properties.
<code>openapi.yaml</code>	URI paths and their respective payload structures.
<code>Resource.yaml</code> and its subsequent versions	All base definitions for resources, resource collections, and common properties, such as <code>Status</code> .

### 11.3.4 openapi.yaml

The `openapi.yaml` file is the starting point for clients to understand the construct of the service.

[Table 33](#) describes the terms that the `openapi.yaml` file contains:

**Table 33 — openapi.yaml terms**

Term	Description
<code>components</code>	Global definitions. For Redfish, contains the format of the <a href="#">Redfish error response</a> .
<code>info</code>	Structure consisting of information about what the <code>openapi.yaml</code> is describing, such as the author of the file and any contact information.
<code>openapi</code>	Version of OpenAPI the document follows.
<code>paths</code>	URIs supported by the document, with possible methods, response bodies, and request bodies.

The service shall return the `openapi.yaml` file, if present in the Redfish service, as a YAML document by using either the `application/yaml` or `application/vnd.oai.openapi` MIME types. The service may append `;charset=utf-8` to the MIME type. Note that while the `application/yaml` type is in common use today, the `application/vnd.oai.openapi` type was recently defined and approved specifically to support OpenAPI. Implementations should use caution when selecting the MIME type as this specification may change in the future to reflect adoption of the OpenAPI-defined MIME type.

The `paths` term shall contain an array of the [possible URLs](#). Each URI shall contain `methods` supported by the URI. Each method shall contain the possible response bodies and request bodies.

Example `paths` entry for a resource:

```
/redfish/v1/Systems/{ComputerSystemId}:
  get:
    parameters:
      - description: The value of the Id property of the ComputerSystem resource
        in: path
        name: ComputerSystemId
        required: true
        schema:
          type: string
    responses:
      '200':
        content:
          application/json:
            schema:
              $ref: http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_6_0.yaml#/components/schemas/
ComputerSystem
          description: The response contains a representation of the ComputerSystem
          resource
      default:
        content:
          application/json:
            schema:
              $ref: '#/components/schemas/RedfishError'
        description: Error condition
```

Example `paths` entry for an action:

```
/redfish/v1/Systems/{ComputerSystemId}/Actions/ComputerSystem.Reset:
  post:
    parameters:
      - description: The value of the Id property of the ComputerSystem resource
        in: path
        name: ComputerSystemId
        required: true
```

```

type: string
requestBody:
  content:
    application/json:
      schema:
        $ref: http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_6_0.yaml#/components/schemas/
ResetRequestBody
  required: true
responses:
  '200':
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/RedfishError'
    description: The response contains the results of the Reset action
  '202':
    content:
      application/json:
        schema:
          $ref: http://redfish.dmtf.org/schemas/v1/Task.v1_4_0.yaml#/components/schemas/Task
    description: Accepted; a task has been generated
  '204':
    description: Success, but no response data
default:
  content:
    application/json:
      schema:
        $ref: '#/components/schemas/RedfishError'
description: Error condition

```

### 11.3.5 OpenAPI file format

With the exception of `openapi.yaml`, each OpenAPI file shall contain a YAML object to describe [resources](#), [resource collections](#), or other definitions for the data model. [Table 34](#) describes the terms that the YAML object contains:

**Table 34 — YAML object terms**

Term	Description
<code>components</code>	Structures, enumerations, and other definitions defined by the schema.
<code>x-copyright</code>	Copyright statement for the organization producing the OpenAPI file.
<code>title</code>	For a schema file that describes a resource or resource collection, the matching <a href="#">type identifier</a> for the resource or resource collection.

## 11.3.6 OpenAPI components body

This clause describes the types of definitions that can be found in the `components` term of a Redfish OpenAPI file.

### 11.3.6.1 Resource definitions in OpenAPI

To satisfy [versioning](#) requirements, the OpenAPI representation of each [resource](#) shall have one unversioned schema file and a set of versioned schema files.

The unversioned definition of a resource shall contain an `anyOf` statement. This statement shall consist of an array of `$ref` terms, which point to the following definitions:

- The OpenAPI definition for a [reference property](#).
- The versioned definitions of the resource.

Example unversioned resource definition in OpenAPI:

```
ComputerSystem:  
  anyOf:  
    - $ref: http://redfish.dmtf.org/schemas/v1/odata.v4_0_3.yaml#/components/schemas/idRef  
    - $ref: http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_0_0.yaml#/components/schemas/  
      ComputerSystem  
    - $ref: http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_0_1.yaml#/components/schemas/  
      ComputerSystem  
    - $ref: http://redfish.dmtf.org/schemas/v1/ComputerSystem.v1_6_0.yaml#/components/schemas/  
      ComputerSystem  
  description: The ComputerSystem schema represents a general purpose machine  
    or system.  
  x-longDescription: This resource shall be used to represent resources that represent  
    a computing system.
```

The versioned definition of a resource shall contain the property definitions for the given version of the resource.

### 11.3.6.2 Enumerations in OpenAPI

[Table 35](#) describes the terms in OpenAPI enumerations:

**Table 35 — OpenAPI enumerations**

Term	Description
<code>enum</code>	String array that contains the possible enumeration values.
<code>type</code>	Because all enumerations in Redfish are strings, the <code>type</code> term always has the value <code>string</code> .

Term	Description
x-enumDescriptions	Object that contains the <a href="#">descriptions</a> for each of the enumerations as name-value pairs.
x-enumLongDescriptions	Object that contains the <a href="#">long descriptions</a> for each enumeration as a name-value pair.
x-enumDeprecated	Object that contains the <a href="#">deprecation guidance</a> for each of the enumerations as name-value pairs.
x-enumVersionDeprecated	Object that contains the <a href="#">deprecation version information</a> for each of the enumerations as name-value pairs.

Example enumeration definition in OpenAPI:

```
IndicatorLED:
  enum:
    - Lit
    - Blinking
    - 'Off'
  type: string
  x-enumDescriptions:
    Blinking: The Indicator LED is blinking.
    Lit: The Indicator LED is lit.
    'Off': The Indicator LED is off.
  x-enumLongDescriptions:
    Blinking: This value shall represent the Indicator LED is in a blinking state
      where the LED is being turned on and off in repetition.
    Lit: This value shall represent the Indicator LED is in a solid on state.
    'Off': This value shall represent the Indicator LED is in a solid off state.
```

### 11.3.6.3 Actions in OpenAPI

Versioned definitions of [resources](#) shall contain a definition called `Actions`. This definition is a container with a set of terms that point to the different [actions](#) supported by the resource. The names of standard actions shall be in the form:

```
#<ResourceType>.<ActionName>
```

Example `Actions` definition:

```
Actions:
  additionalProperties: false
  description: The available actions for this resource.
  properties:
    '#ComputerSystem.Reset':
      $ref: '#/components/schemas/Reset'
  type: object
```

```
x-longDescription: This type shall contain the available actions for this resource.
```

Another definition within the same schema file shall describe the action itself. This definition shall contain property definitions for the `target` and `title` properties shown in response payloads for the resource.

The following example shows a definition of an action:

```
Reset:
  additionalProperties: false
  description: This action resets the system.
  properties:
    target:
      description: Link to invoke action
      format: uri
      type: string
    title:
      description: Friendly action name
      type: string
    type: object
  x-longDescription: This action shall reset the ComputerSystem.
```

The parameters for the action shall be defined in another definition with `RequestBody` appended to the name of the action. This gets mapped from the `openapi.yaml` file for expressing the `POST` method for the URI of the action.

The following example shows a definition of parameters of an action:

```
ResetRequestBody:
  additionalProperties: false
  description: This action resets the system.
  properties:
    ResetType:
      $ref: http://redfish.dmtf.org/schemas/v1/Resource.yaml#/components/schemas/ResetType
      description: The reset type.
    x-longDescription: This parameter shall define the type of reset to perform.
  type: object
  x-longDescription: This action shall reset the ComputerSystem.
```

#### 11.3.6.4 OEM actions in OpenAPI

OEM-specific actions shall be defined by using an action definition in an appropriately named OpenAPI file. For example, the following definition defines the OEM `#ContosoNetworkDevice.Ping` action, assuming it's found in the versioned ContosoNetworkDevice OpenAPI file with a name, such as `ContosoNetworkDevice.v1_0_0.yaml`.

```

Ping:
  additionalProperties: false
  properties:
    target:
      description: Link to invoke action
      format: uri
      type: string
    title:
      description: Friendly action name
      type: string
    type: object
PingRequestBody:
  additionalProperties: false
  properties: {}
  type: object

```

### 11.3.7 OpenAPI terms used by Redfish

Table 36 describes the OpenAPI terms that Redfish uses to provide [schema annotations](#) for Redfish OpenAPI files:

**Table 36 — OpenAPI terms used by Redfish**

OpenAPI term	Related Redfish schema annotation
<code>description</code> <code>x-enumDescriptions</code>	<a href="#">Description</a>
<code>x-longDescription</code> <code>x-enumLongDescriptions</code>	<a href="#">Long description</a>
<code>additionalProperties</code>	<a href="#">Additional properties</a>
<code>readOnly</code> <code>writeOnly</code>	<a href="#">Permissions</a>
<code>required</code>	<a href="#">Required</a>
<code>x-requiredOnCreate</code>	<a href="#">Required on create</a>
<code>x-units</code>	<a href="#">Units of measure</a>
<code>x-autoExpand</code>	<a href="#">Expanded resource</a>
<code>x-uriSegment</code>	<a href="#">URI segment</a>
<code>x-owningEntity</code>	<a href="#">Owning entity</a>
<code>deprecated</code> <code>x-deprecatedReason</code> <code>x-versionDeprecated</code>	<a href="#">Deprecated</a>

## 11.4 Schema modification rules

Schema referenced from the implementation may vary from the canonical definitions of those schema defined by the Redfish schema or other entities, provided they adhere to the following rules. Clients should take this into consideration when attempting operations on the resources defined by schema.

- Modified schema may constrain a [read/write property to be read only](#).
- Modified schema may constrain a property by adding length annotations to properties that do not have those annotations.
- Modified schema may constrain a property by adding a pattern annotation to properties that do not have that annotation.
- Modified schema may constrain the [capabilities of a resource or resource collection](#) to remove support for HTTP operations.
  - Modified schema may change the update capabilities to indicate a client can perform a `PATCH` or `PUT` request on the resource to support writable OEM properties.
- Modified schema may remove [properties](#) that are not [required](#).
- Modified schema may remove [actions](#).
- Modified schema may remove action parameters that are not required.
- Modified schema may change [description annotations](#).
- Modified schema may change any external references to point to Redfish schema that adheres to the modification rules.
- Modified schema may change the [owning entity annotation](#) to specify who made the modifications.
- Modified schema may remove URIs from the [resource URI patterns annotation](#).
- Modified schema may add URIs to the [resource URI patterns annotation](#) to define OEM URIs for standard resources and shall follow the OEM URI rules specified by the [OEM URIs](#) clause.
- Modified schema may alter non-semantic text, such as spaces, tabs, and line breaks.
- Other modifications to the schema shall not be allowed.

Services that collect and surface schema files from managed devices, such as when [aggregating other services](#), shall handle cases where different devices contain different modifications to the same schema file. This specification does not require a specific method to resolve these conflicts, but some options include:

- Merging the schema files from the managed devices.
- Replacing the schema files from the managed devices with unmodified schema files.

# 12 Service details

## 12.1 Eventing

### 12.1.1 Eventing overview

This clause describes how to use the REST-based mechanism to subscribe to and receive event messages.

**Note:** For security implications of eventing, see the [Security details](#) clause.

The Redfish service requires a client or administrator to create [\*subscriptions\*](#) to receive events.

To create a subscription, use one of these methods:

- Directly HTTP [`POST`](#) to the subscription collection.
- Indirectly [open a server-sent events \(SSE\) connection](#) for the event service.

### 12.1.2 POST to subscription collection

To locate the event service, the client traverses the Redfish service interface. The event service is located in the service root, as described in the `ServiceRoot` schema.

After the client discovers the service, they perform an HTTP `POST` on the resource collection URI for `Subscriptions` in the event service to subscribe to events. For the [\*subscription\*](#) body syntax, see the Redfish `EventDestination` schema. This request includes:

- The URI where an event-receiver client expects events to be sent. When an event is triggered within the Redfish service, the service sends an event to that URI.
- The type of events to send.

If the subscription request succeeds, the service shall return:

- An HTTP [`201 Created`](#) status code.
- The `Location` header that contains a URI of the newly created subscription resource.

If the subscription request succeeds, the service should return:

- A response body containing a representation of the subscription resource that conforms to the `EventDestination` schema.

After a subscription is registered with the service, clients begin receiving events. Clients do not receive events retroactively. The service does not retain historical events.

Services shall:

- Support *push* style eventing for all resources that can send events.
- Respond to a request to create a subscription with an error if the body of the request is conflicting. For instance, if parameters in the request are not supported, the service shall return the HTTP `400 Bad Request` status code.
- Retain subscriptions as persistent across service restarts.

Services shall not:

- *Push* events by using HTTP `POST` unless an event subscription has been created. To terminate the event stream at any time, either the client or the service can delete the subscription.
- Send a *push* event payload larger than 1 Mebibyte (MiB). If more than 1 MiB of data is to be sent, the service shall divide the payload on the nearest `Event` entry such that the total payload transmitted to the client is less than 1 MiB. This restriction shall not apply to metric reports.

Services may:

- Terminate a subscription by sending a `SubscriptionTerminated` message from the Base Message Registry as the last event.
- Terminate a subscription if the number of delivery errors exceeds pre-configured thresholds.

To unsubscribe from the events associated with this subscription, the client or administrator shall perform an HTTP `DELETE` request to the subscription's resource URI.

Subsequent requests to subscription resources that have been terminated respond with the HTTP `404 Not Found` status code.

Some configurable properties define the behavior for all event subscriptions. For details, see the Redfish `EventService` schema.

**DEPRECATED:** Previous versions of the specification required services to reject subscription requests if `RegistryPrefixes` and `MessageIds` were in the same request.

### 12.1.3 Open an SSE connection

A service may support the `ServerSentEventUri` property in the `EventService` resource. If a client performs a `GET` request on the URI that the `ServerSentEventUri` contains, an SSE connection opens for the client. For details about this method, see the server-sent events [Event service](#) clause.

### 12.1.4 EventType-based eventing

**DEPRECATED:** `EventType` -based eventing is deprecated in the Redfish schema in favor of using `RegistryPrefix` and `ResourceType`.

---

#### DEPRECATED

[Table 37](#) describes the types of events that Redfish generates:

**Table 37 — EventType-based eventing**

Event	Occurs when	Description
Life cycle	Resources are created, modified, or destroyed.	Not every modification of a resource results in an event. This behavior is similar to when ETags are changed and implementations might not send an event for every resource change.
	Usually indicates that the resource and, optionally, its properties have changed.	For example, if an event is sent for every Ethernet packet that is received or each time that a sensor changes one degree, more events than fit in a scalable interface are generated.
Alert	An event of some significance happens.	Usually adopts a message registry approach similar to extended error handling in that a <code>MessageId</code> is included.
	Depending on the resource, may be generated directly or indirectly.	An example of an alert event is, a chassis is opened, a button is pushed, a cable is unplugged, or a threshold exceeded.  These events usually do not correspond well to life cycle-type events. Therefore, alerts have their own category.
Metric report	The telemetry service generates or updates a metric report.	Generated as specified by the <code>MetricReportDefinition</code> resources found <code>subordinate</code> to the telemetry service. Can occur periodically, on demand, or when changes are detected in the metric properties.  For details, see the Redfish <code>MetricReportDefinition</code> schema.

#### END DEPRECATED

---

### 12.1.5 Subscribing to events

[Table 38](#) describes the properties that a subscriber provides to subscribe to events and filter received messages:

**Table 38 — Subscription properties**

Property	Description
<code>MessageIds</code>	<p>An array of allowable values for <code>MessageId</code> in an event.</p> <p>An event is sent to the subscriber if the <code>MessageId</code> of the event is contained in this array or is found in a message registry referenced by <code>RegistryPrefixes</code>.</p> <p>To not perform inclusive filtering based upon <code>MessageId</code>, provide an empty array.</p> <p>The contents of the array should not include the major or minor version of the message registry. For example, instead of <code>Resource.1.2.ResourceCreated</code>, use <code>Resource.ResourceCreated</code>.</p>
<code>RegistryPrefixes</code>	<p>An array of standard or OEM message registries containing the allowable values for <code>MessageId</code> in an event.</p> <p>An event is sent to the subscriber if the <code>MessageId</code> is found in a message registry in this array or is contained in <code>MessageIds</code>.</p> <p>To not perform inclusive filtering based upon the message registry of the <code>MessageId</code>, provide an empty array.</p> <p>The contents of the array does not include the registry version. For example, instead of <code>Base.1.5.0</code>, use <code>Base</code>.</p>
<code>ResourceTypes</code>	<p>An array of standard or OEM resource types.</p> <p>An event is sent to the subscriber if the <code>OriginOfCondition</code> resource type matches one of the <code>ResourceTypes</code> values.</p> <p>To not perform filtering based upon the resource type of the <code>OriginOfCondition</code>, provide an empty array.</p> <p>The contents of the array does not include the schema version. For example, instead of <code>Task.v1_2_0.Task</code>, use <code>Task</code>.</p>
<code>OriginResources</code>	<p>An array of URIs to resources.</p> <p>An event is sent to the subscriber if the <code>OriginOfCondition</code> property matches one of the URIs listed in <code>OriginResources</code>. To include subordinate resources regardless of depth, set the <code>SubordinateResources</code> property to <code>true</code>.</p> <p>To not perform filtering based upon the URI of the <code>OriginOfCondition</code>, provide an empty array.</p>
<code>ExcludeMessageIds</code>	<p>An array of disallowed values for <code>MessageId</code> in an event.</p> <p>An event is not sent to the subscriber if the <code>MessageId</code> of the event is contained in this array.</p> <p>To not perform exclusive filtering based upon <code>MessageId</code>, provide an empty array.</p> <p>The contents of the array should not include the major or minor version of the message registry. For example, instead of <code>Resource.1.2.ResourceCreated</code>, use <code>Resource.ResourceCreated</code>.</p>

Property	Description
ExcludeRegistryPrefixes	An array of standard or OEM message registries containing the disallowed values for <code>MessageId</code> in an event.
	An event is not sent to the subscriber if the <code>MessageId</code> is found in a message registry in this array.
	To not perform exclusive filtering based upon the message registry of the <code>MessageId</code> , provide an empty array.
EventFormatType	The contents of the array does not include the registry version. For example, instead of <code>Base.1.5.0</code> , use <code>Base</code> .
	The format that can be sent by using the <code>EventFormatTypes</code> property in the event service.
	Represents the format of the payload sent to the event destination.
	If the subscriber omits this value, the payload corresponds to the <code>Event</code> schema.

Clients can read the `EventService` resource of a service in order to determine what control properties in the previous table are supported.

### 12.1.6 Event formats

[Table 39](#) describes the event formats:

**Table 39 — Event formats**

Event format	Description
Metric report message objects	Used when the telemetry service generates a new or updates an existing metric report. Metric report message objects sent to the specified client endpoint shall contain the properties, as described in the Redfish <code>MetricReport</code> schema.

Event format	Description
Event message objects	<p>Used for all other types of events. Event message objects <code>POST</code> ed to the specified client endpoint shall contain the properties as described in the Redfish <code>Event</code> schema. Supports a message registry. In a message registry approach, a message registry lists the <code>MessageIds</code> in a well-known format. These <code>MessageIds</code> are terse in nature and thus they are much smaller than actual messages, making them suitable for embedded environments.</p> <p>The registry also contains a message. The message itself can have arguments and default values for severity and recommended actions. The <code>MessageId</code> property follows the format defined in the <a href="#">MessageId format clause</a></p> <p>Event messages may also have an <code>EventGroupId</code> property, which lets clients know that different messages may be from the same event. For instance, if a LAN cable is disconnected, they may get a specific message from one registry about the LAN cable being disconnected, another message from a general registry about the resource changing, perhaps a message about resource state change, and maybe more. For the client to determine whether these have the same root cause, these messages have the same value for the <code>EventGroupId</code> property.</p>

### 12.1.7 OEM extensions

OEMs can extend both messages and message registries. Any individual message, per the `MessageRegistry` schema definition, define OEM sections. Thus, if OEMs wish to provide additional information or properties, use the OEM section.

OEMs shall not supply additional message arguments beyond those in a standard message registry. OEMs may substitute their own message registry for the standard registry to provide the OEM section within the registry but shall not change the standard values, such as messages, in such registries.

## 12.2 Asynchronous operations

Services that support asynchronous operations implement the `TaskService` and `Task` resources.

The `TaskService` resource describes the service that handles `task` operations. It contains a resource collection of zero or more `Task` resources. Each `Task` resource describes a long-running operation that is spawned when a request takes longer than a few seconds, such as when a service is instantiated.

The `Task` schema defines task structure, including the start time, end time, task state, task status, and zero or more task-associated messages.

Each task has a number of possible states, which are defined in the `Task` schema as the values for the `TaskState` property.

When a client issues a request that results in a long-running operation, the service returns the HTTP `202 Accepted`

status code and a `Location` header that contains a `task monitor` URI and, optionally, the `Retry-After` header that defines the amount of time that the client should wait before querying the status of the operation. The `202 Accepted` response should include a response body. If a response body is provided, it shall contain a representation of the `Task` resource that represents the state of the operation.

The `task monitor` is an opaque, service-generated URI provided to the client that initiated the request. To query the status of an operation and determine when the operation has been completed and whether it succeeded, the client performs a `GET` request on the task monitor URI from the `Location` header of the response from the initial HTTP operation. The client should not include the `application/http` MIME type in the `Accept` header.

The task monitor URI should be in the format:

```
/redfish/v1/TaskService/TaskMonitors/<TaskMonitorId>
```

where

- `<TaskMonitorId>` is an opaque identifier for the task monitor.

As long as the operation is in process, the service shall return the HTTP `202 Accepted` status code when the client performs a `GET` request on the task monitor URI.

If a service supports cancellation of a task, the `Allow` header shall contain `DELETE` for the task monitor. To cancel the operation, the client may perform a `DELETE` request on the task monitor URI. The service determines when to delete the associated `Task` resource. The client may also perform a `DELETE` request on the `Task` resource to cancel the operation. Deleting the `Task` resource may invalidate the associated task monitor. A subsequent `GET` request on the task monitor URI returns either the HTTP `410 Gone` or `404 Not Found` status code.

In the unlikely event that a `DELETE` of the task monitor or `Task` resource returns the HTTP `202 Accepted` status code, an additional task shall not be started and instead the client may monitor the existing `Task` resource for the status of the cancellation request. When the task finally completes cancellation, operations on the task monitor URI and `Task` resources shall return the HTTP `404 Not Found` status code.

After the operation has been completed, the service shall update the `TaskState` in the `Task` resource with the appropriate value. In addition, the task monitor shall return:

- The appropriate HTTP status code, such as but not limited to `200 OK` for most operations or `201 Created` for `POST` to create a resource.
- The headers and response body of the initial operation, as if it had completed synchronously.

If the initial operation fails, the response body shall contain an [error response](#).

If the operation has been completed and the service has already deleted the task, the service may return the HTTP `410 Gone` or `404 Not Found` status code. This situation can occur if the client waits too long to read the task monitor.

To continue to get status information, the client can use the [resource identifier](#) from the [202 Accepted](#) response to directly query the `Task` resource.

- Services that support asynchronous operations shall implement the `Task` resource.
- The response to an asynchronous operation shall return the HTTP [202 Accepted](#) status code and set the `Location` response header to the URI of a task monitor associated with the task. The response may also include the `Retry-After` header that defines the amount of time that the client should wait before polling for status. If a response body is provided, it shall contain a representation of the `Task` resource.
- `GET` requests to either the task monitor or `Task` resource shall return the current status of the operation without blocking.
- HTTP `GET`, `PUT`, and `PATCH` operations should always be synchronous.
- Clients shall be prepared to handle both synchronous and asynchronous responses for HTTP `GET`, `PUT`, `PATCH`, `POST`, and `DELETE` requests.
- Services shall persist pending tasks produced by client requests containing [@Redfish.OperationApplyTime](#) across service restarts, until the task begins execution.
- Tasks that are pending execution should include the [@Redfish.OperationApplyTime](#) property to indicate when the task will start. If the `@Redfish.OperationApplyTime` value is `AtMaintenanceWindowStart` or `InMaintenanceWindowOnReset`, the task should also include the `@Redfish.MaintenanceWindow` property.
  - Services shall reject modification requests to the `@Redfish.MaintenanceWindow` property in the `Task` resource.
  - Changing the maintenance window for a resource may not affect existing tasks.

## 12.3 Resource tree stability

---

The [resource tree](#), which is defined as the set of URIs and array elements within the implementation, should be consistent on a single service across device resets or power cycles, and should withstand a reasonable amount of configuration change, such as adding an adapter to a server.

The resource tree on one service might not be consistent across instances of devices. The client should traverse the data model and discover resources to interact with them.

Some resources might remain very stable from system to system, such as manager network settings. However, the architecture does not guarantee this stability.

- A resource tree should remain stable across service restarts and minor device configuration changes. Thus, the set of URIs and array element indexes should remain constant.
- A client shall not expect the resource tree to be consistent between instances of services.

## 12.4 Discovery

### 12.4.1 Discovery overview

Automatic discovery of managed devices supporting Redfish may be accomplished by using the Simple Service Discovery Protocol (SSDP). This protocol enables network-efficient discovery without resorting to ping-sweeps, router table searches, or restrictive DNS naming schemes. Use of SSDP is optional and, if implemented, shall enable the user to disable the protocol through the `ManagerNetworkProtocol` resource.

The objective of discovery is for client software to locate managed devices that conform to the *Redfish Specification*. Therefore, the primary SSDP functionality is incorporated in the M-SEARCH query. Redfish also follows the SSDP extensions and naming that UPnP uses, where applicable, so that systems that conform to the *Redfish Specification* can also implement UPnP without conflict.

### 12.4.2 UPnP compatibility

For compatibility with general-purpose SSDP client software, primarily UPnP, the service should use UDP port 1900 for all SSDP traffic. In addition, the Time-to-Live (TTL) hop count setting for SSDP multicast messages should default to 2.

### 12.4.3 USN format

The UUID in the USN field of the service shall equal the UUID property in the service root. If multiple or redundant managers exist, the UUID of the service shall remain static regardless of redundancy failover. The unique ID shall be in the canonical UUID format, followed by `::dmtf-org`.

### 12.4.4 M-SEARCH response

The Redfish service Search Target (ST) is defined as:

```
urn:dmtf-org:service:redfish-rest:1
```

The managed device shall respond to M-SEARCH queries for Search Target (ST) of the Redfish service, as well as `ssdp:all`. For UPnP compatibility, the managed device should respond to M-SEARCH queries for Search Target (ST) of `upnp:rootdevice`.

The URN provided in the `ST` header in the reply shall use the `redfish-rest:` service name followed by the major version of the *Redfish Specification*. If the minor version of the *Redfish Specification* to which the service conforms is a non-zero value, the service may append the minor version with a preceding colon ( : ).

**Note:** Previous versions of the specification required the minor version in responses, such as `redfish-rest:1:4` if the service conforms to *Redfish Specification v1.4*. However, this conflicts with UPnP, which requires the version in the response to match the version from the request. This is to allow for backwards compatible devices to present themselves as the requested version.

The managed device shall provide clients with the `AL` header that points to the Redfish service root URL.

For UPnP compatibility, the managed device should provide clients with the `Location` header that points to the UPnP XML descriptor.

The response to an M-SEARCH multicast or unicast query shall use the following format:

```
HTTP/1.1 200 OK
CACHE-CONTROL:max-age=<MaxAgeSeconds>
ST:urn:dmtf-org:service:redfish-rest:1
USN:uuid:<ServiceUUID>::urn:dmtf-org:service:redfish-rest:1
AL:<ServiceRootURI>
EXT:
```

where

- `<MaxAgeSeconds>` is the number of seconds caches can store the response and is at least `1800`.
- `<ServiceUUID>` is the UUID of the Redfish service, such as `92384634-2938-2342-8820-489239905423`.
- `<ServiceRootURI>` is the absolute URI of the Redfish service root, such as `https://192.168.1.50/redfish/v1/`.

A service may provide additional headers for UPnP compatibility.

### 12.4.5 Notify, alive, and shutdown messages

Redfish devices may implement the additional UPnP-defined SSDP messages to announce their availability to software. If implemented, services shall allow the end user to disable the traffic separately from the M-SEARCH response functionality. This capability enables users to use the discovery functionality with minimal amounts of generated network traffic.

## 12.5 Server-sent events

### 12.5.1 General

Server-sent events (SSE), defined by the Web Hypertext Application Technology Working Group (WHATWG), enables a client to open a connection with a web service. The web service can continuously push data to the client, as needed.

Successful resource responses for SSE shall:

- Return the HTTP `200 OK` status code.
- Have a `Content-Type` header set as `text/event-stream` or `text/event-stream; charset=utf-8`.

Unsuccessful resource responses for SSE shall:

- Return an HTTP `400` or greater status code.
- Have a `Content-Type` header set as `application/json` or `application/json; charset=utf-8`.
- Contain a JSON object in the response body, as described in [Error responses](#), which details the error or errors.

A service may occasionally send a comment within a stream to keep the connection alive. Services shall separate events with blank lines. Blank lines should be sent as part of the end of an event, otherwise dispatch may be delayed in conforming consumers.

The following clauses describe how Redfish uses SSE in different Redfish data model contexts. For details about SSE, see the [HTML5 Specification](#).

### 12.5.2 Event service

A service's implementation of the `EventService` resource may contain the `ServerSentEventUri` property. If a client performs a `GET` request on the URI specified by the `ServerSentEventUri` property, the service shall keep the connection open and conform to the [HTML5 Specification](#) until the client closes the socket. Service-generated events shall be sent to the client by using the open connection.

When a client opens an SSE stream for the event service, the service shall create an `EventDestination` resource in the `Subscriptions` collection for the event service to represent the connection. The `Context` property in the `EventDestination` resource shall be a service-generated opaque string.

The service shall delete the corresponding `EventDestination` resource when the connection is closed. The service shall close the connection if the corresponding `EventDestination` resource is deleted.

The service shall use the `id` field in the SSE stream to uniquely identify a payload in the SSE stream. The value of the `id` field is determined by the service. A service should accept the `Last-Event-ID` header from the client to allow a client to restart the event stream in case the connection is interrupted.

The service shall use the `data` field in the SSE stream based on the payload format. The SSE streams have these formats:

- [Metric report SSE stream](#). Services shall use this format when the telemetry service generates or updates a metric report.
- [Event message SSE stream](#). Services shall use this format for all other types of events.

To reduce the amount of data returned to the client, the service should support the `$filter` query parameter in the URI for the SSE stream.

**Note:** The `$filter` syntax shall follow the format in the [\\$filter query parameter](#) clause.

The service should support these properties as filter criteria:

- `EventFormatType`

The service sends events of the matching `EventFormatType`.

Example:

```
https://sseuri?$filter=EventFormatType eq 'Event'
```

Valid values are the `EventFormatType` enumerated string values that the Redfish `EventService` schema defines.

- `EventType`

The service sends events of the matching `EventType`.

Example:

```
https://sseuri?$filter=EventType eq 'StatusChange'
```

Valid values are the `EventType` enumerated string values that the Redfish `Event` schema defines.

- `MessageId`

The service sends events with the matching `MessageId`.

Example:

```
https://sseuri?$filter=MessageId eq 'Contoso.1.0.TempAssert'
```

- `MetricReportDefinition`

The service sends metric reports generated from the `MetricReportDefinition`.

Example:

```
https://sseuri?$filter=MetricReportDefinition eq '/redfish/v1/TelemetryService/MetricReportDefinitions/PowerMetrics'
```

- `OriginResource`

The service sends events for the resource.

Example:

```
https://sseuri?$filter=OriginResource eq '/redfish/v1/Chassis/1/Thermal'
```

- `RegistryPrefix`

The service sends events with messages that are part of the `RegistryPrefix`.

Example:

```
https://sseuri?$filter=(RegistryPrefix eq 'Resource') or (RegistryPrefix eq 'Task')
```

- `ResourceType`

The service sends events for resources that match the `ResourceType`.

Example:

```
https://sseuri?$filter=(ResourceType eq 'Power') or (ResourceType eq 'Thermal')
```

- `SubordinateResources`

When `SubordinateResources` is `true` and `OriginResource` is specified, the service sends events for the resource and its subordinate resources.

Example:

```
https://sseuri?$filter=(OriginResource eq '/redfish/v1/Systems/1') and (SubordinateResources eq true)
```

To allow a client to expand the resource referenced by the `OriginOfCondition` property in the event payload, the service should support the `includeoriginofcondition` query parameter in the URI for the SSE stream. For example:

```
https://sseuri?includeoriginofcondition
```

#### 12.5.2.1 Event message SSE stream

The service shall use the `data` field in the SSE stream to include the JSON representation of the [Event object](#).

The following example payload shows a stream that contains a single event with the `id` field set to `1`, and a `data` field that contains a single `Event` object.

```
id: 1
data:{ 
  data: "@odata.type": "#Event.v1_6_0.Event",
  data: "Id": "1",
  data: "Name": "Event Array",
  data: "Context": "ABCDEFGH",
  data: "Events": [
    data: {
      data: "MemberId": "1",
      data: "EventType": "Alert",
      data: "EventId": "1",
      data: "Severity": "Warning",
      data: "MessageSeverity": "Warning",
      data: "EventTimestamp": "2017-11-23T17:17:42-0600",
      data: "Message": "The LAN has been disconnected",
      data: "MessageId": "Alert.1.0.LanDisconnect",
      data: "MessageArgs": [
        data: "EthernetInterface 1",
        data: "/redfish/v1/Systems/1"
      ],
      data: "OriginOfCondition": {
        data: "@odata.id": "/redfish/v1/Systems/1/EthernetInterfaces/1"
      },
      data: "Context": "ABCDEFGH"
    }
  ]
}
```

#### 12.5.2.2 Metric report SSE stream

The service shall use the `data` field in the SSE stream to include the JSON representation of the [MetricReport object](#).

The following example payload shows a stream that contains a metric report with the `id` field set to `127`, and the `data` field containing the metric report object.

```
id: 127
data:{}
data: "@odata.id": "/redfish/v1/TelemetryService/MetricReports/AvgPlatformPowerUsage",
data: "@odata.type": "#MetricReport.v1_3_0.MetricReport",
data: "Id": "AvgPlatformPowerUsage",
data: "Name": "Average Platform Power Usage metric report",
data: "MetricReportDefinition": {
data:     "@odata.id": "/redfish/v1/TelemetryService/MetricReportDefinitions/AvgPlatformPowerUsage"
data: },
data: "MetricValues": [
data:   {
data:     "MetricId": "AverageConsumedWatts",
data:     "MetricValue": "100",
data:     "Timestamp": "2016-11-08T12:25:00-05:00",
data:     "MetricProperty": "/redfish/v1/Chassis/Tray_1/Power#/0/PowerConsumedWatts"
data:   },
data:   {
data:     "MetricId": "AverageConsumedWatts",
data:     "MetricValue": "94",
data:     "Timestamp": "2016-11-08T13:25:00-05:00",
data:     "MetricProperty": "/redfish/v1/Chassis/Tray_1/Power#/0/PowerConsumedWatts"
data:   },
data:   {
data:     "MetricId": "AverageConsumedWatts",
data:     "MetricValue": "100",
data:     "Timestamp": "2016-11-08T14:25:00-05:00",
data:     "MetricProperty": "/redfish/v1/Chassis/Tray_1/Power#/0/PowerConsumedWatts"
data:   }
data: ]
data:}
```

## 12.6 Update service

### 12.6.1 Overview

This clause covers the mechanism for software updates by using the update service.

### 12.6.2 Software update types

Clients can use these methods to update software through the update service:

- **Simple updates:** The service *pulls* the update from a client-indicated network location.
- **Multipart HTTP push updates:** The client uses HTTP or HTTPS with a multipart-formatted request body to *push*

a software image to the service.

#### 12.6.2.1 Simple updates

A service can support the `SimpleUpdate` action within the `UpdateService` resource. A client can perform a `POST` request on the action target URI to initiate a pull-based update, as defined by the `UpdateService` schema. After a successful `POST`, the service should return the HTTP `202 Accepted` status code with the `Location` header set to the URI of a *task monitor*. Clients can use this *task* to monitor the progress and results of the update, which includes the progress of image transfer to the service, as described in the [Asynchronous operations](#) clause.

#### 12.6.2.2 Multipart HTTP push updates

A service may support the `MultipartHttpPushUri` property within the `UpdateService` resource. A client can perform an HTTP or HTTPS `POST` request on the URI specified by this property to initiate a push-based update.

- Access to this URI shall require the same privilege as access to the update service.
- A client `POST` to this URI shall contain the `Content-Type` HTTP header with the value `multipart/form-data`, with the body formatted as defined by this specification. For more information about `multipart/form-data` HTTP requests, see [RFC7578](#).
- The client `POST` request shall contain the binary image as one of the parts in a `multipart/form-data` request body, as defined by [Table 40](#). In addition, the request shall include action parameters for the update in a JSON formatted part in the same `multipart/form-data` request body, as defined by [Table 40](#). If the request has no action parameters, an empty JSON object shall be used.
- A service may require the `Content-Length` HTTP header for `POST` requests to this URI. In this case, if a client does not include the required `Content-Length` header in the `POST` request, the service shall return the HTTP `411 Length Required` status code.
- A service should return the HTTP `413 Payload Too Large` status code if the size of the binary image is larger than the maximum image size that the service supports, as advertised in `MaxImageSizeBytes` property in the `UpdateService` resource.
- After a successful `POST` to this URI, the service shall return the HTTP `202 Accepted` status code with a `Location` header set to the URI of a *task monitor*. Clients can use this *task* to monitor the progress and results of the update, as described in the [Asynchronous operations](#) clause.
- Upon completion of the requested update, the service shall return the HTTP `200 OK` status code and an `error response`, with a message that indicates success or any additional relevant messages, or the HTTP `204 No Content` status code.
  - The service should return the HTTP `200 OK` status code. If the update was successfully processed and completed without errors, warnings, or other notifications for the client, the service should return the `UpdateSuccessful` message from the Update Message Registry in the `code` property in the response body.

[Table 40](#) describes the requirements of a `multipart/form-data` request body for an HTTP push software update. If the **Required** column contains **Yes**, the service shall require the client to provide the request body part. Services

shall not require clients to provide request body parts that are not marked as required. If the **Required** column contains **No**, the service may allow the client to provide the request body part.

**Table 40 — Multipart HTTP push updates**

Request body part	HTTP headers	Header value and parameters	Required	Description
Action parameters JSON part	Content-Disposition	form-data; name="UpdateParameters"	Yes	JSON-formatted part for passing the action parameters. The value of the <code>name</code> field shall be "UpdateParameters". The format of the JSON shall follow the definition of the <code>UpdateParameters</code> object in the <code>UpdateService</code> schema. Services may allow the inclusion of the <code>@Redfish.OperationApplyTime</code> property in the request body. See <a href="#">Operation apply time</a> .
	Content-Type	application/json; charset=utf-8 or application/json	Yes	Media type format and character set of this request part.
Update file binary part	Content-Disposition	form-data; name="UpdateFile"; filename=string	Yes	Binary file to use for this software update. The value of the <code>name</code> field shall be "UpdateFile". The value of the <code>filename</code> field should reflect the name of the file as loaded by the client.
	Content-Type	application/octet-stream	Yes	Media type format of the binary update file.
OEM specific parts	Content-Disposition	form-data; name="OemXXXX"	No	Optional OEM part. The value of the <code>name</code> field shall start with "Oem". Content-Type is optional and depends on the OEM part type.

This example shows a `multipart/form-data` request to push an update image:

```
POST /redfish/v1/UpdateService/upload HTTP/1.1
Host: <host-path>
Content-Type: multipart/form-data; boundary=-----d74496d66958873e
Content-Length: <computed-length>
Connection: keep-alive
X-Auth-Token: <session-auth-token>

-----d74496d66958873e
Content-Disposition: form-data; name="UpdateParameters"
Content-Type: application/json
```

```
{  
    "Targets": ["/redfish/v1/Managers/1"],  
    "@Redfish.OperationApplyTime": "OnReset",  
    "Oem": {}  
}  
  
-----d74496d66958873e  
Content-Disposition: form-data; name="UpdateFile"; filename="flash.bin"  
Content-Type: application/octet-stream  
  
<software image binary>
```

## 12.7 Outbound connections

### 12.7.1 Overview

There are cases where it may not be possible for a client to directly connect to a Redfish service. One common example is that the Redfish service is behind a firewall. Services may implement WebSocket-based outbound connections to establish a tunnel between the service and a remote client. If supported, services shall implement the `OutboundConnectionCollection` resource subordinate to the `AccountService` resource.

Clients may create new `OutboundConnection` resources to create WebSocket connections to remote clients. Remote clients initiate Redfish requests to the service by encoding Redfish HTTP requests in the WebSocket payload over this connection.

### 12.7.2 Establishing an outbound connection

The service shall establish a persistent [RFC6455](#)-defined WebSocket connection for each `OutboundConnection` resource where `ConnectionEnabled` contains `true` with the configured authentication mechanism to the address specified by the `EndpointURI` property. The service may apply proxy settings to the connection based on the configuration of the `ManagerNetworkProtocol` resource.

The service shall include the HTTP headers specified in the `PreUpgradeHTTPHeaders` property in the initial HTTPS request to establish the WebSocket, prior to the upgrade to a WebSocket. The service shall include the `Sec-WebSocket-Protocol` request header and it shall contain the value `Redfish`. The following example shows an initial HTTPS request from a service prior to the upgrade to a WebSocket:

```
GET /redfish HTTP/1.1  
Host: clientapp.contoso.com:443  
Upgrade: websocket  
Connect: Upgrade  
Sec-WebSocket-Protocol: Redfish
```

```
Sec-WebSocket-Key: abcdefhijklmnop==  
Sec-WebSocket-Version: 13  
PreUpgradeHeader1: PreUpgradeHeader1Value
```

When a connection is established with the remote client, the service shall create a `Session` resource with the role defined in the `OutboundConnection` resource.

If the connection terminates for any reason, any pending Redfish responses are dropped. Services shall re-establish outbound connections:

- Based on the `RetryPolicy` configuration of the `OutboundConnection` resource.
- If the `ConnectionEnabled` property in the `OutboundConnection` resource is set to `true`.

### 12.7.3 MTLS for an outbound connection

MTLS is the preferred authentication mechanism to establish a secure connection. This is specified when the `Authentication` property in the `OutboundConnection` resource contains `MTLS`. There are two sets of certificates in the `OutboundConnection` resource: `Certificates` and `ClientCertificates`.

The certificates found in the collection referenced by the `Certificates` property are the certificates for the remote client's WebSocket service endpoint, typically an HTTPS server with WebSockets enabled. Services shall verify the identity of the WebSocket service during TLS handshaking with the contents of this collection. Services may perform additional verification based on other factors, such as the configuration of the `SecurityPolicy` resource.

The certificates found in the collection referenced by the `ClientCertificate` property are the certificates the service provides to the remote client's WebSocket service during TLS handshaking. In some cases, the service may provide its IDEV-ID or LDEV-ID certificate. To verify the identity Redfish service during TLS handshaking, the remote client hosting the WebSocket will need copies of these certificates.

### 12.7.4 Handling Redfish requests over an outbound connection

Section 5.2 of [RFC6455](#) contains the WebSocket packet definition.

To send a Redfish request, the remote client shall provide the entire HTTP request, including the method, path, version, and headers, in the "Payload data" portion of the WebSocket packet over the established connection. The remote client shall specify either the text ( `0x1` ) opcode or binary ( `0x2` ) opcode in the request. The request should specify the binary ( `0x2` ) opcode if the payload data is binary, such as a firmware image.

To send a Redfish response, the service shall provide the entire HTTP response, including the version, status code, status message, and headers, in the "Payload data" portion of the WebSocket packet over the established connection. The service shall provide the same opcode received in the request from the external client.

Services shall respond to ping ( `0x9` ) opcodes with a pong ( `0xA` ) opcode. Services shall send the ping ( `0x9` ) opcode to the remote client at the interval specified by the `WebSocketPingIntervalMinutes` property.

The following example shows the contents of "Payload data" for a remote client performing an HTTP `GET` request on the `ServiceRoot` resource:

```
GET /redfish/v1/ HTTP/1.1
Host: clientapp.contoso.com:443
Accept: application/json
Content-Type: application/json
```

The following example shows the contents of "Payload data" for a remote client performing an HTTP `POST` request on the `Reset` action for a `ComputerSystem` resource:

```
POST /redfish/v1/Systems/1/Actions/ComputerSystem.Reset HTTP/1.1
Host: clientapp.contoso.com:443
Accept: application/json
Content-Type: application/json
Content-Length: 27

{
    "ResetType": "On"
}
```

## 12.7.5 Closing an outbound connection

The service shall close the WebSocket connection to the remote client if:

- The `Session` resource representing the open connection to the remote client is deleted. The service shall also set the `ConnectionEnabled` property in the `OutboundConnection` resource to `false`.
- The `OutboundConnection` resource representing the outbound connection to the remote client is deleted. The service shall also delete the respective `Session` resource.
- The `ConnectionEnabled` property in the `OutboundConnection` resource is set to `false`. The service shall also delete the respective `Session` resource.
- The close ( `0x8` ) is received over the connection. The service shall also set the `ConnectionEnabled` property in the `OutboundConnection` resource to `false` and delete the respective `Session` resource.

# 13 Security details

## 13.1 Transport Layer Security (TLS) protocol

### 13.1.1 Transport Layer Security (TLS) protocol overview

Implementations shall support the Transport Layer Security (TLS) protocol v1.2 with [RFC7525](#) recommendations or later. Implementations may remove support for older versions for TLS in favor of newer versions.

**DEPRECATED:** Previous versions of this specification allowed for TLS v1.1.

Implementations should support:

- The [Storage Networking Industry Association \(SNIA\) TLS Specification for Storage Systems](#).
- The latest version of the TLS v1.x specification.

### 13.1.2 Cipher suites

Implementations shall only support cipher suites listed as "Recommended" in the **TLS Cipher Suites** table defined by the [IANA TLS Parameters registry](#).

Cipher suites that are listed as mandatory in various RFCs, but are not "Recommended" in the **TLS Cipher Suites** table defined by the [IANA TLS Parameters registry](#), shall not be supported.

Implementations should consider the support of pre-shared key ciphers suites listed as "Recommended" in the **TLS Cipher Suites** table defined by the [IANA TLS Parameters registry](#), which enable authentication and identification without trusted certificates.

---

#### DEPRECATED

Implementations should support AES-256-based ciphers from the TLS suites.

Redfish implementations should consider the support of ciphers, such as the following ciphers, which enable authentication and identification without trusted certificates:

```
TLS_PSK_WITH_AES_256_GCM_SHA384  
TLS_DHE_PSK_WITH_AES_256_GCM_SHA384  
TLS_RSA_PSK_WITH_AES_256_GCM_SHA384
```

The advantage of these recommended ciphers is:

AES-GCM is not only efficient and secure, but hardware implementations can achieve high speeds with low cost and low latency because the mode can be pipelined.

Additionally, Redfish implementations should support the following cipher:

TLS\_RSA\_WITH\_AES\_128\_CBC\_SHA

For more information, see [RFC5487](#) and [RFC5288](#).

**END DEPRECATED**

---

### 13.1.3 Certificates

Redfish implementations shall support replacement of the default certificate if one is provided.

Redfish implementations shall use certificates that conform to X.509-v3, as defined in [RFC5280](#).

## 13.2 Sensitive data

---

Operations that contain sensitive data should use HTTPS only. For example, a `SimpleUpdate` action with a username and password should use HTTPS to protect the sensitive data.

Properties in service responses that represent sensitive data, such as passwords, shall be `null`.

Responses from URIs where the URI itself contains sensitive data in a URI segment may return the HTTP [404 Not Found](#) status code instead of the HTTP [401 Unauthorized](#) status code, the HTTP [403 Forbidden](#) status code, or the HTTP [405 Method Not Allowed](#) status code to prevent attackers from obtaining the sensitive data in the URI.

## 13.3 Authentication

---

### 13.3.1 Authentication overview

Services:

- Shall support both [HTTP Basic authentication](#) and [Redfish session login authentication](#).
- Shall use only connections that conform to TLS to transport the data between any third-party authentication service and clients.

- Shall not require a client that uses HTTP Basic authentication to create a session.
- May implement other authentication mechanisms.

### 13.3.2 Authentication requirements

#### 13.3.2.1 Resource and operation authentication requirements

Services shall authenticate all write requests to Redfish resources. For example:

- `POST`, except to the `Sessions` resource collection for authentication
- `PUT`
- `PATCH`
- `DELETE`

Redfish resources shall not be available as unauthenticated, except for:

- The service root to identify the device and service locations.
- The [Redfish metadata document](#) to get resource types.
- The [OData service document](#) for compatibility with OData clients.
- The Redfish [OpenAPI YAML document](#) for compatibility with OpenAPI clients.
- The `version` object at `/redfish`.

Services may reject requests to the previous resources if invalid credentials are provided by the client.

**Note:** This specification does not cover external services that are linked through external references. These services may have other security requirements.

#### 13.3.2.2 HTTP header authentication requirements

An authentication header shall accompany every request that establishes a secure channel.

Services:

- Shall process HTTP headers for authentication before other headers that may affect the response. For example, `ETag`, `If-Match`, and so on.
- Shall not use HTTP cookies to authenticate any activity, such as `GET`, `POST`, `PUT`, `PATCH`, and `DELETE`.

#### 13.3.2.3 Authentication failure requirements

When authentication fails, extended error messages shall not provide privileged information.

### 13.3.3 HTTP Basic authentication

Services shall support HTTP Basic authentication, as defined by [RFC7617](#), and shall use only connections that conform to TLS to transport the data between any third-party authentication service and clients.

All requests that use HTTP Basic authentication shall require HTTPS.

When multi-factor authentication is enabled, services shall reject HTTP Basic authentication for accounts that are not configured to bypass multi-factor authentication. [Session-based authentication](#) is required in this case.

**Note:** The IETF has highlighted security concerns with HTTP Basic authentication. While HTTPS is required for the usage of HTTP Basic authentication, there are other concerns implementers need to be aware of that [RFC7617](#) documents. This functionality can be restricted or disabled with the `HTTPBasicAuth` property in the `AccountService` resource.

### 13.3.4 Redfish session login authentication

Service shall provide login sessions that conform with this specification.

Session management is determined by the implementation of the Redfish service, which includes orphaned session timeout and the management of the number of simultaneous open sessions.

#### 13.3.4.1 Redfish login sessions

For improved performance and security, a client should use the session management interface to create a Redfish login session. The session service specifies the URI for session management.

To establish a session, find the URI in either:

- The session service's `Sessions` property.
- The service root's [links property](#) under the `Sessions` property.

Both URIs shall be the same.

```
{  
  "SessionService": {  
    "@odata.id": "/redfish/v1/SessionService"  
  },  
  "Links": {  
    "Sessions": {  
      "@odata.id": "/redfish/v1/SessionService/Sessions"  
    }  
  },  
}
```

```
    ...  
}
```

#### 13.3.4.2 Session login

To create a Redfish session without an authentication header, perform an HTTP `POST` request on the session service's `Sessions` resource collection. The `POST` to create a session shall only be supported with HTTPS. If both HTTP and HTTPS are enabled, a `POST` request to create a session through the HTTP port should redirect to the HTTPS port. Include the following `POST` body:

```
POST /redfish/v1/SessionService/Sessions HTTP/1.1  
Host: <host-path>  
Content-Type: application/json; charset=utf-8  
Content-Length: <computed-length>  
Accept: application/json; charset=utf-8  
OData-Version: 4.0  
  
{  
    "UserName": "<username>",  
    "Password": "<password>"  
}
```

Fields in brackets are placeholders for client-specific values.

When a multi-factor authentication type that requires tokens is enabled, services shall require the `Token` property in the `POST` request to the `SessionCollection` resource for accounts that are not configured to bypass multi-factor authentication unless Time-based One-Time Password multi-factor authentication is enabled and the user account does not have a secret key configured. Additional semantics for this condition are described in the [Time-based One-Time Password secret key handling](#) clause. The service shall verify the provided token in addition to verifying the username and password. If the `Token` property is required and not provided by the client, but the client provided a valid `UserName` and `Password` combination, the service shall return the HTTP `401 Unauthorized` status code and an `error response` with the `AuthenticationTokenRequired` message from the Base Message Registry. In addition, if the multi-factor authentication type uses a service-generated one-time passcode, the service shall also return the `OneTimePasscodeSent` message from the Base Message Registry, and send a one-time passcode to the configured delivery address for that account.

To verify that the request has been initiated from an authorized client domain, services should save the `Origin` header in reference to this session creation and compare it to subsequent requests using this session.

The response to the `POST` request to create a session shall include:

- `X-Auth-Token` header. Contains a session authentication token that the client can use in subsequent requests.

- [Location header](#). Contains a hyperlink to the new `Session` resource.
- JSON response body. Contains the full representation of the new `Session` resource.

The following sample response shows a newly created session:

```
HTTP/1.1 201 Created
Location: /redfish/v1/SessionService/Sessions/1
X-Auth-Token: <session-auth-token>

{
    "@odata.id": "/redfish/v1/SessionService/Sessions/1",
    "@odata.type": "#Session.v1_0_0.Session",
    "Id": "1",
    "Name": "User Session",
    "Description": "User Session",
    "UserName": "<username>",
    "Password": null
}
```

The client that sends the session login request should save the session authentication token from the `X-Auth-Token` header and the contents of the [Location header](#) from the response of the login `POST` request.

To authenticate subsequent requests, the client sets the `X-Auth-Token` header to the session authentication token that the `POST` login request returns.

**Note:** The session ID differs from the session authentication token, as follows:

- **Session ID:** The session ID uniquely identifies the `Session` resource. The response data with the last segment of the [Location header](#) URI returns is the session ID. To view active sessions and terminate any session, an administrator with sufficient privileges can use the session ID.
- **Session authentication token:** Only the client that executes the login has the session authentication token.

#### 13.3.4.3 Session lifetime

Unlike some token-based methods that use token expiration times, Redfish sessions time out. As long as a client continues to send requests more frequently than the session timeout period, the session remains open and the session authentication token remains valid. If the session times out, it is automatically terminated.

#### 13.3.4.4 Session termination or logout

When the client logs out, the Redfish session terminates. The session terminates through a `DELETE` request to the `Session` resource defined in either the [Location header](#) URI or the session ID in the response data.

This ability to `DELETE` a session through the `Session` resource enables an administrator with sufficient privileges to terminate other users' sessions from a different session.

When a session is terminated, the service shall not affect independent connections established originally by this session for other purposes, such as connections for [server-sent events](#) or transferring an image for the [update service](#).

### 13.3.5 Client certificate authentication

If client certificate authentication is enabled, the service shall send a client certificate request during the Transport Layer Security (TLS) handshake. When the service obtains a client certificate during the TLS handshake, the service shall verify the certificate with the certificates in `CertificateCollection` resource referenced by the `ClientCertificate` property within the `MFA` property of the `AccountService` resource. The service shall check for certificate revocation before processing the request with any configured Online Certificate Status Protocol (OCSP) servers.

The `RespondToUnauthenticatedClients` property within the `ClientCertificate` property within the `MFA` property of the `AccountService` resource controls the response behavior when an invalid certificate is provided by the client.

- If the property contains `true` or is not supported by the service, the service shall not fail the TLS handshake. This is to allow the service to send error messages or unauthenticated resources to the client.
- If the property contains `false`, the service shall fail the TLS handshake.

## 13.4 Authorization

---

### 13.4.1 Authorization overview

The Redfish authorization subsystem controls which users have access to resources and the type of access that users have. It consists of two parts: the [privilege model](#) and the [operation-to-privilege mapping](#).

The privilege model maps users to roles and maps roles to privileges. A privilege is a permission to complete an operation, such as read or write, within a defined management domain. For example, the `ConfigureUsers` privilege allows adding a user. A user is authorized to access a resource if they have the privileges required for that resource. The operation-to-privilege mapping defines which privileges are required to access any given operation.

Redfish allows vendors to extend the standard privilege model with OEM privileges and custom OEM roles. OEM privileges and custom roles participate in the privilege model the same as Redfish standard privileges and roles. Services may also allow clients to create custom roles. Restricted roles and restricted privileges allow vendors to further refine their authority model.

Services shall enforce the same privilege model for ETag-related activity as is enforced for the data being represented by the ETag. For example, the privilege required to read an ETag shall be the same as the privilege to read the data item that the ETag represents.

### 13.4.2 Privilege model

Each user shall be assigned exactly one role with the `RoleId` property in the `ManagerAccount` resource. The value of the `RoleId` property identifies a `Role` resource in the `RoleCollection` resource, where a role defines a set of privileges. A role shall be assigned to a user when a manager account is created. The client shall provide the `RoleId` property when creating a manager account to select one of the standard or custom roles.

Services shall provide information about all roles through the `RoleCollection` resource. The `AssignedPrivileges` and `OemPrivileges` arrays in the `Role` resource define a set of assigned privileges for the associated role. Two roles with the same privileges shall behave equivalently.

#### 13.4.2.1 Roles

Redfish defines a set of standard roles, allows a service to define custom OEM roles, and allows client-defined custom roles.

A service shall support all of the standard roles in [Table 41](#). The value of the `Id` and `AssignedPrivileges` properties in the `Role` resource for the standard roles shall contain the **Role name** and **Assigned privileges** column values, respectively. The `AssignedPrivileges` property for standard roles shall not be modifiable. The `IsPredefined` property for standard roles shall contain the value `true`.

[Table 41](#) describes the standard roles:

**Table 41 — Required standard roles**

Role name	Assigned privileges
Administrator	<code>Login</code> , <code>ConfigureManager</code> , <code>ConfigureUsers</code> , <code>ConfigureComponents</code> , <code>ConfigureSelf</code>
Operator	<code>Login</code> , <code>ConfigureComponents</code> , <code>ConfigureSelf</code>
ReadOnly	<code>Login</code> , <code>ConfigureSelf</code>

A service may support one or more of the standard roles in [Table 42](#). The value of the `Id` and `AssignedPrivileges` properties in the `Role` resource for the standard roles shall contain the **Role name** and **Assigned privileges** column values, respectively. The `AssignedPrivileges` property for standard roles shall not be modifiable. The `IsPredefined` property for standard roles shall contain the value `true`.

[Table 42](#) describes the optional standard roles:

**Table 42 — Optional standard roles**

Role name	Assigned privileges	Description
SystemAdministrator	Login , ConfigureSelf , AdministrateSystems	Administrator for systems found in the systems collection. Able to manage boot configuration, keys, and certificates for systems.
SystemOperator	Login , ConfigureSelf , OperateSystems	Operator for systems found in the systems collection. Able to perform resets and configure interfaces.
StorageAdministrator	Login , ConfigureSelf , AdministrateStorage	Administrator for storage subsystems and storage systems found in the storage collection and storage system collection respectively.
StorageBackupOperator	Login , ConfigureSelf , OperateStorageBackup	Operator for storage backup functionality for storage subsystems and storage systems found in the storage collection and storage system collection respectively.

A service may define custom OEM roles. The `IsPredefined` property for OEM roles shall contain the value `true`. A service shall not allow users to modify predefined OEM roles. OEM role names should begin with a lowercase character or "Oem" followed by a vendor name to avoid conflict with future Redfish predefined role names.

A service may allow custom client-defined roles to be created, modified, and deleted. If allowed, a user can perform a `POST` request on the `RoleCollection` resource to create a role, indicating privileges in the `AssignedPrivileges` and `OemPrivileges` properties in the `Role` resource. A service may restrict which privileges are allowed. The `IsPredefined` property for client-defined roles shall contain the value `false`. A service shall not allow a client-defined role to be deleted while it is in use, for example, when it is assigned to a local user or an LDAP `RemoteRoleMapping` property.

The value of the `RoleId` property shall be unique across all roles within the `RoleCollection` resource.

Non-Redfish services, such as those enabled by the `AccountTypes` property within the `ManagerAccount` resource, should map the Redfish `RoleId` to their permission system. For example, an SSH user with `Administrator` as the value of the `RoleId` property could map to "root" for the SSH service. However, the privileges specified by the `AssignedPrivileges` and `OemPrivileges` do not necessarily map to non-Redfish services.

### 13.4.2.2 Restricted roles and restricted privileges

Restricted roles and restricted privileges are intended to prevent privilege escalation. Restricted roles and restricted privileges are not less functional, but their usage is restricted to particular users. For example, to have a security administrator have privileges that the administrator does not have, you need to ensure the administrator cannot escalate to the security administrator role. An implementation can help achieve this by restricting the `Administrator` role and providing an alternate administrator role that lacks the security privilege.

A service may restrict any role. The `Restricted` property for restricted roles shall contain the value `true`. When a standard role is restricted, services shall provide the `AlternateRoleId` property to reference a non-restricted custom

role intended for clients to use as an alternate. Services may pre-define or create accounts that are configured with a restricted role.

Services shall not allow:

- A `RoleId` value for a restricted role to be specified when creating or modifying a `ManagerAccount` resource. This ensures administrators cannot create an account for themselves that has a restricted role.
- Modification of `ManagerAccount` resources with a `RoleId` property containing a value for a restricted role, with the exception of the `Enabled` property. This ensures administrators cannot gain access to another account.
- Deletion of `ManagerAccount` resources with a `RoleId` property containing a value for a restricted role.
- A restricted role to be specified in the `LocalRole` property within the `RemoteRoleMapping` property within the `AccountService` and `ExternalAccountProvider` resources.

A service may restrict any privilege, including standard and OEM privileges. The `RestrictedPrivileges` and `RestrictedOemPrivileges` properties in the `AccountService` resource shall specify the restricted privileges. Services shall not allow custom roles to specify restricted privileges. Services may contain predefined roles that are configured with restricted privileges.

#### 13.4.2.3 OEM privileges

OEM privileges allow a service to extend the privilege model by adding additional privileges to have additional control of what operations are allowed. It can be used when a standard privilege is overly broad.

A service may define OEM privileges and may include OEM privileges in any predefined role, including standard and custom OEM roles. The `OemPrivileges` property within the `Role` resource shall contain the OEM privileges that are assigned to the role. The `OemPrivileges` property in the `Role` resource for the predefined roles shall not be modifiable.

A service may allow OEM privileges to be assigned to client-defined roles.

#### 13.4.3 Redfish service operation-to-privilege mapping

For every request that a client makes to a service, the service shall determine that the authenticated identity of the requester has the authorization to complete the requested operation on the resource in the request.

Using the role and privileges authorization model where an authenticated identity context is assigned a role and a role is a set of privileges, the service typically checks an HTTP request against a mapping of the authenticated requesting identity role and privileges to determine whether the identity privileges are sufficient to complete the operation in the request.

A service may perform additional checks based on the identity of the user and remove data from responses. For example, a service might restrict access for non-administrative users to only access their own `ManagerAccount`, `Session`, and `EventDestination` resources.

### 13.4.3.1 Why specify operation-to-privilege mapping?

Initial versions of the *Redfish Specifications* defined several role-to-privilege mappings for standardized roles and normatively identified several privilege labels but did not normatively detail what these privileges or how privilege-to-operations mappings could be specified or represented in a normative fashion.

The lack of a methodology to define which privileges are required to complete a requested operation against the URI in the request puts at risk the interoperability between service implementations that clients may encounter due to variances in privilege requirements between implementations.

Also, a lack of methodology for specifying and representing the operation-to-privilege mapping prevents the Redfish Forum or other governing organizations from normatively defining privilege requirements for a service.

### 13.4.3.2 Representing operation-to-privilege mappings

A service should provide a Privilege Registry in the registry collection. This registry represents the privileges required to complete HTTP operations against resources supported by the service.

The Privilege Registry is a JSON document that contains a `Mappings` array of where an individual entry exists for every resource type that the service supports.

The operation-to-privilege mapping is defined for every resource type and applies to every resource the service implements for the applicable resource type.

In several situations, specific resources or properties may have differing operation-to-privilege mappings than the resource type-level mappings. In these cases, the resource type-level mappings need to be overridden. The `PrivilegeRegistry` schema defines the methodology for resource type-level operation-to-privilege mappings and related overrides.

If a service provides a Privilege Registry, the service shall use the Redfish Forum's Privilege Registry definition as a base operation-to-privilege mapping definition for operations that the service supports to promote interoperability for Redfish clients.

### 13.4.3.3 Operation map syntax

An operation map defines the set of privileges required to complete an operation on a resource-type.

The mapped operations are `GET`, `PUT`, `PATCH`, `POST`, `DELETE`, and `HEAD`. A privilege mapping is defined for each operation, irrespective of whether the service or data model supports the operation on the resource-type.

The privilege labels may be the Redfish standardized labels that the `PrivilegeType` enumeration in the `Privileges` schema defines and they may be OEM-defined privilege labels. The required privileges for an operation are specified using logical AND and OR behavior. For more information, see the [Privilege AND and OR syntax](#) clause.

The following example defines the privileges required for various operations on the `Manager` resource. Unless the implementation defines mapping overrides to the `OperationMap` array, the specified operation-to-privilege mapping represents behavior for all `Manager` resources in a service implementation.

```
{
  "Entity": "Manager",
  "OperationMap": {
    "GET": [
      "Privilege": ["Login"]
    ],
    "HEAD": [
      "Privilege": ["Login"]
    ],
    "PATCH": [
      "Privilege": ["ConfigureManager"]
    ],
    "POST": [
      "Privilege": ["ConfigureManager"]
    ],
    "PUT": [
      "Privilege": ["ConfigureManager"]
    ],
    "DELETE": [
      "Privilege": ["ConfigureManager"]
    ]
  }
}
```

#### 13.4.3.4 Mapping overrides syntax

[Table 43](#) describes the operation-to-privilege mapping, which varies from the resource type-level mapping:

**Table 43 — Mapping overrides syntax**

Situation	Description
Property override	Property has different privilege requirements than the resource in which it resides. For example, the <code>Password</code> property in the <code>ManagerAccount</code> resource requires the <code>ConfigureSelf</code> or <code>ConfigureUsers</code> privilege to change, in contrast to the <code>ConfigureUsers</code> privilege required for the other properties in <code>ManagerAccount</code> resources. If multiple properties with the same name are present in a resource, the property override applies to all property instances.
Subordinate override	Resource is used in context of another resource and the contextual privileges need to govern. For example, the privileges for <code>PATCH</code> operations on <code>EthernetInterface</code> resources depend on whether the resource is <code>subordinate</code> to the <code>Manager</code> resource, where <code>ConfigureManager</code> is required, or the <code>ComputerSystem</code> resource, where <code>ConfigureComponents</code> is required.

Situation	Description
Resource URI override	Resource instance has different privilege requirements for an operation than those defined for the resource type.

The overrides are defined in the context of the operation-to-privilege mapping for a resource type.

If multiple overrides are specified for a single resource type, the following precedence should be used for determining the appropriate override to apply:

- Property override
- Resource URI override
- Subordinate override

#### 13.4.3.5 Property override example

In the following example, the `Password` property on the `ManagerAccount` resource requires the `ConfigureSelf` or `ConfigureUsers` privilege to change, in contrast to the `ConfigureUsers` privilege required for the other properties in `ManagerAccount` resources:

```
{
  "Entity": "ManagerAccount",
  "OperationMap": {
    "GET": [
      {
        "Privilege": ["ConfigureManager"]
      },
      {
        "Privilege": ["ConfigureUsers"]
      },
      {
        "Privilege": ["ConfigureSelf"]
      }
    ],
    "HEAD": [
      {
        "Privilege": ["Login"]
      }
    ],
    "PATCH": [
      {
        "Privilege": ["ConfigureUsers"]
      }
    ],
    "POST": [
      {
        "Privilege": ["ConfigureUsers"]
      }
    ],
    "PUT": [
      {
        "Privilege": ["ConfigureUsers"]
      }
    ],
    "DELETE": [
      {
        "Privilege": ["ConfigureUsers"]
      }
    ]
  }
}
```

```

"PropertyOverrides": [
    "Targets": ["Password"],
    "OperationMap": {
        "PATCH": [
            {
                "Privilege": ["ConfigureUsers"]
            },
            {
                "Privilege": ["ConfigureSelf"]
            }
        ]
    }
]
}

```

#### 13.4.3.6 Subordinate override

The `Targets` property in `SubordinateOverrides` lists a hierarchical representation for when to apply the override. In the following example, the override for an `EthernetInterface` resource is applied when it is subordinate to an `EthernetInterfaceCollection` resource, which in turn is subordinate to a `Manager` resource. If a client were to `PATCH` an `EthernetInterface` resource that matches this override condition, it requires the `ConfigureManager` privilege. Otherwise, the client requires the `ConfigureComponents` privilege.

```

{
    "Entity": "EthernetInterface",
    "OperationMap": {
        "GET": [
            {
                "Privilege": ["Login"]
            }
        ],
        "HEAD": [
            {
                "Privilege": ["Login"]
            }
        ],
        "PATCH": [
            {
                "Privilege": ["ConfigureComponents"]
            }
        ],
        "POST": [
            {
                "Privilege": ["ConfigureComponents"]
            }
        ],
        "PUT": [
            {
                "Privilege": ["ConfigureComponents"]
            }
        ],
        "DELETE": [
            {
                "Privilege": ["ConfigureComponents"]
            }
        ]
    },
    "SubordinateOverrides": [
        {
            "Targets": ["Manager", "EthernetInterfaceCollection"],
            "OperationMap": {
                "PATCH": [

```

```
        "Privilege": ["ConfigureManager"]
    }],
    "POST": [
        "Privilege": ["ConfigureManager"]
    ],
    "PUT": [
        "Privilege": ["ConfigureManager"]
    ],
    "DELETE": [
        "Privilege": ["ConfigureManager"]
    ]
}
}
```

#### **13.4.3.7 Resource URI override**

The following example demonstrates the resource URI override syntax to define operation privilege variations for resource URIs.

The example defines both `ConfigureComponents` and `OEMAdminPriv` privileges as required to make a `PATCH` operation on the two resource URIs listed as targets.

```
{  
    "Entity": "ComputerSystem",  
    "OperationMap": {  
        "GET": [{  
            "Privilege": ["Login"]  
        }],  
        "HEAD": [{  
            "Privilege": ["Login"]  
        }],  
        "PATCH": [{  
            "Privilege": ["ConfigureComponents"]  
        }],  
        "POST": [{  
            "Privilege": ["ConfigureComponents"]  
        }],  
        "PUT": [{  
            "Privilege": ["ConfigureComponents"]  
        }],  
        "DELETE": [{  
            "Privilege": ["ConfigureComponents"]  
        }]  
    },  
    "ResourceURIOVERRIDES": [  
        {  
            "ResourceURI": "ComputerSystem/ConfigureComponents",  
            "ResourceType": "PUT",  
            "Privilege": "ConfigureComponents",  
            "OperationMap": {  
                "PUT": [{  
                    "Privilege": ["ConfigureComponents"]  
                }]  
            }  
        }]  
    ]  
}
```

```
"Targets": ["/redfish/v1/Systems/VM6", "/redfish/v1/Systems/Sys1"],  
"OperationMap": {  
    "GET": [{  
        "Privilege": ["Login"]  
    }],  
    "PATCH": [{  
        "Privilege": ["ConfigureComponents", "OEMSysAdminPriv"]  
    }]  
}  
}
```

#### 13.4.3.8 Privilege AND and OR syntax

The array placement of the privilege labels in the `OperationMap` `GET`, `HEAD`, `PATCH`, `POST`, `PUT`, and `DELETE` operation element arrays define the logical combinations of privileges that are required to call an operation on a resource or property.

For OR logical combinations, the privilege label appears in the operation element array as individual elements.

The following example defines either `Login` or `OEMPrivilege1` privileges that are required to perform a `GET` request.

```
{  
    "GET": [{  
        "Privilege": ["Login"]  
    }, {  
        "Privilege": ["OEMPrivilege1"]  
    }]  
}
```

For logical AND combinations, the privilege label appears in the `Privilege` property array in the operation element.

The following example defines both `ConfigureComponents` and `OEMSysAdminPriv` that are required to perform a `PATCH` request.

```
{  
    "PATCH": [{  
        "Privilege": ["ConfigureComponents", "OEMSysAdminPriv"]  
    }]  
}
```

### 13.4.4 Delegated authorization with OAuth 2.0

Services may support the [RFC6749](#)-defined OAuth 2.0 authorization framework.

#### 13.4.4.1 OAuth 2.0 overview

The OAuth 2.0 authorization framework allows a client to obtain access to a *resource server* from a *resource owner* and an *authorization server*.

Clients request access from a *resource owner* and is given an authorization grant. The authorization grant is then provided to the *authorization server* and an access token is provided to the client. The client provides the access token to the *resource server* in order to access a protected resource.

A Redfish service is considered to be a *resource server* in the OAuth 2.0 authorization framework.

#### 13.4.4.2 OAuth 2.0 data model requirements

Services that support OAuth 2.0:

- Shall support the `OAuth2` property in the `AccountService` resource.
- May support additional OAuth 2.0 servers with `ExternalAccountProvider` resources.

#### 13.4.4.3 OAuth 2.0 access tokens

Access tokens are the credentials the client provides to a service to access a protected resource. Clients provide the access token to the service in the [Authorization request header](#) as a bearer token.

Services that support OAuth 2.0 shall support receiving an [RFC7519](#)-defined JSON Web Token (JWT) in the `Authorization` request header.

JWTs are a compressed JSON structure that contain a JOSE Header, a set of claims that describe the type of access that is granted to a client, and a signature. Each component of a JWT is Base64URL-encoded and concatenated with a `.` to form the token string for the `Authorization` header.

[Table 44](#) describes the JWT JOSE Header parameters and their requirements for services and clients. Any other parameters are outside the scope of this specification.

Services shall process the parameters in [Table 44](#) if the **Service requirement** column contains **Yes**. Services should process other parameters.

The JWT provided by the client shall contain the parameters in [Table 44](#) if the **JWT requirement** column contains **Yes**. The JWT provided by the client may omit other parameters.

**Table 44 — OAuth 2.0 JWT JOSE Header parameters**

Parameter	Service requirement	JWT requirement	Description
<code>typ</code>	Yes	No	Type of token. The string is case-insensitive. If not present, services shall assume the value is <code>JWT</code> .
<code>alg</code>	Yes	Yes	Algorithm for the signature of the token. Services shall not accept the value <code>none</code> .

[Table 45](#) describes the claims and their requirements for services and clients. Any other claims are outside the scope of this specification.

Services shall process the claims in [Table 45](#) if the **Service requirement** column contains **Yes**. Services should process other claims.

The JWT provided by the client shall contain the claims in [Table 45](#) if the **JWT requirement** column contains **Yes**. The JWT provided by the client may omit other claims.

**Table 45 — OAuth 2.0 JWT claims**

Claim	Service requirement	JWT requirement	Description
<code>iss</code>	Yes	Yes	Issuer of the token. Identifies the <i>authorization server</i> that signed the token.
<code>sub</code>	Yes	Yes	Subject of the token. Identifies the client issued the token.
<code>aud</code>	Yes	Yes	Audience of the token. Identifies the <i>resource server</i> intended to accept the token.
<code>exp</code>	Yes	No	Expiration time of the token.
<code>nbf</code>	Yes	No	"Not before" time of the token.
<code>iat</code>	Yes	No	Issued time of the token.
<code>jti</code>	Yes	No	Unique identifier of the token.
<code>scope</code>	Yes	Yes	Type of access the token grants. See the <a href="#">Redfish OAuth2.0 scope usage</a> clause.

Example JOSE Header:

```
{
  "typ": "JWT",
```

```
    "alg": "RS256"  
}
```

Example JWT claims:

```
{  
  "iss": "https://contoso.org/services/oauth2",  
  "sub": "Joe Smith",  
  "aud": "92384634-2938-2342-8820-489239905423",  
  "exp": 1735707600,  
  "scope": "Redfish.Role.Operator",  
  "jti": "97d52311-5f55-4482-b947-8a70c326fdfd"  
}
```

Example token encoded in the `Authorization` request header:

```
Authorization: Bearer mF_9.B5f-4.1JqM
```

**Note:** The previous example does not reflect a real JWT and is provided to show encoding in the `Authorization` request header.

#### 13.4.4.4 Redfish OAuth2.0 scope usage

The value of the `scope` claim is expressed as a list of space-delimited, case-sensitive strings. Each value in the list describes a type of access that was granted to the client.

This specification defines two formats for values in the `scope` claim: [Redfish roles](#) and [Redfish privileges](#). Other formats are outside the scope of this specification.

Redfish roles within the `scope` claim shall be in the form `Redfish.Role.<RoleId>` where `<RoleId>` is the identifier of the Redfish role granted to the client.

Redfish privileges within the `scope` claim shall be in the form `Redfish.Privilege.<PrivilegeId>` where `<PrivilegeId>` is the standard privilege or [OEM privilege](#) granted to the client.

Services shall ignore unsupported values in the `scope` claim. If the token provided by the client is valid, the service shall apply roles and privileges in the `scope` claim to the operation.

## 13.5 Account service

### 13.5.1 Account service overview

- Implementations should store user passwords with one-way encryption techniques.
- Implementations may support exporting user accounts with passwords, but shall do so using encryption methods to protect them.
- User accounts shall support ETags and atomic operations. Implementations may reject requests that do not include an ETag.
- When authentication fails, extended error messages shall not provide privileged information.

### 13.5.2 Password management

A Redfish service provides local user accounts through a collection of `ManagerAccount` resources located under the account service. The `ManagerAccount` resources enable users to manage their own account information, and for administrators to create, delete, and manage other user accounts.

When account properties are changed, the service may close open sessions for this account and require re-authentication.

### 13.5.3 Atomic password changes

If the `RequireChangePasswordAction` property in the `AccountService` resource contains `true`, services shall reject modifications to the `Password` property in `ManagerAccount` resources in `PATCH` or `PUT` operations. Clients are required to invoke the `ChangePassword` action in the `ManagerAccount` resource to update the password for a user account.

This functionality can be used to protect the service from session hijacking. The `SessionAccountPassword` parameter in the `ChangePassword` action allows the service to verify the client's credentials before changing the password of the account. If a user is changing their own password, the `SessionAccountPassword` parameter shall contain the current password of the account. If an administrator is performing the password change for a different user, the `SessionAccountPassword` parameter shall contain the administrator's password.

If the client is performing requests with [HTTP Basic authentication](#), the `SessionAccountPassword` contains the same password encoded in the `Authorization` request header.

### 13.5.4 Password change required handling

The service may require that passwords assigned by the manufacturer be changed by the end user prior to

accessing the service. In addition, administrators may require users to change their account's password upon first access.

The `ManagerAccount` resource contains a `PasswordChangeRequired` boolean property to enable this functionality. Resources that have the property set to `true` shall require the user to change the write-only `Password` property in that resource before access is granted. Manufacturers including user credentials for the service may use this method to force a change to those credentials before access is granted.

When a client accesses the service by using credentials from a `ManagerAccount` resource that has a `PasswordChangeRequired` value of `true`, the service shall allow:

- A session login and include the `@Message.ExtendedInfo` annotation in the response containing the `PasswordChangeRequired` message from the Base Message Registry. This indicates to the client that their session is restricted to performing only the password change operation before access is granted.
- A `GET` operation on the `ManagerAccount` resource associated with the account.
- A `PATCH` operation on the `ManagerAccount` resource associated with the account to update the `Password` property if the `RequireChangePasswordAction` in the `AccountService` resource contains `false` or is not present. If the value of `Password` is changed, the service shall also set the `PasswordChangeRequired` property to `false`.
- A `POST` operation on the `ChangePassword` action on the `ManagerAccount` resource associated with the account to update the `Password` property. If the value of `Password` is changed, the service shall also set the `PasswordChangeRequired` property to `false`.
- A `DELETE` operation on `Session` resources representing open sessions associated with the account.

When a client accesses the service by using credentials from a `ManagerAccount` resource that has a `PasswordChangeRequired` value of `true`, the service may allow `GET` operations on [unauthenticated resources](#), such as the `ServiceRoot` resource.

For all other operations, the service shall respond with the HTTP `403 Forbidden` status code and an [error response](#) with the `PasswordChangeRequired` message from the Base Message Registry.

### 13.5.5 Time-based One-Time Password secret key handling

When RFC6238-defined Time-based One-Time Password (TOTP) multi-factor authentication is enabled, the user account is not configured to bypass Time-based One-Time Password multi-factor authentication, and the user account does not contain a secret key for Time-based One-Time Password multi-factor authentication, the service:

- Shall allow a session login without the `Token` property and include the `@Message.ExtendedInfo` annotation in the response containing the `GenerateSecretKeyRequired` message from the Base Message Registry. This indicates to the client that their session is restricted to performing only the `GenerateSecretKey` action on their `ManagerAccount` resource before access is granted.
- Shall allow a `POST` operation on the `GenerateSecretKey` action on the `ManagerAccount` resource associated with the account to generate a Time-based One-Time Password secret key.
  - Clients shall retain the value of the secret key in the response to generate tokens for future session creation

requests.

- Shall allow a `POST` operation on the `VerifyTimeBasedOneTimePassword` action on the `ManagerAccount` resource associated with the account, if the action is supported by the service. This allows a client to verify they are able to produce valid tokens with their secret key prior to closing the session.
- Shall allow a `DELETE` operation on `Session` resources representing open sessions associated with the account.
- May allow `GET` operations on [unauthenticated resources](#), such as the `ServiceRoot` resource.

For all other operations, the service shall respond with the HTTP `403 Forbidden` status code and an [error response](#) with the `GenerateSecretKeyRequired` message from the Base Message Registry.

When the secret key is generated for the user with the `GenerateSecretKey` action on the `ManagerAccount` resource associated with the account, the service shall require the client to create a new session with the `Token` property to perform other operations.

## 13.6 Asynchronous tasks

---

Irrespective of which user or privileged context starts a [task](#), services shall enforce the privileges described in the privilege registry required to perform operations on the `Task` resource.

## 13.7 Event subscriptions

---

Before pushing event data object to the destination, the service may verify the destination for identity purposes.

Services shall not send events to clients that do not have sufficient privilege to access the resource originating the event. For push-style eventing, this is determined by the client that [performed an HTTP POST to the subscription collection](#). For server-sent events, this is determined by the client that [opened the SSE connection](#).

## 14 Redfish Host Interface

The [\*Redfish Host Interface Specification\*](#) defines how software that runs on a host computer system can interface with a Redfish service that manages the host. For details, see [DSP0270](#).

## 15 Redfish composability

A service may implement the `CompositionService` resource off of `ServiceRoot` to bind resources. One example is disaggregated hardware, which allows for independent components, such as processors, memory, I/O controllers, and drives, to be bound to create logical constructs that operate together. This enables a client to dynamically assign resources for an application.

A service that supports composability shall implement resource blocks, defined by the `ResourceBlock` schema, and resource zones, defined in the `Zone` schema, for the composition service. Resource blocks provide an inventory of components available to the client for building compositions. Resource zones describe the binding restrictions of the resource blocks that the service manages.

The resource zones within the composition service shall include the [collection capabilities annotation](#) in responses. The collection capabilities annotation allows a client to discover which resource collections in the service support compositions, the different [composition request](#) types allowed, how the `POST` request for the resource collection is formatted, and which properties are required.

A service that supports composability and client multi-tenancy shall:

- Implement the `FreePool` and `ActivePool` properties in the `CompositionService` resource.
- Implement the `CompositionReservations` property in the `CompositionService` resource.
- **Filter** `GET` requests for the `ResourceBlocks`, `FreePool`, `ActivePool`, `ResourceZones`, and `CompositionReservations` resource collections where the value of the `Client` property in the `ResourceBlock` resource OR `CompositionReservation` resource matches the client identity.
- Ensure the resources in [composition requests](#) are assigned to the client specified by the `Client` property in the `ResourceBlock` resource or `CompositionReservation` resource.
- Not filter any HTTP operations within the composition service for clients that contain the privilege `ConfigureCompositionInfrastructure` unless specified by [query parameters](#).
- Move resource blocks between the `FreePool` and `ActivePool` resource collections based on the outcome of [composition requests](#).
  - A resource block is moved to the `FreePool` resource collection when it is not contributing to any composed resources.
  - A resource block is moved to the `ActivePool` resource collection when it is contributing to one or more composed resources.

## 15.1 Composition requests

### 15.1.1 Composition requests overview

A service that implements the composition service, as defined by the `CompositionService` schema, shall support one or more of the following types of composition requests:

- Specific composition
- Constrained composition
- Expandable resources

A service that supports the removal of a composed resource shall support the `DELETE` method on the composed resource.

A service may implement the `Compose` action in the `CompositionService` resource for the above composition requests.

### 15.1.2 Specific composition

A specific composition is when a client identifies an exact set of resources in which to build a logical entity.

A service that supports specific compositions shall support a `POST` request that contains an array of hyperlinks to resource blocks. The schema for the resource being composed defines where the resource blocks are specified in the request.

The following example shows a `ComputerSystem` being composed with a specific composition request:

```
POST /redfish/v1/Systems HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
  "Name": "Sample Composed System",
  "Links": {
    "ResourceBlocks": [
      {"@odata.id": "/redfish/v1/CompositionService/ResourceBlocks/ComputeBlock0"},
      {"@odata.id": "/redfish/v1/CompositionService/ResourceBlocks/DriveBlock2"},
      {"@odata.id": "/redfish/v1/CompositionService/ResourceBlocks/NetBlock4"}
    ]
}
```

```
}
```

### 15.1.3 Constrained composition

A constrained composition is when a client has identified a set of criteria, or constraints, in which to build a logical entity. This includes criteria such as quantities of components, or characteristics of components. A service that supports constrained compositions shall support a `POST` request that contains the set of characteristics to apply to the composed resource. The specific format of the request is defined by the schema for the resource being composed. This type of request may include expanded elements of resources [subordinate](#) to the composed resource.

The following constrained composition request composes a `ComputerSystem`:

```
POST /redfish/v1/Systems HTTP/1.1
Content-Type: application/json; charset=utf-8
Content-Length: <computed length>
OData-Version: 4.0

{
  "Name": "Sample Composed System",
  "PowerState": "On",
  "BiosVersion": "P79 v1.00 (09/20/2013)",
  "Processors": {
    "Members": [
      {
        "@Redfish.RequestedCount": 4,
        "@Redfish.AllowOverprovisioning": true,
        "ProcessorType": "CPU",
        "ProcessorArchitecture": "x86",
        "InstructionSet": "x86-64",
        "MaxSpeedMHz": 3700,
        "TotalCores": 8,
        "TotalThreads": 16
      }
    ],
    "Memory": {
      "Members": [
        {
          "@Redfish.RequestedCount": 4,
          "CapacityMiB": 8192,
          "MemoryType": "DRAM",
          "MemoryDeviceType": "DDR4"
        }
      ],
      "SimpleStorage": {
        "Members": [
          {
            "@Redfish.RequestedCount": 6,
            "Devices": [
              {
                "Name": "SSD 1"
              },
              {
                "Name": "SSD 2"
              },
              {
                "Name": "HDD 1"
              },
              {
                "Name": "HDD 2"
              },
              {
                "Name": "HDD 3"
              },
              {
                "Name": "HDD 4"
              }
            ]
          }
        ]
      }
    }
  }
}
```

```
        "CapacityBytes": 322122547200
    }]
}
},
"EthernetInterfaces": {
    "Members": [
        {
            "@Redfish.RequestedCount": 1,
            "SpeedMbps": 1000,
            "FullDuplex": true,
            "NameServers": ["names.redfishspecification.org"],
            "IPv4Addresses": [
                {
                    "SubnetMask": "255.255.252.0",
                    "AddressOrigin": "Dynamic",
                    "Gateway": "192.168.0.1"
                }
            ]
        }
    ]
}
```

#### 15.1.4 Expandable resources

An expandable resource is when a service has a baseline composition that cannot be removed. Instead of a client making requests to create a composed resource, a client can only add or remove resources from the composed resource. A service that supports expandable resources shall support one or more of the update methods that the [Updating a composed resource](#) clause describes.

## 15.2 Updating a composed resource

A service that supports updating a composed resource shall provide one or more of the following methods to update composed resources:

- The `PUT` or `PATCH` methods on the composed resource with a modified list of resource blocks.
- Actions on the composed resource for adding and removing resource blocks.
  - If the actions for adding and removing resource blocks are present in the resource, clients should use this method before attempting `PUT` or `PATCH`.

# 16 Aggregation

Aggregation has been a Redfish concept since its inception. Redfish uses collection for services that can represent more than one system. As the scale of Redfish implementations increase, clients want to operate on Redfish resources in bulk.

Aggregation is the representation of Redfish resources from a variety of sources so that they can be managed, in whole or in part, by a Redfish client. Membership can be heterogeneous and arbitrary, but it is expected that most aggregate members are the same resource type, such as an aggregate of `ComputerSystem` resource, which is represented by an `Aggregate` resource where members of its `Elements` array are exclusively of type `ComputerSystem`. The Redfish service proxies on behalf of the aggregated components to provide common operations. The Redfish service is representing resources on behalf of the components and incoming operations must be tracked by the Redfish service before being accomplished by communicating with the individual resources. Thus, aggregation also allows a Redfish client to act on resources as a group using aggregates.

## 16.1 Classes of aggregators

### 16.1.1 Implicit and complex aggregators

There are at least two classes of Redfish aggregators:

- **Implicit aggregators.** An example of an implicit aggregator is an enclosure manager, such as a manager of blades in an enclosure. This implementation has `ComputerSystem` resources representing blades in the `ComputerSystemCollection` resource, and one or more `Manager` resources in the `ManagerCollection` resource. It also would likely have a `Chassis` resource for each blade and a `Chassis` resource for the enclosure, which would use the `Contains` property in `Links` to express the containment relationship to the individual blades. This class of aggregator has tight coupling with system design and proxies requests to and from the blades to perform management functions.
- **Complex aggregators.** An example of a complex aggregator is a rack-level manager, fabric manager, or a manager of similar scale, especially if it represents resources that it gathers through the proxy of information from other managers, like BMCs. The sources that this manager aggregates are more complex in nature and potentially varying. This manager probably has an interface to the resources and proxies the Redfish service on behalf of each set of resources. At this scale, a Redfish client would prefer to provide common functions, such as resetting a set of systems, to the Redfish service as a whole rather than invoking actions individually to achieve scalability requirements. This class of service also may need assistance in adding members to the service, such as providing address and account information for the aggregator to contact the components and initiate the proxy of Redfish operations.

### 16.1.2 Use cases

Several use cases make explicit aggregator representation necessary. What they have in common is the need for common functions for scalability. There are several classes of these common functions.

One use case is service-type functions. An example is a firmware update on a large number of systems. Rather than invoke actions on individual resources, it is more efficient for a client to specify to which resources to apply the image. In this case, a service already exists in the model so an aggregation service is not needed. Instead, the existing service must be augmented to enable the application of an image to a list of resources.

Another use case is common actions. Examples are the `Reset` or `SetDefaultBootOrder` actions. These actions are defined in the `ComputerSystem` schema, but the Redfish URI structure requires that the action occur on each `ComputerSystem` resource. Thus, an individual operation applies to each resource. It is more efficient for a client to send one action with the list of the resources to which to apply the action. For example, to reset one thousand systems, sending one thousand individual reset operations requires significant overhead as compared to sending a single operation with a list of one thousand systems to reset.

A final use case is changing an attribute on multiple members of a collection. An example is changing the boot order on a large number of systems. This use case requires one operation per system. However, assuming the resources are in the same collection, the [deep PATCH operation](#) meets the requirements of this use case.

## 16.2 Aggregation service

---

### 16.2.1 Aggregation service overview

The `AggregationService` resource represents the Redfish aggregation service, which provides aggregation functions.

The aggregation service contains the group actions that can apply to groups of resources. The `AggregationService` schema defines the common actions that a client can take on groups of resources. These actions take an array of resource URIs as one of the parameters to which the action applies. If all members of the resource array do not support the method, a `4xx` [status code](#) shall be returned and the body shall contain an [error response](#). If at least one member of the resource array successfully completed the action but others did not, the status code should be `200 OK` with `@Message.ExtendedInfo` objects for the failed members.

The aggregation service also contains `Aggregate`, `AggregationSource`, and `ConnectionMethod` resources.

### 16.2.2 Aggregator requirements

By implementing the `AggregationService` resource and including an `AggregationSourceCollection` resource, a complex aggregator shall meet the following requirements:

- Proxy to the aggregated resources on behalf of the service.
- Provide error and state propagation, such as health roll-up, when needed to provide such data to the [parent resource](#).
- Combine resource collections from the aggregated resources.
  - For example, `ComputerSystem` resources that were gathered through proxy shall be in one `ComputerSystemCollection` resource.
  - Services shall complete a *URI fix-up* for all aggregated resources because every system cannot be at `/redfish/v1/Systems/1`.
  - It is advisable for Redfish implementations to use unique values for the `Id` properties. For example, base the `Id` property of a `ComputerSystem` resource on something unique like a UUID or serial number, or the manufacturer MAC address for network adapters, or WWN for Fibre Channel controllers.
- Unify other services.
  - The aggregation implementation hosts only one event service. The implementation shall combine all events into one stream. The implementation also hosts only one sessions service, telemetry service, update service, and other services. Thus the aggregator represents unification of Redfish services with which it communicates and proxies on the client's behalf to the providers of those services and information.

### 16.2.3 Aggregates

The `Aggregate` resource is the grouping mechanism that clients use to indicate to the service that this group of resources can be treated the same for certain functions, such as the actions. Each aggregate contains the list of individual resources that are to be treated as a single unit for operations. For example, if a client wishes to express that a subset of the `ComputerSystemCollection` resource be treated as a single unit for certain operations like reset, reset boot order, or firmware update, it can express the aggregate as the target URI for the operation.

The `Aggregate` schema defines the common actions that a client can make on an aggregate. The `Aggregate` resource contains an `Elements` array that specifies the members of the aggregate. Actions that are supported on an aggregate but not supported on all `Elements`, such as a `Reset` action that is not supported on an individual member of the `Elements` array, are not silently skipped. If all members of the `Elements` array do not support the method, a [4xx status code](#) shall be returned and the body shall contain an [error response](#). If at least one member of the `Elements` array successfully completed the action, but others did not, the status code should be [200 OK](#) with `@Message.ExtendedInfo` objects for the failed members.

### 16.2.4 Aggregation sources and connection methods

The aggregation service model also includes a definition for the information used to access the resources being represented by the aggregator. Two collections of resources are used to represent this. These are the `AggregationSource` and `ConnectionMethod` resources.

The `AggregationSource` resource represents the source of information for the resources being reflected by the aggregator. It typically represents a lower layer service provided by another manager. It contains information needed

to access that source, such as the address and account information. It also has a reference to the `ConnectionMethod` resource used to access it.

The `ConnectionMethod` resource represents the protocol and other semantics required to communicate with the resources being aggregated. Examples of connection methods are Redfish, IPMI, and proprietary access methods. For methods such as IPMI, it's also possible to specify the variations and nuances from multiple vendors.

## 17 ANNEX A (informative) Change log

Version	Date	Description
1.22.1	2025-05-01	Updated the <a href="#">ETags</a> clause to provide stronger guidance to prevent unnecessarily frequent ETag updates.
		Updated the <a href="#">M-SEARCH response</a> clause to align the <code>ST</code> response header version usage with UPnP requirements.
		Added the <a href="#">Extended information implementation notes</a> clause to provide guidance for when to use extended information annotations.
		Corrected the example action info URI in the <a href="#">OEM actions</a> clause.
		Clarified the construction of URI patterns in the <a href="#">Resource URI patterns annotation</a> to show how the schema name is used in the identifiers for the annotation.
1.22.0	2025-02-05	Added <code>includeoriginofcondition</code> query parameter to the <a href="#">Query parameter overview</a> clause.
		Added <code>includeoriginofcondition</code> query parameter to the <a href="#">Server-sent events</a> clause.
		Corrected the <code>\$expand</code> with <code>\$select</code> example in <a href="#">The \$expand query parameter</a> to show the full property path to the selected property is required.
1.21.1	2024-11-27	Updated example text in <a href="#">The \$expand query parameter</a> for clarity.
		Updated <a href="#">OEM-specified object naming</a> to add <code>OpenCompute</code> to the organization exception list.
		Added the <a href="#">Time-based One-Time Password secret key handling</a> clause to describe service behaviors when a user is required to provide an RFC6238-defined Time-based One-Time Password as the <code>Token</code> property during session creation, but a secret key is not configured.
1.20.2	2024-08-01	Updated the <a href="#">Registries</a> clause to correct terminology for attribute registry files.
		Added the <code>/redfish/v1/Registries</code> URI to the <a href="#">Redfish-defined URLs and relative reference rules</a> clause.
		Updated the <a href="#">Schema modification rules</a> clause to give guidance for resolving multiple versions of the same modified schema.
		Corrected terminology in the <a href="#">OEM actions</a> clause to better describe how the action property name is constructed in payloads.
		Updated the <a href="#">Naming rules</a> clause to restrict the usage of <code>ReturnType</code> for action parameter names.
		Updated the <a href="#">Password change required handling</a> clause to clarify that the HTTP <code>403 Forbidden</code> response body follows the error response format.

Version	Date	Description
1.20.1	2024-04-03	Updated the <a href="#">Event subscriptions</a> clause to clarify that privilege enforcement is required when sending events.
		Updated the <a href="#">Query parameter overview</a> clause to optionally allow for case-insensitive verification of query parameter names.
		Updated the <a href="#">Naming rules</a> clause to prevent naming collisions where the only difference is the casing of letters.
		Updated the <a href="#">OEM property format and content</a> clause to give guidance for constructing the value of <code>@odata.type</code> to avoid naming collisions.
		Updated the <a href="#">Schema modification rules</a> clause to allow for non-semantic text changes.
		Added the <a href="#">Transiently unavailable resources</a> clause to describe best practices for showing the status of a device that is in an unavailable state.
		Clarified <a href="#">The \$expand query parameter</a> to show that HTTP <code>200 OK</code> is the proper response code when a resource collection is partially expanded.
		Updated the <a href="#">Status codes</a> clause to remove the recommendation to avoid the HTTP <code>100 Continue</code> status code since this is mandatory to support for chunked client requests from the HTTP 1.1 specifications.
1.20.0	2023-11-30	Added requirement to the <a href="#">Establishing an outbound connection</a> clause to require including the <code>Sec-WebSocket-Protocol</code> HTTP header with the value <code>Redfish</code> .
		Updated the <a href="#">POST (action)</a> clause to require services to accept empty JSON objects for actions that do not have required parameters. <b>This change might require modification to an implementation to support this version of the specification.</b>
1.19.1	2023-11-30	Made several changes for style consistency, grammar, and general clarity. Except for the following additions, no normative changes were made. Any clarifications that inadvertently altered the normative behavior are considered errata and will be corrected in future revisions to the specification.
		Clarified the <a href="#">Password change required handling</a> clause to show that users can delete their own sessions.
		Clarified the <a href="#">Password change required handling</a> clause to show that performing <code>GET</code> operations on unauthenticated resources might be allowed by the service.
		Updated the <a href="#">Action info annotation</a> and <a href="#">OEM actions</a> clauses to give guidance for constructing the URI for the <code>ActionInfo</code> resource.
		Clarified the <a href="#">Property element</a> clause to better describe the usage of <code>Edm.Decimal</code> and <code>Edm.Double</code> .
		Updated the <a href="#">PATCH (update)</a> , <a href="#">POST (create)</a> , and <a href="#">POST (action)</a> clauses to give guidance to not require clients to specify extraneous data in request bodies.

Version	Date	Description
		Clarified the <a href="#">Multipart HTTP push updates</a> to show that clients are not required to provide an OEM form.
		Updated the <a href="#">Next link ( Members@odata.nextLink ) property</a> clause to give guidance for handling a client's original query parameters that affect the number of members returned in the collection.
		Updated the <a href="#">Next link ( Members@odata.nextLink ) property</a> clause to give guidance for using the <code>\$skiptoken</code> query parameter.
		Corrected the <a href="#">PATCH (update)</a> clause to require responding with the HTTP <code>400 Bad Request</code> with the <code>NoOperation</code> message or a modification success response when the client only gives OData annotations in the request body.
		Clarified the <a href="#">Allowable values for strings</a> , <a href="#">Allowable patterns for string values</a> , <a href="#">Allowable values for numbers and durations</a> clauses to show that these terms are used to express client modification restrictions on properties.
		Updated <a href="#">The \$expand query parameter</a> to describe behavior when used with resources containing reference objects with the expanded resource annotation.
		Updated the <a href="#">Response headers</a> clause to show that the <code>WWW-Authenticate</code> is only required if HTTP Basic authentication is configured to advertise its capability.
		Updated the <a href="#">HTTP Basic authentication</a> clause to explain that the <code>HTTPBasicAuth</code> property in the <code>AccountService</code> resource can disable or restrict this functionality.
		Clarified <a href="#">The \$expand query parameter</a> to describe the behavior when a client performs an expansion request, but does not have sufficient privilege to access referenced resources.
		Clarified the <a href="#">Naming rules</a> clause to describe naming pattern for string properties containing URIs.
		Added the <a href="#">URI annotation</a> clause to document how to indicate a string property is expected to contain a URI.
1.19.0	2023-08-03	Added the <a href="#">Atomic password changes</a> clause to define behavior for services that require password changes with the <code>ChangePassword</code> action.
		Updated the <a href="#">Password change required handling</a> clause to allow for a client to invoke the <code>ChangePassword</code> action when a password change is required.
		Added the <a href="#">Outbound connections</a> clause to define how a service creates a WebSocket with a remote client through which Redfish requests are tunneled.
		Updated the <a href="#">Session login</a> clause to define specific messages in responses when a token is required in the request to create a session.
1.18.1	2023-08-03	Clarified the <a href="#">Asynchronous operations</a> clause to state that HTTP <code>202 Accepted</code> response bodies contain a <code>Task</code> resource or are absent.
		Updated the <a href="#">ETags</a> clause to recommend weak ETag comparisons in <code>PUT</code> or <code>PATCH</code> operations.

Version	Date	Description
		Updated the <a href="#">Naming rules</a> clause to clarify all names are alphanumeric and start with letters.
		Updated the <a href="#">Naming rules</a> clause to clarify the naming style for enumeration values.
		Updated the <a href="#">Deep operations</a> clause to show that a deep <code>POST</code> can include different types of resources in the request.
		Updated the <a href="#">Redfish-defined URIs and relative reference rules</a> and <a href="#">Asynchronous operations</a> clauses to give guidance on the recommended URI pattern for task monitors.
		Updated <a href="#">The \$expand query parameter</a> to recommend that all properties in requested resources are included in responses.
1.18.0	2023-04-04	Added <a href="#">URI segment annotation</a> clause to allow schema to describe when URIs do not meet expected naming rules.
1.17.1	2023-04-04	Updated example in <a href="#">M-SEARCH response</a> to better explain the different values a service provides in an <code>M-SEARCH</code> response.
		Updated <a href="#">Writable properties annotation</a> to include an example.
		Updated <a href="#">Modification success responses</a> to provide links to appropriate sections for each operation. Clarified that response bodies for create operations are optional.
		Updated the response table in <a href="#">POST (action)</a> to use consistent terminology for action responses.
		Updated <a href="#">Resource and operation authentication requirements</a> to state that services can optionally reject requests to unauthenticated resources if the provided credentials are invalid.
		Updated <a href="#">Action responses</a> to provide guidance for finding the schema definition of the action response based on the action's name in response payloads.
		Updated <a href="#">The \$select query parameter</a> to clarify that unsupported properties are omitted from responses.
		Updated <a href="#">Response headers</a> to make <code>Access-Control-Allow-Origin</code> an optional response header.
		Clarified <a href="#">Duration values</a> to state that negative durations are not allowed.
1.17.0	2022-12-08	Updated <a href="#">Deprecated annotation</a> to allow for deprecating URIs.
		Updated <a href="#">Protocol details</a> to allow for optional HTTP 2.0 support.
		Updated <a href="#">Permissions annotation</a> to allow for write-only to be specified to enforce services respond with <code>null</code> for their value.
		Updated <a href="#">HTTP Basic authentication</a> and <a href="#">Redfish session login authentication</a> to describe behavior when multi-factor authentication is enabled.
		Added <a href="#">Client certificate authentication</a> clause to allow for client certificates to be authenticated during TLS handshaking.

Version	Date	Description
1.16.1	2022-12-08	<p>Updated <a href="#">GUID and UUID values</a> to correct the format for GUID and UUID properties.</p> <p>Updated <a href="#">The \$select query parameter</a> to clarify the syntax of <code>\$select</code> for arrays.</p> <p>Updated <a href="#">Status codes</a> to remove references to nonexistent HTTP headers.</p> <p>Updated <a href="#">Sensitive data</a> to clarify that the phrase "URIs containing sensitive data" is meant to reflect the URI itself rather than the response body.</p> <p>Corrected the example <code>RelatedProperties</code> property in message objects throughout the specification to remove the leading <code>#</code> to meet syntax specified by RFC6901.</p> <p>Updated <a href="#">MessageId format</a> to better formalize the string tokens used to distinguish the different parts of <code>MessageId</code> values.</p>
1.16.0	2022-08-04	<p>Added optional standard roles to the <a href="#">Roles</a> clause.</p> <p>Added new <a href="#">Allowable patterns for string values</a>, <a href="#">Allowable values for numbers and durations</a>, and <a href="#">Writable properties annotation</a> payload annotations.</p> <p>Renamed allowable values payload annotation to <a href="#">Allowable values for strings</a>.</p>
1.15.2	2022-08-04	<p>Updated <a href="#">URI naming rules</a> to give an example URI when a hyperlink to a subordinate resource is not found at the root of the resource.</p> <p>Updated <a href="#">Modification success responses</a> to not imply a service deficiency if HTTP <code>204 No Content</code> is returned.</p> <p>Clarified the usage of the <code>WWW-Authenticate</code> header in the <a href="#">Response headers</a> clause.</p> <p>Updated <a href="#">Redfish service operation-to-privilege mapping</a> to explain that the service can filter response data based on the user's identity beyond what is assigned as a privilege.</p> <p>Updated <a href="#">Links</a> to state that some design exceptions have been made for putting a hyperlink inside of <code>Links</code> or at the root of a resource.</p> <p>Updated <a href="#">ETags</a> to give guidance to implementers when managing resources that update frequently.</p>
		<p>Updated <a href="#">Properties overview</a> to state that all properties are required to have well-known types, with some exceptions.</p> <p>Updated <a href="#">OEM-specified object naming</a> to allow for exceptions to domain suffix rules for listed organizations.</p> <p>Updated <a href="#">Asynchronous operations</a> to better distinguish <code>Task</code> resources from task monitors to the reader.</p> <p>Updated <a href="#">The \$select query parameter</a> to allow for object-level messages and messages for requested properties even if not specified by the client.</p>

Version	Date	Description
		Clarified <a href="#">The \$select query parameter</a> that error responses act as if <code>\$select</code> was not specified.
1.15.1	2022-04-07	Updated <a href="#">Query parameter overview</a> to recommend the <code>=</code> character is ignored if specified with the <code>only</code> and <code>excerpt</code> query parameters.
		Updated <a href="#">Multipart HTTP push updates</a> to use consistent terminology for action parameters, and to correct the use of HTTP status code 413 instead of 412 to indicate a payload is too large to be processed by the service.
		Updated <a href="#">Schema definition languages</a> to use normative terminology throughout.
		Clarified <a href="#">Asynchronous tasks</a> to explain the allowable operations on a <code>Task</code> resource are controlled by the privilege registry.
		Clarified <a href="#">ETags</a> to describe the behavior of ETags during expansion.
		Deprecated statement in <a href="#">POST to subscription collection</a> with regards to keeping <code>MessageIds</code> and <code>RegistryPrefixes</code> as mutually exclusive.
		Updated <a href="#">Subscribing to events</a> to match the schema updates made to <code>EventDestination</code> .
1.15.0	2021-12-02	Updated the <a href="#">Operation apply time</a> and <a href="#">Multipart HTTP push updates</a> to expand the usage of the <code>@Redfish.OperationApplyTimeSupport</code> annotation to properties referencing URIs for multipart HTTP POST operations.
1.14.2	2021-12-02	Clarified that the <code>Created</code> message from the Base Message Registry is an acceptable response in the <a href="#">POST (action)</a> clause when a new resource is created.
		Clarified the <a href="#">Permissions annotation</a> clause to describe requirements for schemas published by DMTF.
		Updated the <a href="#">Units of measure annotation</a> clause to describe encoding rules for units not covered by UCUM. Provided alternative recommendations for RPM units.
		Clarified <a href="#">Collection capabilities annotation</a> clause that properties marked as required are not required for the resource referenced by the <code>CapabilitiesObject</code> property.
		Clarified that <code>@Redfish.MaintenanceWindow</code> property behavior in <code>Task</code> resources in the <a href="#">Asynchronous operations</a> clause.
1.14.1	2021-10-06	Corrected various examples for <code>@Redfish.OperationApplyTimeSupport</code> , <code>@Redfish.Settings</code> , <code>@Redfish.MaintenanceWindow</code> , and <code>@Redfish.SettingsApplyTime</code> to include <code>@odata.type</code> .
		Clarified the <a href="#">Link header</a> clause to state the <code>Link</code> response header requirement only applies to resources.
1.14.0	2021-09-15	Extended <a href="#">Query parameter overview</a> clause to define how OEM query parameters are constructed.
		Added <a href="#">Delegated authorization with OAuth 2.0</a> clause to define how clients provide OAuth 2.0 tokens to a service as a method of authorization.

Version	Date	Description
1.13.1	2021-08-04	Various clarifications to the <a href="#">Extending standard resources</a> clause to better describe naming rules for OEM resources.
		Added recommended URI for local schema files to the <a href="#">Redfish-defined URIs and relative reference rules</a> clause.
		Clarified the <a href="#">OData \$metadata</a> clause that any absolute or relative URI is allowed for referencing schema files.
		Adding missing statement to the <a href="#">URI naming rules</a> clause that <code>Members</code> cannot be used as the value of a URL segment for resource collections.
		Added missing exceptions to the <a href="#">PATCH (update)</a> clause for when <code>@odata.id</code> is to not be ignored by the service.
		Clarified the <a href="#">ETags</a> clause that both strong and weak ETags are allowed in <code>If-Match</code> and <code>If-None-Match</code> request headers.
		Clarified the <a href="#">Deep operations</a> clause to specify that services ignore resources in the payload if no modifications are requested.
		Added the <a href="#">Enumerations</a> clause to clarify the design patterns for creating enumerations.
		Clarified the <a href="#">Id</a> clause that HTTP unsafe characters are not permitted in the value of the <code>Id</code> property due to its usage in URI construction.
		Clarified the <a href="#">Non-resource reference properties</a> that these properties are strings containing URIs.
1.13.0	2021-04-08	Added client multi-tenancy behavior to the <a href="#">Redfish composability</a> clause. This adds free pool, active pool, and composition reservation constructs to Redfish composability.
		Added <code>Compose</code> action as a method of performing composition requests to the <a href="#">Redfish composability</a> clause.
1.12.1	2021-04-08	International Organization for Standardization (ISO) updates:
		Added paragraph numbering.
		Added <a href="#">Foreword</a> to the table of contents as an unnumbered heading, and placed <a href="#">Acknowledgments</a> inside <a href="#">Foreword</a> .
		Made <a href="#">Scope</a> a level-1 clause.
		<b>Normative references:</b> Removed unused normative references and moved some references into <a href="#">Bibliography</a> . The <a href="#">Bibliography</a> lists, for information, those documents which are cited informatively in the document, as well as other information resources.
		Changed <a href="#">Abstract</a> to <a href="#">Introduction</a> .

Version	Date	Description
		Corrected level-1 clauses to remove hanging paragraphs and to correct the occurrence of the single <b>Use cases</b> and <b>Aggregator requirements</b> sub-clauses.
		<p><b>Terms, definitions, symbols, and abbreviated terms:</b></p> <ul style="list-style-type: none"> <li>Combined <b>Symbols and abbreviated terms</b> clause with <b>Terms and definitions</b> clause into <b>Terms, definitions, symbols, and abbreviated terms</b> clause.</li> <li>Formatted the clause correctly.</li> <li>Added the <b>Hardware terms</b>, <b>Web development terms</b>, and <b>Redfish terms</b> sub-clauses to this clause.</li> <li>Removed may, shall, and should from definitions.</li> <li>Removed these terms: <i>managed system</i>, <i>Redfish event receiver</i>, and <i>Redfish provider</i>.</li> <li>Corrected definitions so none begin with an article.</li> </ul>
		Changed <i>may</i> to <i>can</i> or <i>might</i> where appropriate.
		Changed one <i>must</i> to <i>shall</i> .
		Added numbered captions to tables and changed occurrences of <i>the following table</i> to use precise references to the table numbers.
		Fixed broken cross-references.
		Corrected URLs in the deep <code>PATCH</code> example.
		Fixed several query parameter examples where string values were not properly wrapped with single quotes.
		Corrected <code>Accept-Encoding</code> usage to allow for encoded responses if the client does not provide the header to align with RFC7231.
		Clarified usage of <code>DELETE</code> for the <code>@Redfish.OperationApplyTimeSupport</code> term.
		Removed duplicative clauses for <code>HTTP 405 Method Not Allowed</code> usage in <b>PATCH (update)</b> in favor of more general clauses.
		Replaced exception table in <b>PATCH (update)</b> in favor of text.
		Moved error cases from response table in <b>POST (action)</b> to be with other text that describes error cases.
		Added linkage in the description for <code>HTTP 201 Created</code> to reference response bodies for actions.
		Added informative text regarding the usage of <code>If-Match</code> and <code>If-Match-None</code> headers in <code>GET</code> , <code>PATCH</code> , and <code>PUT</code> clauses.
		Clarified the behavior of <code>\$select</code> when an object property is selected.
		Added introductory text to guide readers to other Redfish documents.

Version	Date	Description
		Clarified the ordering of processing query parameters.
		Clarified that update restrictions for a resource can be modified to support writable OEM properties.
		Clarified the <b>Settings resource</b> clause to show behavior of properties in the active resource and settings resource based on the service's capabilities.
		Corrected behavior for usage of <code>null</code> based on the configuration of a resource and other special situations.
		Clarified OEM naming rules for all OEM definitions to ensure names don't collide.
		Removed the term "namespace" from all non-CSDL related clauses and replaced them with references to a new <i>resource type</i> term.
1.12.0	2020-12-01	<p>Added introductory text to the <b>Authorization</b> clause.</p> <p>Clarified usage of <code>RoleId</code> and how there are standard roles, custom OEM roles, and client-defined custom roles.</p> <p>Added <b>Restricted roles and restricted privileges</b> to describe behavior for when roles and privileges are marked as restricted.</p>
1.11.2	2020-12-01	<p>Clarified that the <code>Accept-Encoding</code> header is used to request compression of response bodies.</p> <p>Corrected the <b>PATCH (update)</b>, <b>PUT (replace)</b>, and <b>DELETE (delete)</b> clauses to leverage all normative statements for successful operations found in the <b>Modification success responses</b> clause.</p> <p>Replaced RFC5988 reference with RFC8288.</p> <p>Updated IETF links to use the "IETF Tools" site.</p> <p>Clarified that insert capabilities is just for resource creation.</p> <p>Fixed ETag examples to be RFC7234-conformant.</p> <p>Clarified that OEM resources can have subordinate resources.</p> <p>Replaced RFC4627 reference with RFC8259.</p> <p>Replaced conflicting statements found in "HTTP redirect authentication requirements" with general clause for enforcing authentication and authorization at the target resource.</p> <p>Clarified behavior of <code>@odata.count</code> when a collection is filtered.</p> <p>Created standalone "MessageId format" clause.</p> <p>Removed duplicative text found in the event format table and referenced the message object clauses as needed.</p>

Version	Date	Description
		Corrected the response body specified for a <code>PATCH</code> operation containing read-only properties.
		Added informative text in the intro to the <b>Data model</b> clause describing the methods for OEM extensions.
		Clarified that sensitive data in URIs can be hidden from unauthorized users by returning HTTP <code>404 Not Found</code> .
		Added embedded links to the <code>Location</code> header entry in the response header table.
		Corrected <code>\$select</code> example in the <b>The \$select query parameter</b> clause.
		Corrected several embedded links to direct to the correct clause.
1.11.1	2020-08-04	Added missing clause requiring sensitive data to be returned as <code>null</code> .
		Clarified that <code>Resolution</code> , <code>Severity</code> , and <code>MessageSeverity</code> in responses can be service-defined and not come from a message registry.
		Relaxed schema rules to require description, long description, URI, and capabilities annotations only for schemas published or republished by DMTF.
		Added clauses to <b>Schema modification rules</b> to allow for properties, actions, parameters, and URIs to be removed, descriptions to be modified, and pattern and length annotations to be added if not specified.
		Relaxed rule for the OData metadata document to not require, but only recommend that all referenced namespaces are included in the document.
		Added clause to clarify the usage of empty strings.
		Clarified behavior of <code>\$skip</code> when the value is greater than or equal to the number of members in a resource collection.
		Corrected the minimum value for <code>\$top</code> to align with OData.
		Clarified behavior of <code>PATCH</code> for partial success scenarios.
		Various clarifications and style fixes to the <b>Aggregation</b> clause.
		Clarified that <code>HEAD</code> requests shall be rejected when a query parameter is provided.
		Removed erroneous requirement for ETags to be strong.
1.11.0	2020-04-30	Added <b>Aggregation</b> clause.
		Clarified that services are allowed use HTTP <code>501 Not Implemented</code> for unsupported HTTP methods.
		Clarified the normative semantics around the term "deprecated".

Version	Date	Description
		Clarified clauses describing the usage of <code>null</code> for properties versus not reporting a property.
1.10.0	2020-03-27	<p>Restructured the <b>Security details</b> clause for ease of reading. Other than the changes listed below, no other changes were intended. Any clarifications that inadvertently altered the normative behavior are considered errata, and will be corrected in future revisions to the specification.</p> <p>Deprecated TLS v1.1, and set the minimum TLS requirement to be TLS v1.2 with RFC7525 recommendations.</p> <p>Deprecated existing cipher suites clause in favor of new clause to leverage IANA recommendations.</p> <p>Added requirement for supporting the <code>/redfish</code> URI.</p> <p>Added support for deep operations.</p>
1.9.1	2020-03-27	<p>Deprecated full ISO8601 duration format in favor of a simplified version that does not contain years, months, and weeks.</p> <p>Added missing normative language for how actions with response bodies are defined in schema.</p> <p>Added HTTP <code>201 Created</code> as valid responses for actions.</p> <p>Clarified the <code>~</code> operator for the <code>\$expand</code> query parameter to expand hyperlinks found in all <code>Links</code> properties.</p> <p>Clarified the <code>*</code> and <code>.</code> operators for the <code>\$expand</code> query parameter to expand hyperlinks found in payload annotations, such as <code>@Redfish.Settings</code>.</p> <p>Clarified usage of action parameters that point to resources; the expectation is a reference object pointing to the resource in question is passed by the client.</p> <p>Clarified that <code>DELETE</code> on a resource likely deletes subordinate resources.</p> <p>Clarified best practices for naming rules, in particular with regards to acronyms.</p> <p>Clarified behavior for when individual members of a resource collection cannot be returned as part of a <code>\$expand</code> request.</p> <p>Clarified usage of <code>@Message.ExtendedInfo</code> in error responses and provided guidance for clients for handling error responses.</p>
1.9.0	2019-12-06	<p>Made change to no longer require the <code>Server</code> response header.</p> <p>Added clause to <b>Schema modification rules</b> to allow for the addition of OEM URIs to standard resources.</p> <p>Loosened requirements on <code>@odata.type</code> within <code>Oem</code> to not require it in arrays where the type is used repeatedly.</p>

Version	Date	Description
1.8.1	2019-12-06	Made many changes for style consistency, grammar, and general clarity. Except for the following additions, no normative changes were made. Any clarifications that inadvertently altered the normative behavior are considered errata, and will be corrected in future revisions to the Specification.
		Clarified SSE with regards to requiring a blank line after each event.
		Clarified order of precedence for resolving multiple operation overrides within the Privilege Registry.
		Clarified cases for property overrides in the Privilege Registry where multiple objects in the same resource contain the same property name.
		Updated references for HTTP Basic authentication to use RFC7617 instead of RFC7235.
		Added <code>text/event-stream</code> , <code>application/yaml</code> , and <code>application/vnd.oai.openapi</code> usage to the <code>Accept</code> and <code>Content-Type</code> header table entries.
		Added clause that provides guidance on service behavior when <code>null</code> is a property value in <code>POST</code> (create) operations.
		Loosened requirements on SSE <code>id</code> based on client usage.
		Added documentation for settings, settings apply time, operation apply time, operation apply time support, maintenance window, collection capabilities, requested count, allow over-provisioning, zone affinity, supported certificates, and deprecated terms to the <b>Payload annotations</b> clause.
		Added clauses that document responses for actions with a response body defined in schema.
		Clarified the allowable values payload annotation to show it can be used for both properties and action parameters.
1.8.0	2019-08-08	Added clause for using <code>/redfish/v1/openapi.yaml</code> as the well-known URI for the OpenAPI document.
		Added clause that specifies non-resource reference properties with <code>uri</code> in the name are accessed using Redfish protocol semantics.
		Added <code>SubordinateResources</code> <code>\$filter</code> parameter for SSE.
		Added <b>Update service</b> clause that describes requirements for the <code>SimpleUpdate</code> action and the <code>MultipartHttpPushUri</code> property.
1.7.1	2019-08-08	Added statements about the <code>owning entity</code> annotation term and its usage in schema modifications.
		Clarified SSE <code>id</code> from <code>Id</code> in an event payload and <code>EventId</code> within an event record.
		Fixed recommended sequencing of the SSE <code>id</code> to be related to <code>EventId</code> within an event record.
		Clarified that services are allowed to close sessions for an account when its password has changed.

Version	Date	Description
		Corrected the <b>Password management</b> clause to describe how a user can <code>GET</code> their respective account resources when a password change is required.
		Clarified that registries are not required to return <code>@odata.id</code> .
		Clarified that services should use HTTP <code>400 Bad Request</code> for invalid query requests.
		Clarified that services should use HTTP <code>400 Bad Request</code> when the <code>only</code> query is being combined with other query parameters.
		Clarified that services should use HTTP <code>400 Bad Request</code> when query parameters are used on non-GET operations.
		Added clause about how to construct enumeration values.
		Clarified references to specific messages to also reference their message registry.
		Added language about the construction of action names in payloads.
		Added informative text for how OEM actions can be defined.
		Added guidance for using HTTPS whenever sensitive data is being transmitted.
		Added clause restricting the maximum size of an event payload to be 1MiB.
		Clarified that auto expanded resource collections can use paging.
		Clarified error response format for SSE.
		Clarified that <code>charset=utf-8</code> is not required within the <code>Content-Type</code> header for SSE.
		Added clause about how URI patterns are constructed.
		Added Excerpt term.
1.7.0	2019-05-16	Made many changes for style consistency, grammar, and general clarity. Except for the following additions, no normative changes were made. Any clarifications that inadvertently altered the normative behavior are considered errata, and will be corrected in future revisions to the Specification.
		Added normative statements about how to handle array properties and <code>PATCH</code> operations on arrays.
		Separated data model and schema language clauses.
		Added clauses that describe how JSON Schema and OpenAPI files are formatted.
		Added clause that describes the schema versioning methodology.
		Added clause about how URI patterns are constructed based on the resource tree and property hierarchy.

Version	Date	Description
		Added dictionary file naming rules and repository locations.
		Enhanced localization definitions and defined repository locations.
		Added statement about SSE to the <b>Eventing mechanism</b> clause.
		Added <b>Constrained composition</b> and <b>Expandable resources</b> clauses to Redfish Composability.
		Added clause about requiring event subscriptions to be persistent across service restarts.
		Added clause about persistence of tasks generated as a result of using <code>@Redfish.OperationApplyTime</code> across service restarts.
		Added clause about using <code>@Redfish.OperationApplyTime</code> and <code>@Redfish.MaintenanceWindow</code> within task responses.
		Removed <code>@odata.context</code> property from example payloads.
		Added <b>Password management</b> clause to describe functional behavior for restricting access when an account requires a password change.
		Added clause around the usage of the HTTP <code>403 Forbidden</code> status code when an account requires a password change.
1.6.1	2018-12-13	Added clause about percent encoding being allowed for query parameters.
		Changed <code>\$expand</code> example to use <code>SoftwareInventory</code> instead of <code>LogEntry</code> .
		Added clause about the use of a separator for multiple query parameters.
		Fixed <code>\$filter</code> examples to use <code>/</code> instead of <code>.</code> for property paths.
		Clarified the usage of messages in a successful action response; provided an example.
		Added clarification about services supporting a subset of HTTP operations on resources specified in schema.
		Added clarification about services implementing writable properties as read only.
		Added clarification about session termination not affecting connections opened by the session.
		Added <i>Redfish Provider</i> term definition.
		Updated JSON Schema references to point to Draft 7 of the <i>JSON Schema Specification</i> .
		Added clarifications about scenarios for when a request to add an event subscription contains conflicting information and how services respond.
		Removed language about ignoring the <code>Links</code> property in <code>PATCH</code> requests.
		Clarified usage of ETags to show that a client is not supposed to <code>PATCH @odata.etag</code> when attempting to use ETag protection for a resource.

Version	Date	Description
		Clarified usage of the <code>only</code> query parameter to show it's not to be combined with <code>\$expand</code> and not to be used with singular resources.
		Clarified the usage of the HTTP status codes with task monitors.
		Made various spelling and grammar fixes.
1.6.0	2018-08-23	<p>Added methods of using <code>\$filter</code> on the SSE URI for the event service.</p> <p>Added support for the OpenAPI Specification v3.0. This allows OpenAPI-conforming software to access Redfish service implementations. <b>This change might require modification to an implementation to support this version of the specification.</b></p> <p>Added strict definitions for the URI patterns used for Redfish resources to support OpenAPI. Each URI is now constructed using a combination of fixed, defined path segments and the values of <code>Id</code> properties for resource collections. Also added restrictions on usage of unsafe characters in URIs. Implementations reporting support for Redfish v1.6.0 conform to these URI patterns.</p> <p>Added support for creating and naming Redfish schema files in the OpenAPI YAML-based format.</p> <p>Added URI construction rules for OEM extensions.</p> <p>Changed ETag usage to require strong ETag format.</p> <p>Added requirement for HTTP <code>Allow</code> header as a response header for <code>GET</code> and <code>HEAD</code> operations.</p> <p>Added metric reports as a type of event that can be produced by a Redfish service. Added support for SSE streaming of metric reports in support of new telemetry service.</p> <p>Added registry, resource, origin, or <code>EventFormatType</code>-based event subscription methods as detailed in the Specification and schema. Added an <code>EventFormatType</code> to enable additional payload types for subscription-based or streaming events. Deprecated <code>EventType</code>-based event subscription mechanism.</p> <p>Added event message grouping capability.</p> <p>Provided guidance for defining and using OEM extensions for messages and Message Registries.</p> <p>Added <code>excerpt</code> and <code>only</code> query parameters.</p> <p>Clarified requirements for resource collection responses, which includes required properties that were expected, but not listed explicitly in the Specification.</p> <p>Changed the requirement for the <code>@odata.context</code> annotation to be optional.</p> <p>Removed requirement for clients to include the <code>odata-Version</code> HTTP header in all requests.</p>
1.5.1	2018-08-10	Added clarifications to required properties in structured properties derived from <code>ReferenceableMembers</code> .

Version	Date	Description
		Reorganized <b>Eventing</b> clause to break out the different subscription methods to differentiate pub-sub from SSE.
		Removed statements referencing OData conformance levels.
		Clarified terminology to explain usage of absolute versus relative reference throughout.
		Clarified client-side HTTP <code>Accept</code> header requirements.
		Added evaluation order for supported query parameters and clarified examples.
		Clarified handling of annotations in response payloads when used with <code>\$select</code> queries.
		Clarified service handling of annotations in <code>PATCH</code> requests.
		Clarified handling of various <code>PATCH</code> request error conditions.
		Clarified ability to create resource collection members by <code>POST</code> operations to the resource collection or the <code>Members</code> array within the resource.
		Corrected several examples to show required properties in payload.
		Clarified usage of the <code>Link</code> header and values of <code>rel=describedBy</code> .
		Clarified that the HTTP status code table only describes Redfish-specific behavior and that unless specified, all other usage follows the definitions within the appropriate RFCs.
		Added entry for the HTTP <code>431 Request Header Fields Too Large</code> status code.
		Added statement that the HTTP <code>503 Service Unavailable</code> status code can be used during reboot or reset of a service to indicate that the service is temporarily unavailable.
		Clarified usage of the <code>@odata.type</code> annotation within embedded objects.
		Added statements about the required <code>Name</code> , <code>Id</code> , and <code>MemberId</code> properties, and the common <code>Description</code> property, which have always been shown as required in schema files, but which the Specification did not mention.
		Added guidance for the value of time-date properties when time is unknown.
		Added the <code>title</code> property description in actions.
		Clarified usage of the <code>@odata.nextLink</code> annotation at the end of resource collections.
		Added additional guidance for naming properties and enumeration values that contain "OEM" or that include acronyms.
		Corrected requirements for description and long description annotations.
		Corrected name of <code>ConfigureComponents</code> in the <b>Operation-to-privilege mapping</b> clause.

Version	Date	Description
		Various typographical errors and grammatical improvements.
1.5.0	2018-04-05	Added support for server-sent eventing for streaming events to web-based GUIs or other clients.
		Added <code>@Redfish.OperationApplyTime</code> annotation to provide a mechanism for specifying deterministic behavior for the application of Create, Delete or Action (POST) operations.
1.4.1	2018-04-05	Updated name of the forum from <i>SPMF</i> to <i>Redfish Forum</i> .
		Consistently used the term, <i>hyperlink</i> .
		Added example to clarify usage of <code>\$select</code> query parameter with <code>\$expand</code> , and clarified expected results when using <code>AutoExpand</code> . Corrected order of precedence for <code>\$filter</code> parameter options.
		Corrected terminology for OEM-defined actions removing "custom" in favor of OEM, and clarified that the action <code>target</code> property is always required for an action, along with its usage.
		Corrected location header values for responses to data modification requests that create a task ( <code>Task</code> resource vs. task monitor). Clarified error handling of <code>DELETE</code> operations on <code>Task</code> resources.
		Removed references to obsolete and unused <code>Privilege</code> annotation namespace.
		Clarified usage of the <code>Base.1.0.GeneralError</code> message in the Base Message Registry.
		Added durable URIs for registries and profiles, and clarified intended usage for each folder in the repository. Added file naming conventions for registries and profiles, and clarified file naming for schemas.
		Added statement to clarify that additional headers may be added to M-SEARCH responses for SSDP to enable UPnP compatibility.
		Clarified assignment requirements for predefined or custom roles when new manager account instances are created, using the <code>RoleId</code> property.
1.4.0	2017-11-17	Added support for optional query parameters ( <code>\$expand</code> , <code>\$filter</code> , and <code>\$select</code> ) on requests to enable more efficient retrieval of resources or properties from a Redfish service.
		Clarified HTTP status and payload responses after successful processing of data modification requests. This includes <code>POST</code> operations to complete actions, and other <code>POST</code> , <code>PATCH</code> , or <code>PUT</code> requests.
		Added entries for the HTTP <code>428 Precondition Required</code> and <code>507 Insufficient Storage</code> status codes to clarify the proper response to certain error conditions. Added reference links to the HTTP status code table throughout.
		Updated the <b>Abstract</b> to reflect the current state of the specification.
		Added reference to RFC6585 and clarified expected behavior when ETag support is used in conjunction with <code>PUT</code> or <code>PATCH</code> operations.

Version	Date	Description
		Added definition for <i>Property</i> term and updated text to use term consistently.
		Added <b>Client requirement</b> column and information for HTTP headers on requests.
		Clarified the usage and expected format of the <code>@odata.context</code> property value.
		Added clause to describe how to revise structured properties and resolve their definitions in schema.
		Added more descriptive definition for the settings resource. Added an example for the <code>SettingsObject</code> . Added description and example for using the <code>@Redfish.SettingsApplyTime</code> annotation.
		Added Action example using the <code>ActionInfo</code> resource in addition to the simple <code>@Redfish.AllowableValues</code> example. Updated example to show a proper subset of the available enumerations to reflect a real-world example.
		Added statement explaining the updates required to <code>TaskState</code> upon task completion.
1.3.0	2017-08-11	Added support for a service to optionally reject a <code>PATCH</code> or <code>PUT</code> operation if the <code>If-Match</code> or <code>If-Match-None</code> HTTP header is required by returning the HTTP <code>428 Precondition Required</code> status code.
		Added support for a service to describe when the values in the settings object for a resource are applied via the <code>@Redfish.SettingsApplyTime</code> annotation.
1.2.1	2017-08-10	Clarified wording of the <code>Oem</code> object definition.
		Clarified wording of the <b>Partial resource results</b> clause.
		Clarified behavior of a service when receiving a <code>PATCH</code> with an empty JSON object.
		Added statement about other uses of the HTTP <code>503 Service Unavailable</code> status code.
		Clarified format of URI fragments to conform to RFC6901.
		Clarified use of absolute and relative URIs.
		Clarified definition of the <code>target</code> property as originating from OData.
		Clarified distinction between <i>hyperlinks</i> and the <i>links property</i> .
		Corrected the JSON example of the privilege map.
		Clarified format of the <code>@odata.context</code> property.
		Added clauses about the schema file naming conventions.
		Clarified behavior of a service when receiving a <code>PUT</code> with missing properties.
		Clarified valid values in the <code>Accept</code> header to include wildcards per RFC7231.

Version	Date	Description
		Corrected <code>ConfigureUser</code> privilege to be spelled <code>ConfigureUsers</code> .
		Corrected the <b>Session login</b> clause to include normative language.
1.2.0	2017-04-14	<p>Added support for the Redfish composable service.</p> <p>Clarified service handling of the <code>Accept-Encoding</code> header in a request.</p> <p>Improved consistency and formatting of example requests and responses throughout.</p> <p>Corrected usage of the <code>@odata.type</code> property in response examples.</p> <p>Clarified usage of the required annotation.</p> <p>Clarified usage of <code>SubordinateOverrides</code> in the Privilege Registry.</p>
1.1.0	2016-12-09	<p>Added <b>Redfish service operation-to-privilege mapping</b> clause. This functionality enables a service to present a resource or even property-level mapping of HTTP operations to roles and privileges.</p> <p>Added references to the Redfish Host Interface Specification (DSP0270).</p>
1.0.5	2016-12-09	<p>Errata release. Various typographical errors.</p> <p>Corrected the use of <i>collection</i>, <i>resource collection</i>, and <i>members</i> throughout.</p> <p>Added glossary entries for <i>resource collection</i> and <i>members</i>.</p> <p>Corrected certificate requirements to reference definitions and requirements in RFC5280 and added a normative reference to RFC5280.</p> <p>Clarified usage of the HTTP <code>POST</code> and <code>PATCH</code> operations.</p> <p>Clarified usage of the HTTP status codes and error responses.</p>
1.0.4	2016-08-28	<p>Errata release. Various typographical errors.</p> <p>Added example of an HTTP <code>Link</code> Header and clarified usage and content.</p> <p>Added the <b>Schema modification</b> clause, which describes the allowed usage of the schema files.</p> <p>Added recommendation to use TLS 1.2 or later, and to follow the SNIA TLS Specification. Added reference to the SNIA TLS Specification. Added additional recommended <code>TLS_RSA_WITH_AES_128_CBC_SHA</code> cipher suite.</p> <p>Clarified that the <code>Id</code> property of a <code>Role</code> resource matches the role name.</p>
1.0.3	2016-06-17	<p>Errata release. Fixed the missing numbering in the table of contents and clauses. Corrected URL references to external specifications. Added missing normative references. Corrected typographical error in ETag example.</p> <p>Clarified examples for <code>@Message.ExtendedInfo</code> to show arrays of messages.</p>

Version	Date	Description
		Clarified that a <code>POST</code> to session service to create a new session does not require authorization headers.
1.0.2	2016-03-31	<p>Errata release. Various typographical errors.</p> <p>Corrected normative language for M-SEARCH queries and responses.</p> <p>Corrected <code>Cache-Control</code> and <code>USN</code> format in M-SEARCH responses.</p>
		Corrected schema namespace rules to conform to OData namespace requirements and updated examples throughout the document to conform to this format. Specifically, <code>&lt;namespace&gt;.&lt;n&gt;.&lt;n&gt;.&lt;n&gt;</code> becomes <code>&lt;namespace&gt;.v&lt;n&gt;_&lt;n&gt;_&lt;n&gt;</code> . File naming rules for JSON Schema and CSDL (XML) schemas were also corrected to match this format and to enable future major (v2) versions to coexist.
		Added clause that details the location of the schema repository and lists the durable URLs for the repository.
		Added definition for the value of the Units annotation, using the definitions from the UCUM Specification. Updated examples throughout to use this standardized form.
		Modified the naming requirements for <code>oem</code> property naming to avoid future use of colon <code>:</code> and period <code>.</code> in property names, which can produce invalid or problematic variable names when used in some programming languages or environments. Both separators have been replaced with underscore <code>(_)</code> , with colon <code>(:)</code> and period <code>(.)</code> usage now deprecated (but valid).
		Removed duplicative or out-of-scope subclauses from the <b>Security</b> clause, which made unintended requirements on Redfish service implementations.
		Added the requirement that property names in resource responses match the casing (capitalization) as specified in schema.
		Updated normative references to current HTTP RFCs and added clause references throughout the document where applicable.
		Clarified ETag header requirements.
		Clarified that no authentication is required for accessing the service root.
		Clarified description of retrieving resource collections.
		Clarified usage of <code>charset=utf-8</code> in the HTTP <code>Accept</code> and <code>Content-Type</code> headers.
		Clarified usage of the <code>Allow</code> HTTP response header and added a table entry for the <code>Retry-After</code> header usage.
		Clarified normative usage of the <code>type</code> property and <code>context</code> property, explaining the ability to use two URL forms, and corrected the <code>@odata.context</code> URL examples throughout.
		Corrected inconsistent terminology throughout the resource collection response clause.

Version	Date	Description
		Corrected name of normative resource <code>Members</code> property ( <code>Members</code> , not <code>value</code> ).
		Clarified that error responses may include information about multiple error conditions.
		Corrected name of <code>Measures.Unit</code> annotation term as used in examples.
		Corrected outdated reference to Core OData Specification in annotation term examples.
		Added the <code>Members</code> property to the <b>Common Redfish resource properties</b> clause.
		Clarified terminology and usage of the task monitor and related operations in the <b>Asynchronous operations</b> clause.
		Clarified that implementation of the SSDP protocol is optional.
		Corrected typographical error in the SSDP <code>usn</code> field's string definition (now <code>::dmtf-org</code> ).
		Added the <code>OPTIONS</code> method to the allowed HTTP methods list.
		Fixed nullability in example.
1.0.1	2015-09-17	Errata release. Various grammatical corrections.
		Clarified normative use of long description in schema files.
		Clarified usage of the <code>rel-describedby</code> <code>Link</code> header.
		Corrected text in example of "Select List" in OData context property.
		Clarified <code>Accept-Encoding</code> request header handling.
		Deleted duplicative and conflicting statement on returning extended error resources.
		Clarified relative URI resolution rules.
		Clarified USN format.
1.0.0	2015-08-04	Initial release.

## 18 Bibliography

---

- R. Fielding, 2000, *Architectural Styles and the Design of Network-based Software Architectures*, <https://www.ics.uci.edu/%7Efielding/pubs/dissertation/top.htm>
- IETF RFC5288, J. Salowey et al., *AES Galois Counter Mode (GCM) Cipher Suites for TLS*, <https://tools.ietf.org/html/rfc5288>
- IETF RFC5487, M. Badra et al., *Pre-Shared Key Cipher Suites for TLS with SHA-256/384 and AES Galois Counter Mode*, <https://tools.ietf.org/html/rfc5487>
- IETF RFC5789, L. Dusseault et al., *PATCH Method for HTTP*, <https://tools.ietf.org/html/rfc5789>
- IETF RFC6906, E. Wilde, *The 'profile' Link Relation Type*, <https://tools.ietf.org/html/rfc6906>
- 28 October 1999, *Simple Service Discovery Protocol/1.0 Operating without an Arbiter*, <https://tools.ietf.org/html/draft-cai-ssdp-v1-03>
- 10 March 2016, *OData Version 4.0 Plus Errata 03: Core Vocabulary*, <https://docs.oasis-open.org/odata/odata/v4.0/errata03/csd01/complete/vocabularies/Org.OData.Core.V1.xml>
- 24 February 2014, *OData JSON Format Version 4.0*, <https://docs.oasis-open.org/odata/odata-json-format/v4.0/os/odata-json-format-v4.0-os.html>
- 24 February 2014, *OData Version 4.0 Part 2: URL Conventions*, <https://docs.oasis-open.org/odata/odata/v4.0/os/part2-url-conventions/odata-v4.0-os-part2-url-conventions.html>